

Machine Learning: A Probabilistic Perspective Solution Manual Version 2.1

Fangqi Li,
Shanghai Jiao Tong University,
P. R. China.

Contents

1	Preface	9
1.1	The Second Edition	9
1.2	The First Edition	12
1.3	Updating log	13
2	Probability	14
2.1	Probability are sensitive to the form of the question that was used to generate the answer	14
2.2	Legal reasoning	15
2.3	Variance of a sum	15
2.4	Bayes rule for medical diagnosis	15
2.5	The Monty Hall problem(The dilemma of three doors)	16
2.6	Conditional Independence	16
2.7	Pairwise independence does not imply mutual independence .	17
2.8	Conditional independence iff joint factorizes	18
2.9	Conditional independence	19
2.10	Deriving the inverse gamma density	20
2.11	Normalization constant for a 1D Gaussian	21
2.12	Expressing mutual information in terms of entropies	21
2.13	Mutual information for correlated normals	22

2.14	A measure of correlation	23
2.15	MLE minimizes KL divergence to the empirical distribution	23
2.16	Mean, mode, variance for the beta distribution	24
2.17	Expected value of the minimum	24
3	Generative models for discrete data	27
3.1	MLE for the Bernoulli/binomial model	27
3.2	Marginal likelihood for the Beta-Bernoulli model	28
3.3	Posterior predictive for Beta-Binomial model	30
3.4	Beta updating from censored likelihood	30
3.5	Uninformative prior for log-odds ratio	30
3.6	MLE for the Poisson distribution	31
3.7	Bayesian analysis of the Poisson distribution	32
3.8	MLE for the uniform distribution	32
3.9	Bayesian analysis of the uniform distribution	33
3.10	Taxicab problem	34
3.11	Bayesian analysis of the exponential distribution	35
3.12	MAP estimation for the Bernoulli with non-conjugate priors	36
3.13	Posterior predictive distribution for a batch of data with the Dirichlet-multinomial model	38
3.14	Posterior predictive for Dirichlet-multinomial	38
3.15	Setting the hyper-parameters I	39
3.16	Setting the beta hyper-parameters II	39
3.17	Marginal likelihood for beta-binomial under uniform prior	40
3.18	Bayes factor for coin tossing	41
3.19	Irrelevant features with naive Bayes	41
3.20	Class conditional densities for binary data	43
3.21	Mutual information for naive Bayes classifiers with binary features	43
3.22	Fitting a naive Bayesian spam filter by hand	44
4	Gaussian models	45
4.1	Uncorrelated does not imply independent	45

4.2	Uncorrelated and Gaussian does not imply independent unless jointly Gaussian	46
4.3	Correlation coefficient is between -1 and 1	47
4.4	Correlation coefficient for linearly related variables is 1 or -1	47
4.5	Normalization constant for a multidimensional Gaussian	48
4.6	Bivariate Gaussian	49
4.7	Conditioning a bivariate Gaussian	49
4.8	Whitening vs standardizing	50
4.9	Sensor fusion with known variances in 1d	50
4.10	Derivation of information form formulae for marginalizing and conditioning	51
4.11	Derivation of the NIW posterior	51
4.12	BIC for Gaussians	53
4.13	Gaussian posterior credible interval	55
4.14	MAP estimation for 1d Gaussians	56
4.15	Sequential(recursive) updating of covariance matrix	57
4.16	Likelihood ratio for Gaussians	57
4.17	LDA/QDA on height/weight data	58
4.18	Naive Bayes with mixed features	59
4.19	Decision boundary for LDA with semi tied covariances	59
4.20	Logistic regression vs LDA/QDA	60
4.21	Gaussian decision boundaries	61
4.22	QDA with 3 classes	62
4.23	Scalar QDA	62
5	Bayesian statistics	64
5.1	Proof that a mixture of conjugate priors is indeed conjugate	64
5.2	Optimal threshold on classification probability	65
5.3	Reject option in classifiers	65
5.4	More reject options	66
5.5	Newsvendor problem	67
5.6	Bayes factors and ROC curves	67
5.7	Bayes model averaging helps predictive accuracy	67
5.8	MLE and model selection for a 2d discrete distribution	68

5.9	Posterior median is optimal estimate under L1 loss	70
5.10	Decision rule for trading off FPs and FNs	71
6	Frequentist statistics	72
6.1	Pessimism of LOOCV	72
6.2	James Stein estimator for Gaussian means	73
6.3	$\hat{\sigma}_{\text{MLE}}^2$ is biased	74
6.4	Estimation of σ^2 when μ is known	74
7	Linear regression	76
7.1	Behavior of training set error with increasing sample size . .	76
7.2	Multi-output linear regression	77
7.3	Centering and ridge regression	78
7.4	MLE for σ^2 for linear regression	79
7.5	MLE for the offset term in linear regression	79
7.6	MLE for simple linear regression	79
7.7	Sufficient statistics for online linear regression	80
7.8	Bayesian linear regression in 1d with known σ^2	82
7.9	Generative model for linear regression	84
7.10	Bayesian linear regression using the g-prior	85
8	Logistic regression	87
8.1	Spam classification using logistic regression	87
8.2	Spam classification using naive Bayes	87
8.3	Gradient and Hessian of log-likelihood for logistic regression .	87
8.4	Gradient and Hessian of log-likelihood for multinomial logis- tic regression	88
8.5	Symmetric version of l2 regularized multinomial logistic re- gression	89
8.6	Elementary properties of l2 regularized logistic regression . .	90
8.7	Regularizing separate terms in 2d logistic regression	91
9	Generalized linear models and the exponential family	93
9.1	Conjugate prior for univariate Gaussian in exponential family form	93

9.2 The MVN is in the exponential family	94
10 Directed graphical models(Bayes nets)	95
10.1 Marginalizing a node in a DGM	95
10.2 Bayes Ball	96
10.3 Markov blanket for a DGM	97
10.4 Hidden variables in DGMs	98
10.5 Bayes nets for a rainy day	98
10.6 Fishing nets	99
10.7 Removing leaves in BN20 networks	100
10.8 Handling negative findings in the QMR network	101
10.9 Moralization does not introduce new independence statements	101
11 Mixture models and the EM algorithm	103
11.1 Student T as infinite mixture of Gaussian	103
11.2 EM for mixture of Gaussians	104
11.3 EM for mixtures of Bernoullis	105
11.4 EM for mixture of Student distributions	106
11.5 Gradient descent for fitting GMM	107
11.6 EM for a finite scale mixture of Gaussians	109
11.7 Manual calculation of the M step for a GMM	109
11.8 Moments of a mixture of Gaussians	110
11.9 K-means clustering by hand	111
11.10 Deriving the K-means cost function	112
11.11 Visible mixtures of Gaussians are in exponential family . . .	112
11.12 EM for robust linear regression with a Student t likelihood .	113
11.13 EM for EB estimation of Gaussian shrinkage model	114
11.14 EM for censored linear regression	115
11.15 Posterior mean and variance of a truncated Gaussian	116
12 Latent linear models	118
12.1 M-step for FA	118
12.2 MAP estimation for the FA model	120
12.3 Heuristic for assessing applicability of PCA	120

12.4 Deriving the second principal component	120
12.5 Deriving the residual error for PCA	122
12.6 Derivation of Fisher's linear discriminant	123
12.7 PCA via successive deflation	124
12.8 Latent semantic indexing	125
12.9 Imputation in a FA model	125
12.10 Efficiently evaluating the PPCA density	126
12.11 PPCA vs FA	126
13 Sparse linear models	129
13.1 Partial derivative of the RSS	129
13.2 Derivation of M-step for EB for linear regression	130
13.3 Derivation of fixed point updates for EB for linear regression	131
13.4 Marginal likelihood for linear regression	132
13.5 Reducing elastic net to lasso	132
13.6 Shrinkage in linear regression	133
13.7 Prior for the Bernoulli rate parameter in the spike and slab model	134
13.8 Deriving E step for GSM prior	135
13.9 EM for sparse probit regression with Laplace prior	135
13.10 GSM representation of group lasso	136
13.11 Projected gradient descent for l1 regularized least squares . .	137
13.12 Subderivative of the hinge loss function	138
13.13 Lower bounds to convex functions	139
14 Kernels	140
14.1 Fitting an SVM classifier by hand	140
14.2 Linear separability	140
15 Gaussian processes	141
15.1 Reproducing property	141
16 Adaptive basis function models	142
16.1 Nonlinear regression for inverse dynamics	142

17 Markov and hidden Markov models	143
17.1 Derivation of Q function for HMM	143
17.2 Two filter approach to smoothing in HMMs	143
17.3 EM for HMMs with mixture of Gaussian observations	145
17.4 EM for HMMs with tied mixtures	146
18 State space models	147
18.1 Derivation of EM for LG-SSM	147
18.2 Seasonal LG-SSM model in standard form	149
19 Undirected graphical models(Markov random fields)	150
19.1 Derivation of the log partition function	150
19.2 CI properties of Gaussian graphical models	150
19.3 Independencies in Gaussian graphical models	152
19.4 Cost of training MRFs and CRFs	152
19.5 Full conditional in an Ising model	153
20 Exact inference for graphical models	154
20.1 Variable elimination	154
20.2 Gaussian times Gaussian is Gaussian	155
20.3 Message passing on a tree	156
20.4 Inference in 2D lattice MRFs	157
21 Variational inference	159
21.1 Laplace approximation to $p(\mu, \log \sigma \mathcal{D})$ for a univariate Gaussian	160
21.2 Laplace approximation to normal-gamma	161
21.3 Variational lower bound for VB for univariate Gaussian . . .	161
21.4 Variational lower bound for VB for GMMs	163
21.5 Derivation of $\mathbb{E}[\log \pi_k]$ under a Dirichlet distribution	165
21.6 Alternative derivation of the mean field updates for the Ising model	165
21.7 Forwards vs reverse KL divergence	166
21.8 Derivation of the structured mean field updates for FHMM .	167
21.9 Variational EM for binary FA with sigmoid link	168

21.10 VB for binary FA with probit link	169
22 More variational inference*	171
23 Monte Carlo inference	172
23.1 Sampling from a Cauchy	172
23.2 Rejection sampling from a Gamma using a Cauchy proposal .	172
23.3 Optimal proposal for particle filtering with linear-Gaussian measurement model	173
24 Markov chain Monte Carlo (MCMC) inference	175
24.1 Gibbs sampling from a 2D Gaussian	175
24.2 Gibbs sampling for a 1D Gaussian mixture model	175
24.3 Gibbs sampling from the Potts model	177
24.4 Full conditionals for hierarchical model of Gaussian means .	177
24.5 Gibbs sampling for robust linear regression with a Student t likelihood	179
24.6 Gibbs sampling for probit regression	180
24.7 Gibbs sampling for logistic regression with the Student ap- proximation	180
25 Clustering	182
26 Graphical model structure learning	183
26.1 Causal reasoning in the sprinkler network	183
27 Latent variable models for discrete data	184
27.1 Partition function for an RBM	184
28 Deep learning	185

1 Preface

1.1 The Second Edition

The tide of artificial intelligence (AI) has swept and reformed so many disciplines and pushed forward the borderline of the state-of-the-art. Such a situation has resulted in the positive feedback that drives even more attention and effort into the study and research of AI, together with more unsettled regret and pities.

I have participated in this study, with equal passion that any student who has not formed an exclusive view of the world should have when engaging in a booming subject.

It has been three years and four months since I started the first edition of this solution manual. Although I have received no patron and have no intention of finding one, I received several grateful and advisory messages, from which I felt more pleased than having any of my technical papers published online. For the convenience of them, I would gladly edit this manuscript again, be the time goes back to 2017. In the second edition, I tried to be more concrete in deduction so readers can follow up easily. More graphical or numerical examples were provided to increase the overall readability. At the beginning part of each chapter/after some exercises, I left some remarks, which I thought could help.

The purpose of this manuscript is, as its first edition, to complete the textbook *Machine Learning, A Probabilistic Perspective* as a closed collection of knowledge as far as I could, and to save those who lose themselves in the ocean of deduction and symbols in ML, whom any talent mind could have become for some times in his/her course with this textbook. I hope that this manuscript can help, be it ever so little, to any reader who purportedly or accidentally finds it.

I take responsibility for any typo or mistake in β -reductions.

Fangqi Li,

Shanghai Jiao Tong University,

Shanghai, P.R.China.

January the 4th, 2021.

Contact me by: solour_lfq@sjtu.edu.cn or 1524587011@qq.com.

My homepage is at: <https://solour-lfq.github.io/>.

第二版序

这篇文档的主体由英文书写，这一方面是因为英文是学术上最泛用的语言，有利于本文档的传播；另一方面是我认为有能力阅读 MLaPP 原教材的中国学生、汉语母语学生基本上也能畅通无阻地阅读本文档中的英文。

希望中国学者和其他可以以中文为语言写作的学者一同努力，提升中文期刊、会议的质量和中文区科研院所的硬实力、影响力，让越来越多的学者乐于用中文叙述自己的观点。

第二版修订了第一版的一些文本问题，补充了第一版在习题推理上比较缺少的理念连接，同时增加了一些示例以提升可读性。文档内的所有排版错误、推导错误由我一人负责。

李方圻

上海交通大学

2021 年 1 月 4 日

邮箱: solour_lfq@sjtu.edu.cn, 1524587011@qq.com。

主页: <https://solour-lfq.github.io/>。

1.2 The First Edition

This document provides detailed solutions to almost all exercises in the textbook MLaPP from Chapter One to Chapter Twenty-one. A reader is assumed to find support from this document when he/she is teaching himself/herself an introductory or advanced course with MLaPP.

There are two classes for problems in MLaPP: theoretical ones and practical ones. We provide the solution to most theoretical problems. Practical problems, which base on a Matlab toolbox, are beyond the scope of this document.

I started reading MLaPP after selecting a machine learning course, but I failed to find any free compiled solution manuals. Although several publicly available projects have started working on it, the velocity has been too slow. In the end, I hope that readers can provide comments and opinions. Apart from correcting the wrong answers, those who good at using MATLAB, Latex typesetting or those who are willing to participate in the improvement of the document are always welcome to contact me.

22/10/2017

Fangqi Li

Munich, Germany

solour_lfq@sjtu.edu.cn

ge72bug@tum.de

1.3 Updating log

22/10/2017 First Chinese compilation.

02/03/2018 English version.

06/01/2020 The second edition begins.

2 Probability

The probability theory for ML is usually a small subset of elementary probability. More involved topics in probability beginning from Kolmogorov's theory to the martingale and Markov process is usually beyond the scope of an ordinary statistical ML textbook. Readers are encouraged to refer to Shiryaev's *Probability, 3rd edition*, whose first chapter gives a comprehensive summary of elementary probability theory. For supplementary materials, readers can refer to Thomas's *Elements of Information Theory* for a solid introduction.

2.1 Probability are sensitive to the form of the question that was used to generate the answer

Denote two children by A and B . The space of all experiment results, Ω is composed of:

$$\begin{aligned}\omega_1 : A \text{ is a girl, } B \text{ is a girl,} \\ \omega_2 : A \text{ is a boy, } B \text{ is a boy,} \\ \omega_3 : A \text{ is a girl, } B \text{ is a boy,} \\ \omega_4 : A \text{ is a boy, } B \text{ is a girl.}\end{aligned}$$

With uniform probability measure. Denote the σ -algebra on Ω as 2^Ω .

In question (a), with the knowledge *there is at least one boy*, Ω is modified into:

$$\Omega' = \{\omega_2, \omega_3, \omega_4\}.$$

The event that one child is a girl is $\{\omega_3, \omega_4\}$, whose probability is:

$$\frac{|\{\omega_3, \omega_4\}|}{|\Omega'|} = \frac{2}{3}.$$

In question (b), with the knowledge that A is a boy, the reduced experiment space is:

$$\Omega'' = \{\omega_2, \omega_4\}.$$

Then the probability that B is a girl is:

$$\frac{|\{\omega_4\}|}{|\Omega''|} = \frac{1}{2}.$$

The difference in the form of the question is reflected in that Ω is reduced to different forms and the desired events vary with our questions.

2.2 Legal reasoning

Given the assertion that the criminal has the special blood type, the space of all possibilities contains $800,000 \times \frac{1}{100} = 8,000$ samples. A sample $\omega_i \in \Omega$ denotes that the i -th person with this blood type committed the crime. Let the suspect be the j -th person with this special blood type, the event that he/she was the criminal is

$$\frac{1}{|\Omega|} = \frac{1}{8,000}.$$

For question (a): The probability that an innocent person has this blood type is almost 1%, whose opposite event is *the probability that an innocent person has another blood type*, which is 99%. This event is different from *the suspect is the criminal*.

For question (b): Justice is not measured by probability. More evidence from forensics might increase this probability to unity or reduce it to zero.

2.3 Variance of a sum

Calculate this straightforwardly:

$$\begin{aligned} \text{var}[X + Y] &= \mathbb{E}[(X + Y)^2] - \mathbb{E}^2[X + Y] \\ &= \mathbb{E}[X^2] - \mathbb{E}^2[X] + \mathbb{E}[Y^2] - \mathbb{E}^2[Y] + 2\mathbb{E}[XY] - 2\mathbb{E}[X][Y] \\ &= \text{var}[X] + \text{var}[Y] + 2\text{cov}[X, Y]. \end{aligned}$$

Using the definition of operators var , cov and the linearity of expectation should yield this result easily.

2.4 Bayes rule for medical diagnosis

Let ill and positive denote the event that you are infected and are tested positive for this disease respectively. Let health denote the opposite event

of ill. Apply Bayes's rules:

$$\begin{aligned}\Pr(\text{ill}|\text{positive}) &= \frac{\Pr(\text{ill}, \text{positive})}{\Pr(\text{positive})} \\ &= \frac{\Pr(\text{ill})\Pr(\text{positive}|\text{ill})}{\Pr(\text{ill})\Pr(\text{positive}|\text{ill}) + \Pr(\text{health})\Pr(\text{positive}|\text{health})} \\ &= 0.0098\end{aligned}$$

2.5 The Monty Hall problem(The dilemma of three doors)

The answer is (b). Use prize_i , choose_i , open_i to denote the event that the prize is in/the player chooses/the host opens the i -th box. Apply Bayes's rules:

$$\begin{aligned}\Pr(\text{prize}_1|\text{choose}_1, \text{open}_3) &= \frac{\Pr(\text{choose}_1)\Pr(\text{prize}_1)\Pr(\text{choose}_3|\text{prize}_1, \text{choose}_1)}{\Pr(\text{choose}_1)\Pr(\text{open}_3|\text{choose}_1)} \\ &= \frac{\Pr(\text{prize}_1)\Pr(\text{choose}_3|\text{prize}_1, \text{choose}_1)}{\Pr(\text{open}_3|\text{choose}_1)} \\ &= \frac{\frac{1}{3} \cdot \frac{1}{2}}{\frac{1}{3} \cdot \frac{1}{2} + \frac{1}{3} \cdot 0 + \frac{1}{3} \cdot 1} = \frac{1}{3}\end{aligned}$$

In the last step, we summarize the potential location of the prize. This is a classical example of anti-intuition results from probability.

2.6 Conditional Independence

In question (a), we have:

$$\Pr(H|e_1, e_2) = \frac{\Pr(H)\Pr(e_1, e_2|H)}{\Pr(e_1, e_2)}$$

Thus the answer is (ii).

For question (b), we have the further decomposition:

$$\Pr(H|e_1, e_2) = \frac{\Pr(H)\Pr(e_1|H)\Pr(e_2|H)}{\Pr(e_1, e_2)}$$

So both (i) and (ii) are sufficient. Moreover, we have:

$$\begin{aligned}\Pr(e_1, e_2) &= \sum_H \Pr(e_1, e_2, H) \\ &= \sum_H \Pr(H)\Pr(e_1|H)\Pr(e_2|H)\end{aligned}$$

so (iii) is sufficient as well since we can calculate $p(e_1, e_2)$ from scratch.

2.7 Pairwise independence does not imply mutual independence

Consider three boolean variables ξ_1, ξ_2, ξ_3 , ξ_1 and ξ_2 take values in 0 or 1 with equal possibility independently and $x_3 = \text{XOR}(x_1, x_2)$. It is easy to prove that x_3 is independent with x_1 or x_2 , but given both x_1 and x_2 , the value of x_3 is determined and thereby the mutual independence fails. For a detailed examination, denote the space of experiment outcomes by

$$\Omega = \{00, 01, 10, 11\},$$

with the first component denote the value of ξ_1 , the second is for ξ_2 . Then ξ_1 generates the σ -algebra:

$$\mathcal{F}_1 = \{\emptyset, \Omega, \{00, 01\}, \{10, 11\}\}.$$

ξ_2 generates:

$$\mathcal{F}_2 = \{\emptyset, \Omega, \{00, 10\}, \{01, 11\}\}.$$

ξ_3 generates:

$$\mathcal{F}_3 = \{\emptyset, \Omega, \{00, 11\}, \{01, 10\}\}.$$

One can easily check that each pair out of the triplet $\mathcal{F}_1, \mathcal{F}_2$ and \mathcal{F}_3 is a pair of independent σ -algebra, hence meets the pairwise independence.

However, each pair out of the triplet $\mathcal{F}_1, \mathcal{F}_2$ and \mathcal{F}_3 can span the entire 2^Ω . Then we would have counter examples, e.g., consider

$$A = \{00, 01\} \in \mathcal{F}_1,$$

$$B = \{11\} \in \sigma(\xi_2, \xi_3).$$

Then

$$\Pr(A \cap B) = 0 \neq \Pr(A) \cdot \Pr(B) = \frac{1}{8}.$$

Hence the mutual independence does not hold.

This example comes from cryptography. The only theoretical secure (defined by Shannon) encryption system is the one-pad cipher book. With the message denoted by m in binary code, the encryption is done by XOR m with a binary key k drawn from a cipher book. The result is ciphertext

c . From a statistical point of view, c is equally likely to be the ciphertext of any message, hence the adversary cannot break the security (since the ciphertext can be equally understood as any message, so no specific message prevails). But the encryption is a deterministic process (by using a strictly one-pad cipher book, this cipher is resistant to chosen-plaintext attack as well), so the mutual independence fails.

2.8 Conditional independence iff joint factorizes

We prove that (2.129) is tantamount to (2.130). One direction is trivial by denoting:

$$g(x, z) = p(x|z)$$

$$h(y, z) = p(y|z)$$

Conversely, we have:

$$\begin{aligned} p(x|z) &= \sum_y p(x, y|z) \\ &= \sum_y g(x, z)h(y, z) \\ &= g(x, z) \sum_y h(y, z). \end{aligned}$$

And vice versa,

$$p(y|z) = h(y, z) \sum_x g(x, z).$$

Moreover, for any z :

$$\begin{aligned} 1 &= \sum_{x,y} p(x, y|z) \\ &= \left(\sum_x g(x, z) \right) \left(\sum_y h(y, z) \right) \end{aligned}$$

Thus:

$$\begin{aligned} p(x|z)p(y|z) &= g(x, z)h(y, z) \left(\sum_x g(x, z) \right) \left(\sum_y h(y, z) \right) \\ &= g(x, z)h(y, z) \\ &= p(x, y|z) \end{aligned}$$

2.9 Conditional independence

For question (a), the antecedent $(X \perp W|Z, Y)$ means that $\forall x \in \sigma(X)$, $\forall w \in \sigma(W)$ and $\forall v \in \sigma(Y, Z)$ we have

$$\Pr(x \cap w|v) = \Pr(x|v) \cdot \Pr(w|v).$$

The antecedent $(X \perp Y|Z)$ can be translated to that $\forall x \in \sigma(X)$, $\forall y \in \sigma(Y)$ and $\forall z \in \sigma(Z)$,

$$\Pr(x \cap y|z) = \Pr(x|z) \cdot \Pr(y|z).$$

What we desire to obtain is $\forall x \in \sigma(X)$, $\forall w \in \sigma(W)$ and $\forall z \in \sigma(Z)$,

$$\Pr(x \cap w|z) = \Pr(x|z) \cdot \Pr(w|z).$$

This is correct by having v in the first equation taking values in $\sigma(Z)$ solely. Since $\sigma(Z) \subset \sigma(Y, Z)$.

For question (b), we have the premises: $\forall x \in \sigma(X)$, $\forall y \in \sigma(Y)$, $\forall z \in \sigma(Z)$ and $\forall w \in \sigma(W)$:

$$\Pr(x \cap y|z) = \Pr(x|z) \cdot \Pr(y|z).$$

$$\Pr(x \cap y|w) = \Pr(x|w) \cdot \Pr(y|w).$$

The desired result if $\forall v \in \sigma(Z, W)$,

$$\Pr(x \cap y|v) = \Pr(x|v) \cdot \Pr(y|v).$$

Let x and y be two disjoint events, z and w be another pair of disjoint, and none of $x \cap z$, $x \cap w$, $y \cap z$, $y \cap w$ is empty w.r.t. the underlying probability measure. Finally, let $v = w \cup z$. One can then check that the equation above does not hold for this setting, hence the deduction in (b) is false.

(b) is intuitively false. A straightforward example is a cryptography example: group signature with three participants. The group signature is a protocol for encryption/verification that ensures a series of security requirements including:

- Anyone participant solely cannot pass the verification.
- Two participants can pass the verification.

For example, let Alice and Bob each hold half of the secret key denoted by Z and W respectively, Y denotes the ciphertext, and X denotes the plaintext. Good group encryption would meet both antecedents in (b) but fails the conclusion. An example of a naive group signature is to use a quadratic function as the secret key:

$$f(t) = x \cdot t^2 + y \cdot t + c.$$

And provide two different points $z = (t_1, f_1)$, $w = (t_2, f_2)$ from f to Alice and Bob. Both antecedents in (b) are satisfied (by correctly translating the density) since the value of x yields no information about y . However, given both z and w then y is a deterministic function of x :

$$y = \frac{f_1 - f_2}{t_1 - t_2} - x \cdot (t_1 + t_2),$$

and the independence no longer holds. Note that a third participant is necessary to reveal the entire secret key (x, y, c) . In practice, the calculation is usually done on an algebraic field/group, e.g. the elliptic curves, to deal with the problem with the density of x, y, c which is usually not uniform in the real number field.

2.10 Deriving the inverse gamma density

According to the change of variables formula:

$$p(y) = p(x) \left| \frac{dx}{dy} \right|.$$

We have:

$$\begin{aligned} \text{IG}(y) &= \text{Ga}(x) \cdot y^{-2} \\ &= \frac{b^a}{\Gamma(a)} \left(\frac{1}{y} \right)^{(a-1)+2} e^{-\frac{b}{y}} \\ &= \frac{b^a}{\Gamma(a)} (y)^{-(a+1)} e^{-\frac{b}{y}}. \end{aligned}$$

The change of variables formula is a simplified version of the Lebesgue-Rikdon Theorem, which formally addresses the transform between probability measures defined on the same space. In the simplified version, we

take the existence of the derivative $\frac{dx}{dy}$ for granted. In the general case, such differential is obtained by taking the limit of simple functions that meet the dominance condition. The Lebesgue Theorem was developed to properly define the *differential of one probability measure w.r.t. another probability measure*. The derived general differential is usually denoted by $\frac{d\mu_1(x)}{d\mu_2(x)}$ where μ_i are probability measures.

2.11 Normalization constant for a 1D Gaussian

We have:

$$\begin{aligned} C &= \int_0^{2\pi} \int_0^\infty r \cdot \exp \left\{ -\frac{r^2}{2\sigma^2} \right\} dr d\theta \\ &= 2\pi\sigma^2 \cdot \int_0^\infty \exp \{-u\} du \\ &= 2\pi\sigma^2. \end{aligned}$$

For multivariate Gaussian, the trick is to diagonalize the covariance matrix and integrate each component independently.

2.12 Expressing mutual information in terms of entropies

We have:

$$\begin{aligned} I(X; Y) &= \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p(x)p(y)} \\ &= \sum_{x,y} p(x, y) \log \frac{p(x|y)}{p(x)} \\ &= \sum_{x,y} p(x, y) \log p(x|y) - \sum_x \left(\sum_y p(x, y) \right) \log p(x) \\ &= -H(X|Y) + H(X) \end{aligned}$$

Inversing X and Y yields another formula. One can proceed to show that $I(X; Y) = H(X) + H(Y) - H(X, Y)$.

2.13 Mutual information for correlated normals

We have:

$$\begin{aligned}
 I(X_1; X_2) &= H(X_1) - H(X_1|X_2) \\
 &= H(X_1) + H(X_2) - H(X_1, X_2) \\
 &= \frac{1}{2} \log 2\pi\sigma^2 + \frac{1}{2} \log 2\pi\sigma^2 + \frac{1}{2} \log(2\pi)^2 \sigma^4 (1 - \rho^2) \\
 &= -\frac{1}{2} \log(1 - \rho^2)
 \end{aligned}$$

Here we incorporate a comprehensive deduction on (2.138) and (2.139), which shall not be taken for granted. The differential entropy for a 1D-Gaussian with density function:

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{x^2}{2\sigma^2}\right\}$$

is

$$\begin{aligned}
 & - \int \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{x^2}{2\sigma^2}\right\} \ln\left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{-\frac{x^2}{2\sigma^2}\right\}\right) dx \\
 &= \ln(\sqrt{2\pi\sigma^2}) + \frac{1}{2\sigma^2} \mathbb{E}[x^2] \\
 &= \frac{1}{2} \ln(2\pi e\sigma^2).
 \end{aligned}$$

For the multi-dimensional case, we begin by diagonalizing the covariance matrix/decoupling the components and integrating along each independent component. Under this new set of coordinates v_1, \dots, v_d , the logarithm of the density can be decomposed into

$$C + \sum_{i=1}^d -\frac{v_i^2}{2\sigma_i^2},$$

where σ_i^2 is the i -th diagonal component in the transformed covariance matrix. The product of all diagonal components is exactly $\det\Sigma$, hence proving (2.138).

2.14 A measure of correlation

For question (a), we only have to borrow the conclusion from Exercise 2.12.:

$$\begin{aligned} r &= 1 - \frac{H(Y|X)}{H(X)} = \frac{H(X) - H(Y|X)}{H(X)} \\ &= \frac{H(Y) - H(Y|X)}{H(X)} \\ &= \frac{I(X;Y)}{H(X)} \end{aligned}$$

For question (b), we have $0 \leq r \leq 1$ in question b for $I(X;Y) \geq 0$ and $H(X|Y) \geq 0$.

For question (c), $r = 0$ iff X and Y are independent so the distance between $p(x, y)$ and $p(x) \cdot p(y)$ is zero regarding KL-divergence.

For question (d), $r = 1$ iff X is determined by, but not necessarily equal to, Y .

2.15 MLE minimizes KL divergence to the empirical distribution

Expand the KL divergence:

$$\begin{aligned} \theta &= \arg \min_{\theta} \{ \mathbb{KL}(p_{\text{emp}} || q(\theta)) \} \\ &= \arg \min_{\theta} \left\{ \mathbb{E}_{p_{\text{emp}}} \left[\log \frac{p_{\text{emp}}}{q(\theta)} \right] \right\} \\ &= \arg \min_{\theta} \left\{ -H(p_{\text{emp}}) - \sum_{\mathbf{x} \in \text{dataset}} (\log q(\mathbf{x}; \theta)) \right\} \\ &= \arg \max_{\theta} \left\{ \sum_{\mathbf{x} \in \text{dataset}} \log p(\mathbf{x}; \theta) \right\} \end{aligned}$$

We use the weak law of large numbers in the third step and drop the entropy of empirical distribution, which is independent of θ , in the last step. The other direction of optimization is $\arg \min_{\theta} \{ \mathbb{KL}(q(\theta) || p_{\text{emp}}) \}$. It contains an expectation term w.r.t. $q(\theta)$ and is harder to solve.

2.16 Mean, mode, variance for the beta distribution

Firstly, we derive the mode for beta distribution by differentiating the pdf:

$$\frac{d}{dx} x^{a-1}(1-x)^{b-1} = [(1-x)(a-1) - (b-1)x]x^{a-2}(1-x)^{b-2}$$

Setting this to zero yields:

$$\text{mode} = \frac{a-1}{a+b-2}$$

Secondly, derive the moment in beta distribution:

$$\begin{aligned} \mathbb{E}[x^N] &= \frac{1}{B(a, b)} \int x^{a+N-1}(1-x)^{b-1} dx \\ &= \frac{B(a+N, b)}{B(a, b)} \\ &= \frac{\Gamma(a+N)\Gamma(b)}{\Gamma(a+N+b)} \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \end{aligned}$$

Setting $N = 1, 2$:

$$\begin{aligned} \mathbb{E}[x] &= \frac{a}{a+b} \\ \mathbb{E}[x^2] &= \frac{a(a+1)}{(a+b)(a+b+1)} \end{aligned}$$

Where we have used the properties of the Gamma function. Finally:

$$\begin{aligned} \text{mean} &= \mathbb{E}[x] = \frac{a}{a+b} \\ \text{variance} &= \mathbb{E}[x^2] - \mathbb{E}[x]^2 = \frac{ab}{(a+b)^2(a+b+1)} \end{aligned}$$

2.17 Expected value of the minimum

Let m denote the location of the left most point, we have:

$$\begin{aligned} p(m > t) &= p([X > t] \text{ and } [Y > t]) \\ &= p(X > t)p(Y > t) \\ &= (1-t)^2 \end{aligned}$$

Therefore:

$$\begin{aligned}\mathbb{E}[m] &= \int_0^1 t \cdot p(m = t) dt \\ &= \int_0^1 p(m > t) dt \\ &= \int_0^1 (1 - t)^2 dt \\ &= \frac{1}{3}\end{aligned}$$

For the second equation, note that:

$$p(m \geq t) = \int_t^{1-t} p(m = t') dt',$$

therefore

$$\int_0^1 \int_t^{1-t} p(m = t') dt' dt = \int_0^1 \int_0^{t'} p(m = t') dt dt' = \int_0^1 t \cdot p(m = t') dt'.$$

There are perhaps more intuitive solutions to this problem, for example, plotting the value of X , Y and $\min(X, Y)$ into one graph: The height of the




Figure. 1. Exercise 2.17, P1.

cyan pyramid at (x, y) marks the value of $\min(x, y)$, so the expectation of the statistics equals the average height, in this case also the volume of the pyramid, $\frac{1}{3}$. One can also graphically compute the average distance between X and Y from the following plot: Since the average distance between X and Y is $\frac{1}{3}$, so is that between $\min(X, Y)$ and $\max(X, Y)$. Moreover, we have $\mathbb{E}[\min(X, Y) + \max(X, Y)] = \mathbb{E}[X] + \mathbb{E}[Y] = 1$, hence $\min(X, Y) = \frac{1}{3}$.




Figure. 2. Exercise 2.17, P2

However, the graphical method, although entertaining and inspiring, should not be considered as a reliable option in proving probability properties. The dependency can complicate the underlying topology (e.g., the $X - Y$ plane might have another geometry other than the Euclidean one, if X and Y are not independent), resulting in confusions and fallacies.

For example, if X is subject to a uniform distribution on $[0, 1]$ while Y is uniformly distributed in $[\max(0, X - 0.2), \min(1, X + 0.2)]$. Then the pyramid in Fig. 1 is left with the region along the diagonal line in the $X - Y$ plane. This generalization is insignificant since it only changes the region where the integral shall be done.

Consider the case where X and Y are independent random variables subject to truncated Gaussian centered at $\frac{1}{2}$ on $[0, 1]$. The plot for visualizing $\min(X, Y)$ is the same as Fig. 1. However, to compute the expectation of $\min(X, Y)$, one cannot simply calculate the volume of the pyramid since the geometry of the $X - Y$ plane has changed.

3 Generative models for discrete data

The Bayesian paradigm follows the following steps:

- Writing down the likelihood as a function of the parameters to be learned from the formulation of the problem.
- Choosing a corresponding prior distribution from the likelihood function such that the increment of data can be turned into easier operators.
- Writing down the posterior distribution as a function of the hyperparameters of the prior distribution and the observed data.
- If the task is to learn the (distribution of the) parameters: maximizing the posterior density at the observed data w.r.t. the hyperparameters.
- If the task is to predict: integrate out the posterior distribution conditioned on the observed data.

The Bayesian statistics beginning from this chapter provides an interesting and intuitive perspective into understanding and predicting the world. I learned Bayesian statistics from Bishop's *Pattern Recognition and Machine Learning*.

In conducting Bayesian analysis to examples provided in the exercises, we compute an extra term, the *evidence*. The evidence is the probability that a dataset being generated from a set of hyperparameters and is usually implicitly absorbed into the normalization term of the posterior distribution. The evidence is of practical value in empirical Bayesian, where we manage to select the optimal hyperparameters. In models such as variational inference, evidence plays an important role.

3.1 MLE for the Beroulli/binomial model

We begin with (3.11), which is the likelihood function of a collection of the outcomes in a coin-toss experiment \mathcal{D} w.r.t. the parameter θ , the probability of heads:

$$p(\mathcal{D}|\theta) = \theta^{N_1}(1 - \theta)^{N_0},$$

where N_0 and N_1 are the number of tails/heads respectively.

To decompose the differential into term-independent forms, taking logarithm:

$$\ln p(\mathcal{D}|\theta) = N_1 \ln \theta + N_0 \ln(1 - \theta).$$

Setting its derivative to zero:

$$\frac{\partial}{\partial \theta} \ln p(\mathcal{D}|\theta) = \frac{N_1}{\theta} - \frac{N_0}{1 - \theta} = 0,$$

yields (3.22):

$$\theta = \frac{N_1}{N_1 + N_0} = \frac{N_1}{N},$$

where N is the size of \mathcal{D} .

Of course one need not turn to the logarithmic field. Differentiating $p(\mathcal{D}|\theta)$ w.r.t. θ directly gives the same result. But taking a logarithm almost always simplifies the form and the deduction procedure.

3.2 Marginal likelihood for the Beta-Bernoulli model

This exercise continues the discussion of the toy coin-toss experiment, so we borrow all symbols from the exercise above. The likelihood takes the form:

$$p(\mathcal{D}|\theta) = \theta^{N_1} (1 - \theta)^{N_0}.$$

The prior distribution of θ takes the form:

$$p(\theta|a, b) = \text{Beta}(\theta|a, b) \propto \theta^{a-1} (1 - \theta)^{b-1} = C_1(a, b) \cdot \theta^{a-1} (1 - \theta)^{b-1},$$

where we adopt $C_1(a, b)$ in the hope of eliminating the ambiguity of using \propto , which, although simplifies the symbolization, results in countless errors.

The posterior distribution takes the form:

$$\begin{aligned} p(\theta|\mathcal{D}, a, b) &= \frac{p(\theta|a, b) \cdot p(\mathcal{D}|\theta, a, b)}{p(\mathcal{D}|a, b)} \\ &= \frac{p(\theta|a, b) \cdot p(\mathcal{D}|\theta)}{p(\mathcal{D}|a, b)} \\ &= \frac{C_1(a, b)}{p(\mathcal{D}|a, b)} \cdot \theta^{N_1+a-1} \cdot (1 - \theta)^{N_0+b-1}. \end{aligned}$$

The first step is the straightforward Bayesian rule, the second is the Markov property. In the last step, we adopt the equations before. Since $p(\theta|\mathcal{D}, a, b)$ should be normalized w.r.t. θ , it has to be a Beta distribution with hyperparameters $N_1 + a, N_0 + b$. We can now derive the *evidence* of \mathcal{D} w.r.t. a and b explicitly. The normalization of $p(\theta|\mathcal{D}, a, b)$ indicates that:

$$\frac{C_1(a, b)}{p(\mathcal{D}|a, b)} = C_1(N_1 + a, N_0 + b),$$

so:

$$p(\mathcal{D}|a, b) = \frac{C_1(a, b)}{C_1(N_1 + a, N_0 + b)},$$

where $C_1(\cdot, \cdot)$ is the normalization factor for the Beta distribution. This is enough for deriving (3.80) by recalling the normalization of Beta distribution. The value of $p(\mathcal{D}|a, b)$ can help us select proper hyperparameters.

As for prediction:

$$\begin{aligned} p(x_{\text{new}} = 1|\mathcal{D}, a, b) &= \int p(x_{\text{new}} = 1|\theta, a, b) \cdot p(\theta|\mathcal{D}, a, b) d\theta \\ &= \int p(x_{\text{new}} = 1|\theta) \cdot p(\theta|\mathcal{D}, a, b) d\theta \\ &= \int \theta \cdot p(\theta|\mathcal{D}, a, b) d\theta \\ &= \mathbb{E}_{\text{Beta}(N_1+a, N_0+b)}(\theta) = \frac{N_1 + a}{N_1 + a + N_0 + b}. \end{aligned}$$

The first step is the Bayesian rule, the second is Markov property. The rest is straightforward algebra.

Concretely, we calculate $p(\mathcal{D})$ where $\mathcal{D} = \{1, 0, 0, 1, 1\}$:

$$\begin{aligned} p(\mathcal{D}) &= p(x_1)p(x_2|x_1)p(x_3|x_2, x_1) \dots p(x_N|x_{N-1}, x_{N-2}, \dots, x_1) \\ &= \frac{a}{a+b} \frac{b}{a+b+1} \frac{b+2}{a+b+2} \frac{a+1}{a+b+3} \frac{a+2}{a+b+4}. \end{aligned}$$

Rename the variables $\alpha = a + b, \alpha_1 = a, \alpha_0 = b$, we have (3.83). To derive (3.80), we make use of:

$$[(\alpha_1) \dots (\alpha_1 + N_1 - 1)] = \frac{(\alpha_1 + N_1 - 1)!}{(\alpha_1 - 1)!} = \frac{\Gamma(\alpha_1 + N_1)}{\Gamma(\alpha_1)}.$$

3.3 Posterior predictive for Beta-Binomial model

Straightforward algebra (recall (2.61)) gives:

$$\begin{aligned} \text{Bb}(1|\alpha'_1, \alpha'_0, 1) &= \frac{B(\alpha'_1 + 1, \alpha'_0)}{B(\alpha'_1, \alpha'_0)} \\ &= \frac{\Gamma(\alpha'_0 + \alpha'_1)}{\Gamma(\alpha'_0 + \alpha'_1 + 1)} \frac{\Gamma(\alpha'_1 + 1)}{\Gamma(\alpha'_1)} \\ &= \frac{\alpha'_1}{\alpha'_1 + \alpha'_0}. \end{aligned}$$

The hint provided in the textbook is incorrect by mistaking

$$\Gamma(a) = (a - 1) \cdot \Gamma(a - 1)$$

for

$$\Gamma(a) = a \cdot \Gamma(a - 1).$$

3.4 Beta updating from censored likelihood

The derivation is straightforward:

$$\begin{aligned} p(\theta, X < 3) &= p(\theta) \cdot p(X < 3|\theta) \\ &= p(\theta) \cdot \left(\sum_{i=0}^2 p(X = i|\theta) \right) \\ &= \text{Beta}(\theta|1, 1) \cdot \left(\sum_{i=0}^2 \text{Bin}(i|5, \theta) \right), \end{aligned}$$

with

$$\text{Bin}(m|n, \theta) = \binom{n}{m} \cdot \theta^m \cdot (1 - \theta)^{n-m}$$

is the probability that m heads appear in n times of experiments with the probability of head θ . The posterior distribution over θ , in this case, becomes much more involved.

3.5 Uninformative prior for log-odds ratio

Since:

$$\phi = \log \frac{\theta}{1 - \theta}.$$

By using change of variables formula:

$$p(\theta) = p(\phi) \cdot \left| \frac{d\phi}{d\theta} \right| \propto \frac{1}{\theta(1-\theta)},$$

hence

$$p(\theta) = \text{Beta}(\theta|0, 0).$$

That is to say, we can generate samples subject to a Beta distribution by transforming samples drawn from a uniform distribution. This trick is of significant practical value. Direct sampling from a Beta distribution requires inverting its cumulative probability function, which involves too much computation.

3.6 MLE for the Poisson distribution

The Poisson distribution plays a central role in the stochastic process, e.g., the queueing theory. If data are assumed to be generated from a similar process then the Bayesian analysis of the Poisson distribution derived in this exercise and the next can be applied directly. The likelihood of data for a Poisson distribution is (assuming i.i.d.):

$$p(\mathcal{D}|\lambda) = \prod_{n=1}^N \text{Poi}(x_n|\lambda) = \exp(-\lambda N) \cdot \lambda^{\sum_{n=1}^N x_n} \cdot \frac{1}{\prod_{n=1}^N x_n!}.$$

Setting the derivative of the likelihood w.r.t. λ to zero:

$$\frac{\partial}{\partial \lambda} p(\mathcal{D}|\lambda) = \frac{\exp(-\lambda N) \cdot \lambda^{(\sum_{n=1}^N x_n)-1}}{\prod_{n=1}^N x_n!} \left\{ -N\lambda + \sum_{n=1}^N x_n \right\}.$$

Thus:

$$\lambda_{\text{MLE}} = \frac{\sum_{n=1}^N x_n}{N}.$$

The formulation could be made easier by taking logarithm (since the Poisson distribution can be considered an element of the exponential family as well):

$$\log p(\mathcal{D}|\lambda) = -\lambda \cdot N + \left(\sum_{n=1}^N x_n \right) \cdot \log \lambda,$$

where we have omitted the term independent of λ .

3.7 Bayesian analysis of the Poisson distribution

The conjugate prior for the Poisson distribution is the Gamma distribution:

$$\text{Ga}(\lambda|a, b) = \frac{b^a}{\Gamma(a)} \cdot \lambda^{a-1} \cdot \exp(-\lambda \cdot b).$$

The posterior for a Bayesian Poisson model reads:

$$\begin{aligned} p(\lambda|\mathcal{D}, a, b) &= \frac{p(\mathcal{D}|\lambda) \cdot p(\lambda|a, b)}{p(\mathcal{D}|a, b)} \\ &= \frac{b^a}{\prod_{n=1}^N x_n! \cdot p(\mathcal{D}|a, b) \cdot \Gamma(a)} \cdot \lambda^{a+\sum_{n=1}^N x_n-1} \cdot \exp(-\lambda \cdot (N+b)) \\ &= \text{Ga}(a + \sum_{n=1}^N x_n, N+b), \end{aligned}$$

in which the last step follows the normalization condition. We now have the evidence:

$$p(\mathcal{D}|a, b) = \frac{b^a \cdot \Gamma(a + \sum_{n=1}^N x_n)}{\prod_{n=1}^N x_n! \cdot (N+b)^{a+\sum_{n=1}^N x_n} \cdot \Gamma(a)}.$$

Finally, be a and b approximate zero, the posterior mean approaches $\frac{\sum_{n=1}^N x_n}{N}$, the same as the MLE, i.e., $a = 0, b = 0$ is a non-informative prior. One should note that this property does not hold for all Bayesian analysis, setting all hyperparameters to zero does not necessarily gracefully degenerate the posterior mean to the MLE. Since the names and definitions of those symbols might differ.

3.8 MLE for the uniform distribution

The Bayesian analysis for the uniform distribution seems to be of less significance since uniform distribution appears to appear less frequently than other continuous distributions. But the exercises remain good introductory examples.

The likelihood for the uniform distribution is a truncated function, whose domain is $[-a, a]$, so we must have $a \geq \max_i \{|x_i| \in \mathcal{D}\}$. Then the likelihood looks like:

$$p(\mathcal{D}|a) = \prod_{i=1}^n \frac{1}{2a},$$

or generally:

$$p(\mathcal{D}|a) = \mathbb{I}[a \geq \max_i \{|x_i| \in \mathcal{D}\}] \cdot (2a)^{-n}.$$

For question (a), in order to maximize this value with $a \geq \max_n \{|x_n| \in \mathcal{D}\}$, the outcome is:

$$a_{\text{MLE}} = \max_i \{|x_i| \in \mathcal{D}\}.$$

For question (b), if $|x_{n+1}| > \max_{i=1}^n \{|x_i|\}$ then $p(x_{n+1})$ is zero. Otherwise the probability is $\frac{1}{2 \cdot a_{\text{MLE}}}$.

For question (c), we believe that MLE for the uniform distribution is *not fluent enough* since when x_{n+1} passes $\pm \max_{i=1}^n \{|x_i|\}$, the predicted probability drops as a step function, which is undesired for a continuous distribution.

3.9 Bayesian analysis of the uniform distribution

The conjugate prior for uniform distribution is the Pareto distribution, whose density function is defined by:

$$p(\theta|K, b) = \text{Pa}(\theta|K, b) = K \cdot b^K \cdot \theta^{-(K+1)} \cdot \mathbb{I}[\theta \geq b].$$

Let $m = \max \{|x_i|\}_{i=1}^n$, the joint distribution of θ and \mathcal{D} is:

$$\begin{aligned} p(\theta, \mathcal{D}|K, b) &= p(\theta|K, b) \cdot p(\mathcal{D}|\theta) \\ &= K \cdot b^K \cdot \theta^{-(K+1)} \cdot \mathbb{I}[\theta \geq b] \cdot \mathbb{I}[\theta \geq m] \cdot (\theta)^{-n} \\ &= K \cdot b^K \cdot \theta^{-(K+n+1)} \cdot \mathbb{I}[\theta \geq \max(b, m)]. \end{aligned}$$

Now $p(\theta, \mathcal{D}|K, b) = p(\mathcal{D}|K, b) \cdot p(\theta|\mathcal{D}, K, b)$, hence the posterior distribution depends on θ through:

$$\theta^{-(K+n+1)} \cdot \mathbb{I}[\theta \geq \max(b, m)].$$

So the posterior distribution is another Pareto distribution with hyperparameters $K + n, \max(b, \max \{|x_i|\}_{i=1}^n)$.

The evidence is computed from the Bayesian rule:

$$\begin{aligned} p(\mathcal{D}|K, b) &= \int_0^\infty p(\mathcal{D}, \theta|K, b) d\theta \\ &= \int_{\max(b, m)}^\infty \frac{K \cdot b^K}{\theta^{K+n+1}} d\theta. \end{aligned}$$

The rest is trivial calculus.

3.10 Taxicab problem

Some similar entertaining problems are *guessing the number of piano tuners from the average time for a tuner to arrive in one guest's house*, etc.

For question (a), we begin with hyperparameters $K = 0$, $b = 0$, which is improper since the Pareto distribution cannot normalize. With $\mathcal{D} = \{100\}$, we have the posterior distribution another Pareto distribution with $K = 1$ and $b = 100$, i.e.,

$$p(\theta|\mathcal{D}) = \frac{100}{\theta^2} \cdot \mathbb{I}[\theta \geq 100].$$

For question (b), we firstly derive the distribution of the taxi index:

$$\begin{aligned} p(x|\mathcal{D}, K, b) &= \int_0^\infty p(x, \theta) d\theta \\ &= \int_0^\infty p(x|\theta) \cdot p(\theta|\mathcal{D}, K, b) d\theta \\ &= \int_{100}^\infty \mathbb{I}[x \leq \theta] \cdot \frac{1}{\theta} \cdot \frac{100}{\theta^2} d\theta \\ &= \int_{\max(x, 100)}^\infty \frac{100}{\theta^3} d\theta \\ &= 50 \cdot \max(x, 100)^{-2}, \end{aligned}$$

whose plots look very much similar to that of electrical potential along an axis that penetrates the center of a conductor sphere with radius 100, through declines exponentially faster.

The posterior mode of x is any number in $[0, 100]$.

The posterior mean of x is:

$$\mathbb{E}(x) = \sum_{x=0}^{100} \frac{x}{200} + \sum_{x=100}^{\infty} \frac{50}{x},$$

whose second term diverges, so the posterior mean does not exist.

The posterior median is 99.5, since:

$$\sum_{x=0}^{99} \frac{1}{200} < 0.5 < \sum_{x=0}^{100} \frac{1}{200}.$$

Question (c) is identical to (b), as we have adopted a Bayesian treatment for (b).

For question (d), we have:

$$\begin{aligned} p(x = 100|\mathcal{D}, K, b) &= \frac{1}{200}, \\ p(x = 50|\mathcal{D}, K, b) &= \frac{1}{200}, \\ p(x = 150|\mathcal{D}, K, b) &= \frac{1}{450}. \end{aligned}$$

For question (e), we might adopt better K and b with expert knowledge and collect more samples.

3.11 Bayesian analysis of the exponential distribution

The exponential distribution is also crucial for the queueing theory. The log-likelihood for an exponential distribution with density:

$$p(x|\theta) = \theta \cdot \exp(-\theta \cdot x)$$

is:

$$\ln p(\mathcal{D}|\theta) = N \cdot \ln \theta - \theta \cdot \sum_{n=1}^N x_n,$$

whose derivative is:

$$\frac{\partial}{\partial \theta} \ln p(\mathcal{D}|\theta) = \frac{N}{\theta} - \sum_{n=1}^N x_n$$

Thus for question (a), we have:

$$\theta_{\text{MLE}} = \frac{N}{\sum_{n=1}^N x_n}.$$

For question (b), $\theta_{\text{MLE}} = 5$.

For question (c), we begin with an exponential prior distribution:

$$p(\theta|\lambda) = \lambda \cdot \exp(-\lambda \cdot \theta),$$

whose expectation is:

$$\int_0^\infty \lambda \cdot \theta \cdot \exp(-\lambda \cdot \theta) d\theta.$$

Integration by parts (or resort to the normalization term of the Gamma distribution) yields:

$$\mathbb{E}(\theta) = \frac{1}{\lambda}.$$

So $\hat{\lambda} = 3$.

For question (d), the posterior distribution is:

$$\begin{aligned} p(\theta|\mathcal{D}, \lambda) &= \frac{p(\mathcal{D}|\theta) \cdot p(\theta|\lambda)}{p(\mathcal{D}|\lambda)} \\ &= \frac{1}{p(\mathcal{D}|\lambda)} \cdot \theta^N \cdot \exp(-\theta \cdot \sum_{n=1}^N x_n) \cdot \lambda \cdot \exp(-\lambda \cdot \theta) \\ &= \frac{\theta^N \cdot \lambda}{p(\mathcal{D}|\lambda)} \cdot \exp\left(-\theta \cdot \left(\lambda + \sum_{n=1}^N x_n\right)\right). \end{aligned}$$

Hence the posterior is a Gamma distribution with hyperparameters:

$$\begin{aligned} a &= N + 1, \\ b &= \lambda + \sum_{n=1}^N x_n. \end{aligned}$$

The evidence is given by: $\frac{\lambda \cdot \Gamma(a)}{b^a}$, a function of λ and \mathcal{D} . Hence the exponential distribution is not the conjugate distribution of itself, answering question (e).

For question (f), the posterior mean is the mean of the Gamma distribution:

$$\frac{a}{b} = \frac{N + 1}{\lambda + \sum_{n=1}^N x_n}.$$

Compared with the MLE, the posterior mean has additional terms for both the numerator and the denominator as basic knowledge when N is relatively small. The influence of using this prior is tantamount to introducing a prior sample with value λ .

3.12 MAP estimation for the Bernoulli with non-conjugate priors

For question (a), we adopt the different prior:

$$p(\theta) = \begin{cases} 0.5, & \text{if } \theta = 0.5, \\ 0.5, & \text{if } \theta = 0.4, \end{cases}$$

The posterior distribution now reads:

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}|\theta) \cdot p(\theta)}{p(\mathcal{D})},$$

whose support is $\{0.4, 0.5\}$, so the MAP is:

$$\max_{\theta \in \{0.4, 0.5\}} \{ \theta^{N_1} \cdot (1 - \theta)^{N_0} \}.$$

For question (b), it is intuitive that the non-conjugate has better performance when N is small. But the conjugate Bayesian method prevails with N grows. For a solid verification, consider the case where N is large, the probability that $\frac{N_1}{N}$ deviates ϵ from 0.41 can be bounded by the Chernoff bounding. Let $\{\xi_n\}_{n=1}^N$ be a collection of i.i.d. Bernoulli random variables with distribution:

$$\xi_n = \begin{cases} 1, & \text{with probability } 0.41, \\ 0, & \text{with probability } 0.59, \end{cases}$$

Denote $X = \sum_{n=1}^N \xi_n$ as the random variable marks their summation. Then:

$$\begin{aligned} \Pr(X \geq N \cdot (0.41 + \epsilon)) &= \Pr(e^{\lambda \cdot X} \geq e^{N\lambda(0.41 + \epsilon)}) \\ &\leq \frac{\mathbb{E}[e^{\lambda \cdot X}]}{e^{N\lambda(0.41 + \epsilon)}} \\ &= \frac{(\mathbb{E}[e^{\lambda \cdot \xi_0}])^N}{e^{N\lambda(0.41 + \epsilon)}} \\ &= \frac{(\mathbb{E}[e^{\lambda \cdot \xi_0}])^N}{e^{N\lambda(0.41 + \epsilon)}} \\ &= \frac{(0.41e^\lambda + 0.59)^N}{e^{N\lambda(0.41 + \epsilon)}}. \end{aligned}$$

Where λ can be an arbitrary positive number. The probability that X exceeds $N \cdot (0.41 + \epsilon)$, denoted by P_1 is bounded by the lower bound of the last line in the deduction above. With $N = 1000$, $\epsilon = 0.05$, the probability is numerically bounded by 0.00602. The other side of error $\Pr(X \leq N \cdot (0.41 - \epsilon))$, whose probability is P_2 , can be derived in a similar way. The probability that the Bayesian way is dominated by the non-uniform prior is no higher than $P_1 + P_2$. Taking $\epsilon \rightarrow 0.1$ and $N \rightarrow \infty$, this bound remains negligible.

3.13 Posterior predictive distribution for a batch of data with the dirichlet-multinomial model

The likelihood for Dirichlet-multinomial model is:

$$p(\mathcal{D}|\theta) = \prod_{k=1}^K \theta_k^{N_k^{\text{old}}},$$

following the symbols defined in the textbook. The conjugate prior is the Dirichlet distribution:

$$p(\theta|\alpha) = \frac{1}{B(\alpha)} \cdot \prod_{k=1}^K \theta_k^{\alpha_k - 1},$$

where θ is a K -dimension simplex. The (3.37) in the textbook mistake θ for \mathbf{x} .

The posterior distribution is another Dirichlet distribution with update:

$$\alpha_k + N_k^{\text{old}} \leftarrow \alpha_k.$$

To predict a new batch of data $\tilde{\mathcal{D}}$, we begin with one sample $x \in \tilde{\mathcal{D}}$:

$$\begin{aligned} p(x = k|\mathcal{D}, \alpha) &= \int_{\theta} p(x = k|\theta) \cdot p(\theta|\mathcal{D}, \alpha) d\theta \\ &= \mathbb{E}_{\text{Dir}}[\theta_k], \end{aligned}$$

where the expectation is computed w.r.t. the posterior Dirichlet distribution, hence is:

$$\frac{\alpha_k + N_k^{\text{old}}}{\sum_{t=1}^K \alpha_t + N_t^{\text{old}}}.$$

Finally,

$$\begin{aligned} p(\tilde{\mathcal{D}}|\mathcal{D}, \alpha) &= \prod_{x \in \tilde{\mathcal{D}}} p(x|\mathcal{D}, \alpha) \\ &= \prod_{k=1}^K \left(\frac{\alpha_k + N_k^{\text{old}}}{\sum_{t=1}^K \alpha_t + N_t^{\text{old}}} \right)^{N_k^{\text{new}}}. \end{aligned}$$

3.14 Posterior predictive for Dirichlet-multinomial

For question (a). In this concrete case we have $K = 27$, $N = 2,000$, any component of α be 10, and $N_e^{\text{old}} = 260$. To derive $p(x_{2001} = e|\mathcal{D})$, we

resort to the deduction in Exercise 3.13:

$$p(x_{2001} = e | \mathcal{D}) = \frac{10 + 260}{10 \times 27 + 2000} \approx 0.1189.$$

For question (b), the independence between characters is still ignored, bringing no significant change to the computation:

$$p(x_{2001} = p, x_{2002} = a | \mathcal{D}) = \frac{10 + 87}{2270} \cdot \frac{10 + 100}{2270} \approx 0.0021.$$

3.15 Setting the hyper-parameters I

Solve for:

$$\begin{aligned} \frac{\alpha_1}{\alpha_1 + \alpha_2} &= m, \\ \frac{\alpha_1 \cdot \alpha_2}{(\alpha_1 + \alpha_2)^2 \cdot (\alpha_1 + \alpha_2 + 1)} &= v. \end{aligned}$$

We have:

$$\begin{aligned} \alpha_2 &= \frac{m \cdot (1 - m)^2}{v} + m - 1, \\ \alpha_1 &= \alpha_2 \cdot \frac{m}{1 - m}. \end{aligned}$$

3.16 Setting the beta hyper-parameters II

```

1 import math
2 m=0.15
3 l=0.05
4 u=0.3
5 MC=1000
6 delta=(u-l)/MC
7 def pm(a2):
8     a1=a2*m/(1-m)
9     pivot=1
10    mass=0
11    B=math.gamma(a1)*math.gamma(a2)/math.gamma(a1+a2)
12    for i in range(MC):
13        pivot=pivot+delta
14        mass=mass+pivot**(a1-1)*(1-pivot)**(a2-1)
15    mass=mass*delta/B
16    return mass

```




Figure. 3. Exercise 3.16.

The result of which is better demonstrated through the graph: So the optimal choice is $\alpha = 26$. This is tantamount to adopt 32 extra samples.

3.17 Marginal likelihood for beta-binomial under uniform prior

The marginal likelihood is given by:

$$p(N_1|N) = \int_0^1 p(N_1, \theta|N) d\theta = \int_0^1 p(N_1|\theta, N) \cdot p(\theta) d\theta.$$

Plug in:

$$p(N_1|\theta, N) = \text{Bin}(N_1|\theta, N),$$

$$p(\theta) = \text{Beta}(\theta|1, 1).$$

Thus:

$$\begin{aligned} p(N_1|N) &= \int_0^1 \binom{N}{N_1} \cdot \theta^{N_1} \cdot (1 - \theta)^{N - N_1} d\theta \\ &= \binom{N}{N_1} \cdot B(N_1 + 1, N - N_1 + 1) \\ &= \frac{N!}{N_1! \cdot (N - N_1)!} \cdot \frac{N_1! \cdot (N - N_1)!}{(N + 1)!} \\ &= \frac{1}{N + 1}. \end{aligned}$$

Where B is the regularizer for a Beta distribution:

$$B(a, b) = \frac{\Gamma(a) \cdot \Gamma(b)}{\Gamma(a + b)}.$$

The physics behind this setting is that if no prior information is introduced then all $N + 1$ possibilities are equally likely to appear.

3.18 Bayes factor for coin tossing

The Bayes factor for hypothesis test is defined by:

$$\text{BF}_{1,0} = \frac{p(\text{data}|H_1)}{p(\text{data}|H_0)},$$

where H_0 is the null hypothesis.

We have:

$$p(\text{data}|H_0) = \text{Bin}(9|0.5, 10) = \binom{10}{9} \cdot 0.5^{10} \approx 0.00977.$$

And

$$p(\text{data}|H_1) = \frac{1}{10 + 1} \approx 0.09091,$$

according to Exercise 3.17. The Bayes factor is approximately 9.3.

When $N = 100$ and $N_1 = 90$, the Bayes factor is:

$$\frac{\frac{1}{\binom{100}{90} \cdot 0.5^{100}}}{\frac{1}{101^{11}}} > \frac{2^{100}}{101^{11}} \approx 113622530.$$

When $\frac{N}{N_1}$ remains a constant deviated from 0.5, the larger N is, the more likely that the coin is biased. This is an intuitive conclusion from the law of large numbers.

3.19 Irrelevant features with naive Bayes

The log-likelihood is defined by:

$$\log p(\mathbf{x}_i|c, \theta) = \sum_{w=1}^W x_{iw} \cdot \log \frac{\theta_{cw}}{1 - \theta_{cw}} + \sum_{w=1}^W \log(1 - \theta_{cw}).$$

In a succinct way:

$$\log p(\mathbf{x}_i|c, \theta) = \phi(\mathbf{x}_i)^T \beta_c,$$

where:

$$\phi(\mathbf{x}_i) = (\mathbf{x}_i, 1)^T,$$

$$\beta_c = \left(\log \frac{\theta_{c1}}{1 - \theta_{c1}}, \dots, \sum_{w=1}^W \log(1 - \theta_{cw}) \right)^T.$$

For question (a):

$$\begin{aligned} \log \frac{p(c=1|\mathbf{x}_i)}{p(c=2|\mathbf{x}_i)} &= \log \frac{p(c=1) \cdot p(\mathbf{x}_i|c=1)}{p(c=2) \cdot p(\mathbf{x}_i|c=2)} \\ &= \log \frac{p(\mathbf{x}_i|c=1)}{p(\mathbf{x}_i|c=2)} \\ &= \phi(\mathbf{x}_i)^T (\beta_1 - \beta_2). \end{aligned}$$

For question (b), with:

$$\log \frac{p(c=1|\mathbf{x}_i)}{p(c=2|\mathbf{x}_i)} = \log \frac{p(c=1)}{p(c=2)} + \phi(\mathbf{x}_i)^T (\beta_1 - \beta_2),$$

a word w will not affect this posterior measure as long as:

$$x_{iw}(\beta_{1,w} - \beta_{2,w}) = 0$$

Hence if:

$$\theta_{c=1,w} = \theta_{c=2,w},$$

then it cannot affect the classification decision. That is to say, w appear in class 1 and 2 with the same frequency.

For question (c), we have:

$$\hat{\theta}_{1,w} = 1 - \frac{1}{2 + N_1},$$

$$\hat{\theta}_{2,w} = 1 - \frac{1}{2 + N_2}.$$

They are different when $N_1 \neq N_2$ so the bias effect remains. However, this bias reduces when N grows large.

For question (d), using information theory would be a solid option.

3.20 Class conditional densities for binary data

For question (a), we have:

$$p(\mathbf{x}|y = c) = \prod_{i=1}^D p(x_i|y = c, x_1, \dots, x_{i-1}).$$

The number of parameter in this case is:

$$C \cdot \sum_{i=0}^{D-1} 2^i = C \cdot (2^{D+1} - 2) = \mathcal{O}(C \cdot 2^D).$$

For question (b) and (c), the overfitting is generally assumed to decline when N grows. The dependence within the variables generally hinder generalization when N is small, but it could correctly capture the dependence, be it exist.

For question (d), fitting each parameter for the naive Bayes model requires averaging all components in the samples belong to one class, which is of complexity order $\mathcal{O}(N)$. The total complexity is thus of order $\mathcal{O}(C \cdot D \cdot N)$. For the full Bayes classifier, the complexity is of order $\mathcal{O}(C \cdot 2^D \cdot N)$.

For question (e), the cost of inference in the naive Bayes model is $\mathcal{O}(C \cdot D)$ for each sample. While that for full Bayes model is $\mathcal{O}(C \cdot 2^D)$.

For question (f), we have:

$$p(y|\mathbf{x}_v, \theta) \propto p(y|\theta) \cdot p(\mathbf{x}_v|\theta) \propto \sum_{\mathbf{x}_h} p(\mathbf{x}_v, \mathbf{x}_h|\theta),$$

where we assumed a uniform prior on all classes since it is not the bottleneck of complexity. The complexity for the naive model is then:

$$\mathcal{O}(C \cdot v \cdot 2^h),$$

while that for the full model is:

$$\mathcal{O}(C \cdot 2^v \cdot 2^h).$$

3.21 Mutual information for naive Bayes classifiers with binary features

By definition:

$$I(X; Y) = \sum_{x_j} \sum_y p(x_j, y) \log \frac{p(x_j, y)}{p(x_j) \cdot p(y)}.$$

For binary features, the value of x_j is either zero or one. Given $\pi_c = p(y = c)$, $\theta_{jc} = p(x_j = 1|y = c)$, $\theta_j = p(x_j = 1)$, we have the mutual information between x_j and Y be:

$$\begin{aligned} I_j &= \sum_c p(x_j = 1, c) \log \frac{p(x_j = 1, c)}{p(x_j = 1) \cdot p(c)} \\ &\quad + \sum_c p(x_j = 0, c) \log \frac{p(x_j = 0, c)}{p(x_j = 0) \cdot p(c)} \\ &= \sum_c \pi_c \theta_{jc} \log \frac{\theta_{jc}}{\theta_j} + (1 - \theta_{jc}) \pi_c \log \frac{1 - \theta_{jc}}{1 - \theta_j}, \end{aligned}$$

which ends in (3.76).

3.22 Fitting a naive Bayesian spam filter by hand

$$\begin{aligned} \theta_{\text{spam}} &\text{ is } \frac{3}{3+4}. \\ \theta_{\text{secret}|\text{spam}} &\text{ is } \frac{2}{3}. \\ \theta_{\text{secret}|\text{non-spam}} &\text{ is } \frac{1}{4}. \\ \theta_{\text{sports}|\text{non-spam}} &\text{ is } \frac{1}{2}. \\ \theta_{\text{dollar}|\text{spam}} &\text{ is } \frac{1}{3}. \end{aligned}$$

4 Gaussian models

Though being a family of models for continuous variables, traditional Gaussian models do not cast a more important role than models covered in the chapter before. Since in most scenarios, the assumption that the data are subject to a normal distribution is unrealistic.

However, Gaussian models are crucial for latent space analysis. Even for complex objects such as images or videos, the assumption that they can be encoded into features under a Gaussian distribution turns out to be reliable. Moreover, Gaussian models pave the way to general variational inference and Gaussian process, an important kernel machine.

For the reason above one cannot overestimate the significance of Gaussian models, be their history so long. Most of the mathematical difficulties with the Gaussian model have been covered in sections (marked with stars) in the textbook.

4.1 Uncorrelated does not imply independent

The mean for Y is:

$$\int_{-1}^1 X^2 dX = \mathbb{E}[X^2].$$

Calculate the covariance of X and Y :

$$\begin{aligned} \text{cov}(X, Y) &= \int \int (X - \mathbb{E}(X)) \cdot (Y - \mathbb{E}(Y)) \cdot p(X, Y) dX dY \\ &= \int_{-1}^1 X(X^2 - \mathbb{E}[X^2]) dX = 0, \end{aligned}$$

whose value is zero since we are integrating an odd function in range $[-1, 1]$, hence:

$$\rho(X, Y) = \frac{\text{cov}(X, Y)}{\sqrt{\text{var}(X)\text{var}(Y)}} = 0.$$

Independence is a much stronger condition than uncorrelation. The former exerts constraints on the σ -algebra that random variables generate while the latter only regulates the value of the expectation of a new random

variable. Decomposition $p(X, Y) = p(X) \cdot p(Y)$ is sufficient for reducing the covariance to zero, but not necessary.

4.2 Uncorrelated and Gaussian does not imply independent unless jointly Gaussian

For question (a). The p.d.f. for Y is:

$$p(Y \in [a, a+da]) = 0.5 \cdot p(X \in [a, a+da]) + 0.5 \cdot p(X \in [-a-da, -a]) = p(X \in [a, a+da]),$$

since X is symmetric. So Y subject to a normal distribution $(0, 1)$.

For question (b), we have:

$$\begin{aligned} \text{cov}(X, Y) &= \mathbb{E}(XY) - \mathbb{E}(X) - \mathbb{E}(Y) \\ &= \mathbb{E}_W(\mathbb{E}(XY|W)) - 0 \\ &= 0.5 \cdot \mathbb{E}(X^2) + 0.5 \cdot \mathbb{E}(-X^2) = 0. \end{aligned}$$

So they are uncorrelated.

To disprove dependence (in case of confusion), let:

$$a = \Phi^{-1}\left(\frac{1}{4}\right),$$

where Φ is the c.d.f of X , i.e.:

$$\int_{-\infty}^a \mathcal{N}(x|0, 1)dx = \frac{1}{4}.$$

Let $R_1 = (-\infty, a]$, $R_2 = (a, 0]$. The space of experiment results for $X \times Y$ is \mathcal{R}^2 . Let $R_1 \times R_2$ be a Borel set in \mathcal{R}^2 . Be X and Y independent, its probability measure should be $\frac{1}{16}$. However, when $X \in R_1$, it is impossible for Y to take a value from R_2 . Hence the independency fails.

The rule of iterated expectation is but the Bayes rule:

$$\begin{aligned} \mathbb{E}[XY] &= \int_X \int_Y XY \cdot p(X, Y) dX dY \\ &= \int_X \int_Y XY \cdot \left(\int_W p(X, Y, W) dW \right) dX dY \\ &= \int_W \left(\int_X \int_Y XY \cdot p(X, Y|W) \right) p(W) dW. \end{aligned}$$

4.3 Correlation coefficient is between -1 and 1

Without loss of generality, assume $\mathbb{E}[X] = \mathbb{E}[Y] = 0$. The statement:

$$-1 \leq \rho(X, Y) \leq 1,$$

is equal to

$$|\rho(X, Y)| \leq 1.$$

Hence we are to prove:

$$|\text{cov}(X, Y)|^2 \leq \text{var}(X) \cdot \text{var}(Y)$$

Which can be drawn straightforwardly from Cauchy–Schwarz inequality. Let

$$g(t) = t^2 \cdot \text{var}(X) + 2t \cdot \text{cov}(X, Y) + \text{var}(Y) = \mathbb{E}[(tX + Y)^2] \geq 0.$$

Taking g 's discriminator finishes the proof.

4.4 Correlation coefficient for linearly related variables is 1 or -1

When $Y = aX + b$:

$$\mathbb{E}(Y) = a \cdot \mathbb{E}(X) + b,$$

$$\text{var}(Y) = a^2 \cdot \text{var}(X).$$

Therefore:

$$\begin{aligned} \text{cov}(X, Y) &= \mathbb{E}(XY) - \mathbb{E}(X) \cdot \mathbb{E}(Y) \\ &= a \cdot \mathbb{E}(X^2) + b \cdot \mathbb{E}(X) - a \cdot \mathbb{E}^2(X) - b \cdot \mathbb{E}(X) \\ &= a \cdot \text{var}(X). \end{aligned}$$

Meanwhile:

$$\text{var}(X) \cdot \text{var}(Y) = a^2 \cdot \text{var}(X).$$

These two are sufficient to derive:

$$\rho(X, Y) = \frac{a}{|a|}.$$

4.5 Normalization constant for a multidimensional Gaussian

Assume $\mu = \mathbf{0}$ w.l.o.g. If the covariance matrix already takes a diagonal form:

$$\Sigma = \begin{pmatrix} \lambda_1^{-1} & \cdots \\ \cdots & \cdots \\ \cdots & \lambda_d^{-1} \end{pmatrix},$$

then:

$$\begin{aligned} \int \exp\left(-\frac{1}{2}\mathbf{x}^T \Sigma \mathbf{x}\right) d\mathbf{x} &= \int \exp\left(-\frac{1}{2}\left(\sum_{i=1}^d \frac{x_i^2}{\lambda_i}\right)\right) d\mathbf{x} \\ &= \prod_{i=1}^d \int \exp\left(-\frac{x_i^2}{2 \cdot \lambda_i}\right) dx_i \\ &= (2\pi\lambda_i)^{-\frac{d}{2}}. \end{aligned}$$

Plugging in $|\Sigma| = \prod_{i=1}^d \lambda_i^{-1}$ yields the desired normalization constant. In the second equation, using the distribution law (though somewhat intimidating).

For the general case, we begin by diagonalizing Σ into:

$$\Sigma = U^T \Lambda U,$$

where Λ is a diagonal matrix with components $\lambda_1^{-1} \cdots \lambda_d^{-1}$ and U is a orthogonal matrix. The integral now becomes:

$$\int \exp\left(-\frac{1}{2}(U\mathbf{x})^T \Lambda (U\mathbf{x})\right) d\mathbf{x}.$$

Since $|U| = 1$ uniformly, we can directly rewrite the integral into:

$$\int \exp\left(-\frac{1}{2}\mathbf{u}^T \Lambda \mathbf{u}\right) d\mathbf{u}.$$

The rest is repeating the diagonal case.

4.6 Bivariate Gaussian

We have:

$$\begin{aligned} p(x_1, x_2) &= \frac{1}{2\pi\sqrt{|\Sigma|}} \cdot \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma (\mathbf{x} - \mu)\right) \\ &= \frac{1}{2\pi\sigma_1\sigma_2\sqrt{1-\rho^2}} \cdot \exp\left(-\frac{1}{2}\begin{pmatrix} x_1 - \mu_1 & x_2 - \mu_2 \end{pmatrix} \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix} \begin{pmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \end{pmatrix}\right). \end{aligned}$$

4.7 Conditioning a bivariate Gaussian

For question (a), we begin with the form from Exercise 4.6.:

$$p(x_1, x_2) = \frac{\exp\left(-\frac{1}{2(1-\rho^2)}\left(\frac{(x_1-\mu_1)^2}{\sigma_1^2} + \frac{(x_2-\mu_2)^2}{\sigma_2^2} - 2\rho\frac{(x_1-\mu_1)(x_2-\mu_2)}{\sigma_1\sigma_2}\right)\right)}{2\pi\sigma_1\sigma_2\sqrt{1-\rho^2}}.$$

By the Bayes rule:

$$p(x_2|x_1) = \frac{p(x_1, x_2)}{p(x_1)}.$$

So we only need to compute:

$$p(x_1) = \int p(x_1, x_2) dx_2.$$

This can be done by *completing the square* inside $p(x_1, x_2)$ w.r.t. x_2 :

$$\begin{aligned} 2\pi\sigma_1\sigma_2\sqrt{1-\rho^2} \cdot p(x_1, x_2) &= \exp\left(-\frac{1}{2(1-\rho^2)}\frac{(x_1-\mu_1)^2}{\sigma_1^2}\right) \\ &\quad \cdot \exp\left(-\frac{1}{2(1-\rho^2)}\left(\frac{(x_2-\mu_2)^2}{\sigma_2^2} - 2\rho\frac{(x_1-\mu_1)(x_2-\mu_2)}{\sigma_1\sigma_2} + \frac{\rho^2(x_1-\mu_1)^2}{\sigma_1^2}\right)\right) \\ &\quad \cdot \exp\left(\frac{1}{2(1-\rho^2)}\frac{\rho^2(x_1-\mu_1)^2}{\sigma_1^2}\right) \\ &= \exp\left(-\frac{(x_1-\mu_1)^2}{2\sigma_1^2}\right) \\ &\quad \cdot \exp\left(-\frac{1}{2\sigma_2^2(1-\rho^2)}\left((x_2-\mu_2) - \frac{\sigma_2\rho(x_1-\mu_1)}{\sigma_1}\right)^2\right). \end{aligned}$$

Now we are ready to perform the integrating:

$$\begin{aligned} \int p(x_1, x_2) dx_2 &= \frac{\exp\left(-\frac{(x_1-\mu_1)^2}{2\sigma_1^2}\right)}{2\pi\sigma_1\sigma_2\sqrt{1-\rho^2}} \cdot \int \exp\left(-\frac{1}{2\sigma_2^2(1-\rho^2)}(x_2-\mu)^2\right) dx_2 \\ &= \frac{\exp\left(-\frac{(x_1-\mu_1)^2}{2\sigma_1^2}\right)}{2\pi\sigma_1\sigma_2\sqrt{1-\rho^2}} \cdot \sqrt{2\pi\sigma_2^2(1-\rho^2)}, \end{aligned}$$

where

$$\sigma^2 = \sigma_2^2(1 - \rho^2),$$

$$\mu = \mu_2 + \frac{\sigma_2 \rho (x_1 - \mu_1)}{\sigma_1}.$$

Finally,

$$p(x_2|x_1) = \frac{\exp\left(\frac{(x_1 - \mu_1)^2}{2\sigma_1^2}\right) \cdot \exp\left(-\frac{1}{2(1-\rho^2)}\left(\frac{(x_1 - \mu_1)^2}{\sigma_1^2} + \frac{(x_2 - \mu_2)^2}{\sigma_2^2} - 2\rho\frac{(x_1 - \mu_1)(x_2 - \mu_2)}{\sigma_1\sigma_2}\right)\right)}{\sqrt{2\pi\sigma_2^2(1 - \rho^2)}}.$$

For question (b), we can further simplify the numerator of $p(x_2|x_1)$ into:

$$p(x_2|x_1) = \frac{\exp\left(-\frac{1}{2(1-\rho^2)}(\rho(x_1 - \mu_1) - (x_2 - \mu_2))^2\right)}{\sqrt{2\pi(1 - \rho^2)}}.$$

4.8 Whitening vs standardizing

Standardizing is a gold standard step for complex data, e.g., the batch normalization layer. Whitening, although fancy, is harder to carry out since it involves solving an eigen problem.

4.9 Sensor fusion with known variances in 1d

Denote the two observed datasets by $Y^{(1)}$ and $Y^{(2)}$, with size N_1, N_2 , the likelihood is:

$$p(Y^{(1)}, Y^{(2)}|\mu) = \prod_{n_1=1}^{N_1} p(Y_{n_1}^{(1)}|\mu) \cdot \prod_{n_2=1}^{N_2} p(Y_{n_2}^{(2)}|\mu)$$

$$\propto \exp\{A \cdot \mu^2 + B \cdot \mu\},$$

where we have dropped terms independent from μ and used:

$$A = -\frac{N_1}{2v_1} - \frac{N_2}{2v_2},$$

$$B = \frac{1}{v_1} \sum_{n_1=1}^{N_1} Y_{n_1}^{(1)} + \frac{1}{v_2} \sum_{n_2=1}^{N_2} Y_{n_2}^{(2)}.$$

Differentiate the likelihood w.r.t. μ and set it to zero, we have:

$$\mu_{\text{MLE}} = -\frac{B}{2A}$$

The conjugate prior of this model must have a form proportional to $\exp\{A \cdot \mu^2 + B \cdot \mu\}$, so it is a normal distribution:

$$p(\mu|a, b) \propto \exp\{a \cdot \mu^2 + b \cdot \mu\}.$$

The posterior distribution is:

$$p(\mu|Y) \propto \exp\{(A + a) \cdot \mu^2 + (B + b) \cdot \mu\}.$$

Hence we have the MAP estimation:

$$\mu_{\text{MAP}} = -\frac{B + b}{2(A + a)}.$$

It is noticable that the MAP converges to ML estimation when observation times grow:

$$\mu_{\text{MAP}} \rightarrow \mu_{\text{ML}}.$$

The posterior distribution is another normal distribution, with:

$$\sigma_{\text{MAP}}^2 = -\frac{1}{2(A + a)}.$$

For non-informative prior, we have $a = b = 0$ so $p(\mu|a, b)$ is uniform in the domain, then the MAP estimation is the same as MLE.

4.10 Derivation of information form formulae for marginalizing and conditioning

Plugging (4.93) and (4.95) into proven lines of (4.69) yields the information form formulae.

4.11 Derivation of the NIW posterior

We begin with the likelihood for a MVN:

$$p(\mathbf{X}|\mu, \Sigma) = (2\pi)^{-\frac{ND}{2}} |\Sigma|^{-\frac{N}{2}} \exp\left\{-\frac{1}{2} \sum_{n=1}^N (\mathbf{x}_i - \mu)^T \Sigma^{-1} (\mathbf{x}_i - \mu)\right\}.$$

By (4.195), which can be proven by:

$$\begin{aligned}
\sum_{n=1}^N (\mathbf{x}_i - \mu)^T \Sigma^{-1} (\mathbf{x}_i - \mu) &= \sum_{n=1}^N (\bar{\mathbf{x}} - \mu + (\mathbf{x}_i - \bar{\mathbf{x}}))^T \Sigma^{-1} (\bar{\mathbf{x}} - \mu + (\mathbf{x}_i - \bar{\mathbf{x}})) \\
&= N(\bar{\mathbf{x}} - \mu)^T \Sigma^{-1} (\bar{\mathbf{x}} - \mu) + \sum_{n=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})^T \Sigma^{-1} (\mathbf{x}_i - \bar{\mathbf{x}}) \\
&= N(\bar{\mathbf{x}} - \mu)^T \Sigma^{-1} (\bar{\mathbf{x}} - \mu) + \text{tr} \left\{ \Sigma^{-1} \sum_{n=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \right\} \\
&= N(\bar{\mathbf{x}} - \mu)^T \Sigma^{-1} (\bar{\mathbf{x}} - \mu) + \text{tr} \{ \Sigma^{-1} \mathbf{S}_{\bar{\mathbf{x}}} \},
\end{aligned}$$

where we have used the fact that $\text{tr}(\mathbf{Y}^T \mathbf{Z}) = \text{tr}(\mathbf{Z} \mathbf{Y}^T)$, with \mathbf{Y} the shifted design matrix and $\mathbf{Z} = \Sigma^{-1} \mathbf{Y}$.

The conjugate prior for MVN's parameters (μ, Σ) is Normal-inverse-Wishart(NIW) distribution defined by:

$$\begin{aligned}
\text{NIW}(\mu, \Sigma | \mathbf{m}_0, k_0, v_0, \mathbf{S}_0) &= \mathcal{N}(\mu | \mathbf{m}_0, \frac{1}{k_0} \Sigma) \cdot \text{IW}(\Sigma | \mathbf{S}_0, v_0) \\
&= \frac{1}{Z} |\Sigma|^{-\frac{v_0 + D + 2}{2}} \cdot \exp \left\{ -\frac{k_0}{2} (\mu - \mathbf{m}_0)^T \Sigma^{-1} (\mu - \mathbf{m}_0) - \frac{1}{2} \text{tr} \{ \Sigma^{-1} \mathbf{S}_0 \} \right\}.
\end{aligned}$$

Hence the posterior reads (where we have omitted the condition on hyperparameters):

$$p(\mu, \Sigma | \mathbf{X}) \propto |\Sigma|^{-\frac{v_{\mathbf{X}} + D + 2}{2}} \exp \left\{ -\frac{k_{\mathbf{X}}}{2} (\mu - \mathbf{m}_{\mathbf{X}})^T \Sigma^{-1} (\mu - \mathbf{m}_{\mathbf{X}}) - \frac{1}{2} \text{tr} \{ \Sigma^{-1} \mathbf{S}_{\mathbf{X}} \} \right\},$$

where $v_{\mathbf{X}}$, $k_{\mathbf{X}}$, $\mathbf{m}_{\mathbf{X}}$ and $\mathbf{S}_{\mathbf{X}}$ are variables whose values are to be decided. Only terms that dependent on μ and Σ can explicitly enter the terms on the r.h.s.

Firstly, by comparing the exponential for $|\Sigma|$, we have:

$$v_{\mathbf{X}} = v_0 + N.$$

Secondly, compare the coefficient for the term $\mu^T \Sigma^{-1} \mu$ inside the exponential and we have:

$$k_{\mathbf{X}} = k_0 + N.$$

Thirdly, check the coefficient for μ^T so we have:

$$N \Sigma^{-1} \bar{\mathbf{x}} + k_0 \Sigma^{-1} \mathbf{m}_0 = k_{\mathbf{X}} \Sigma^{-1} \mathbf{m}_{\mathbf{X}},$$

therefore:

$$\mathbf{m}_{\mathbf{X}} = \frac{N\bar{\mathbf{x}} + k_0\mathbf{m}_0}{k_{\mathbf{X}}}.$$

Finally, recall that for an arbitrary column vector A :

$$A^T \Sigma^{-1} A = \text{tr}(A^T \Sigma^{-1} A) = \text{tr}(\Sigma^{-1} A A^T).$$

The terms that solely dependent on Σ^{-1} should equal to each other, so:

$$\text{tr}(\Sigma^{-1}(k_0\mathbf{m}_0\mathbf{m}_0^T + \mathbf{S}_0)) + \text{tr}(\Sigma^{-1}(N\bar{\mathbf{x}}\bar{\mathbf{x}}^T \mathbf{S}_{\bar{\mathbf{x}}})) = \text{tr}(\Sigma^{-1}(k_{\mathbf{X}}\mathbf{m}_{\mathbf{X}}\mathbf{m}_{\mathbf{X}}^T + \mathbf{S}_{\mathbf{X}})).$$

Having arrived in:

$$N\bar{\mathbf{x}}\bar{\mathbf{x}}^T + \mathbf{S}_{\bar{\mathbf{x}}} + k_0\mathbf{m}_0\mathbf{m}_0^T + \mathbf{S}_0 = k_{\mathbf{X}}\mathbf{m}_{\mathbf{X}}\mathbf{m}_{\mathbf{X}}^T + \mathbf{S}_{\mathbf{X}},$$

we obtain:

$$\mathbf{S}_{\mathbf{X}} = N\bar{\mathbf{x}}\bar{\mathbf{x}}^T + \mathbf{S}_{\bar{\mathbf{x}}} + k_0\mathbf{m}_0\mathbf{m}_0^T + \mathbf{S}_0 - k_{\mathbf{X}}\mathbf{m}_{\mathbf{X}}\mathbf{m}_{\mathbf{X}}^T.$$

Recall the definition for mean we ends in (4.214) since:

$$\mathbf{S} = \sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^T = \mathbf{S}_{\bar{\mathbf{x}}} + N\bar{\mathbf{x}}\bar{\mathbf{x}}^T.$$

This finishes proving that the posterior distribution for MVN takes the form:

NIW($\mathbf{m}_{\mathbf{X}}, k_{\mathbf{X}}, v_{\mathbf{X}}, \mathbf{S}_{\mathbf{X}}$).

4.12 BIC for Gaussians

For question (a), recall that the maximum likelihood estimation for a MVN model is:

$$\mu_{\text{MLE}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n,$$

$$\Sigma_{\text{MLE}} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \mu_{\text{MLE}})(\mathbf{x}_n - \mu_{\text{MLE}})^T.$$

So the likelihood reads:

$$\begin{aligned} p(\mathcal{D} | \mu_{\text{MLE}}, \Sigma_{\text{MLE}}) &= \prod_{n=1}^N p(\mathbf{x}_n | \mu_{\text{MLE}}, \Sigma_{\text{MLE}}) \\ &= \prod_{n=1}^N (2\pi)^{-\frac{D}{2}} \cdot |\Sigma_{\text{MLE}}|^{-\frac{1}{2}} \cdot \exp \left(-\frac{1}{2} (\mathbf{x}_n - \mu_{\text{MLE}})^T \Sigma_{\text{MLE}}^{-1} (\mathbf{x}_n - \mu_{\text{MLE}}) \right) \\ &= (2\pi)^{-\frac{ND}{2}} \cdot |\Sigma_{\text{MLE}}|^{-\frac{N}{2}} \cdot \exp \left(-\frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n - \mu_{\text{MLE}})^T \Sigma_{\text{MLE}}^{-1} (\mathbf{x}_n - \mu_{\text{MLE}}) \right). \end{aligned}$$

Denote:

$$\mathbf{Y} = \begin{pmatrix} \mathbf{x}_1 - \mu_{\text{MLE}} & \cdots & \mathbf{x}_N - \mu_{\text{MLE}} \end{pmatrix},$$

then $\Sigma_{\text{MLE}} = \frac{1}{N} \mathbf{Y} \mathbf{Y}^T$, while the term in the exponential of the likelihood is:

$$-\frac{1}{2} \cdot \text{tr}(\mathbf{Y}^T \Sigma_{\text{MLE}}^{-1} \mathbf{Y}) = -\frac{1}{2} \cdot \text{tr}(\Sigma_{\text{MLE}}^{-1} \mathbf{Y} \mathbf{Y}^T) = -\frac{ND}{2}.$$

Thus the BIC is:

$$-\frac{ND}{2} \cdot \log(2\pi e) - \frac{N}{2} \cdot \log |\Sigma_{\text{MLE}}| - \frac{D + \frac{D(D+1)}{2}}{2} \cdot \log N.$$

For question (b), the fitting of a diagonal MVN model is tantamount to fitting D independent 1d Gaussian models simultaneously, thus the d -th diagonal component of Σ_{MLE}^d is:

$$\frac{1}{N} \sum_{n=1}^N \mathbf{x}_{n,d}^2,$$

where we have assumed $\mathbf{x} = \mathbf{0}$ w.l.o.g. Thus the term inside the exponential of the likelihood remains $-\frac{ND}{2}$. So the BIC in this case is:

$$-\frac{ND}{2} \cdot \log(2\pi e) - \frac{N}{2} \cdot \log |\Sigma_{\text{MLE}}^d| - D \cdot \log N.$$

We observe that if all D components are mutually independent, i.e., Σ_{MLE} is diagonal then the BIC for diagonal MVN model is strictly larger than that for general MVN, hence the diagonal version is always preferred. In cases there exists dependence among components, the BIC for general MVN is still not necessarily larger than that of diagonal MVN. This is a reflection of the trade-off between complexity and generalization.

The Bayesian information criterion is an approximation of a model's evidence, $p(\mathcal{D})$. Let us start from:

$$p(\mathcal{D}) = \int p(\mathcal{D}|\theta) \cdot p(\theta) d\theta,$$

where θ is the collection of all parameters within the current model. The trick here is to expand $\log p(\mathbf{x}|\theta)$ as a function of θ and taking approximation to the second order at θ_{MAP} so the first order gradient vanishes:

$$\log p(\mathbf{x}|\theta) \approx \log p(\mathbf{x}|\theta_0) - \frac{1}{2}(\theta - \theta_0)^T \mathbf{H}(\theta - \theta_0),$$

where \mathbf{H} is the Hessian matrix at $\log p(\mathbf{x}|\theta_0)$. Thus we have:

$$p(\mathbf{x}|\theta) \approx p(\mathbf{x}|\theta_0) \cdot \exp\left(-\frac{1}{2}(\theta - \theta_0)^T \mathbf{H}(\theta - \theta_0)\right).$$

We are now ready to perform the integral, with:

$$p(\mathcal{D}|\theta) = p(\mathbf{x}|\theta_0)^N \cdot \exp\left(-\frac{N}{2}(\theta - \theta_0)^T \mathbf{H}(\theta - \theta_0)\right),$$

conducting the integral on the neighbour of $\theta_0 = \theta_{\text{MAP}}$:

$$\begin{aligned} \int p(\mathcal{D}|\theta) \cdot p(\theta) d\theta &\approx p(\mathcal{D}|\theta_{\text{MAP}}) \cdot p(\theta_{\text{MAP}}) \cdot \int \exp\left(-\frac{N}{2}(\theta - \theta_{\text{MAP}})^T \mathbf{H}(\theta - \theta_{\text{MAP}})\right) d\theta \\ &= p(\mathcal{D}|\theta_{\text{MAP}}) \cdot p(\theta_{\text{MAP}}) \cdot (2\pi)^{\frac{d}{2}} |N^{-1} \mathbf{H}^{-1}|^{\frac{1}{2}} \\ &= p(\mathcal{D}|\theta_{\text{MAP}}) \cdot p(\theta_{\text{MAP}}) \cdot (2\pi)^{\frac{d}{2}} \cdot N^{-\frac{d}{2}} \cdot |\mathbf{H}^{-1}|^{\frac{1}{2}}, \end{aligned}$$

where d is the number of components in θ . Taking the logarithm of both side of the evidence yields the BIC. One can see how many compromises and assumptions have been applied in deriving an analytic form of the evidence, which is arguably the most complex variable for Bayesian analysis.

See also *PRML*, Section 4.4.

4.13 Gaussian posterior credible interval

Assume the prior distribution for an 1d normal distribution:

$$p(\mu) = \mathcal{N}(\mu|\mu_0, \sigma_0^2 = 9).$$

And the likelihood is:

$$p(x) = \mathcal{N}(x|\mu, \sigma^2 = 4).$$

Having observed n variables, we want that the probability measure of μ 's posterior distribution is no less than 0.95 within an interval no longer than 1. The posterior for μ is:

$$\begin{aligned} p(\mu|D) &\propto p(\mu) \cdot p(D|\mu) = \mathcal{N}(\mu|\mu_0, \sigma_0^2) \cdot \prod_{i=1}^n \mathcal{N}(x_i|\mu, \sigma^2) \\ &\propto \exp\left\{-\frac{1}{2\sigma_0^2}(\mu - \mu_0)^2\right\} \cdot \prod_{i=1}^n \exp\left\{-\frac{1}{2\sigma^2}(x_i - \mu)^2\right\} \\ &= \exp\left\{\left(-\frac{1}{2\sigma_0^2} - \frac{n}{2\sigma^2}\right)\mu^2 + \dots\right\}, \end{aligned}$$

where we have dropped the terms irrelevant with μ . The posterior variance of μ is determined by the coefficient of μ^2 in the exponential of the posterior distribution:

$$\sigma_{\text{post}}^2 = \frac{\sigma_0^2 \sigma^2}{\sigma^2 + n \sigma_0^2}.$$

Since 0.95 of the probability mass for a normal distribution lies within -1.96σ and 1.96σ , we have:

$$n \geq 611.$$

4.14 MAP estimation for 1d Gaussians

Assume that the variance for this distribution σ^2 is known, and the mean μ is subject to a normal distribution with mean m and variance s^2 . Similiar to the question before, the posterior takes the form:

$$p(\mu|X) \propto p(\mu) \cdot p(X|\mu).$$

So the posterior is another normal distribution, by comparing the coefficient for μ^2 in the exponential:

$$-\frac{1}{2s^2} - \frac{N}{2\sigma^2},$$

and that for μ :

$$\frac{m}{s^2} + \frac{\sum_{n=1}^N x_n}{\sigma^2},$$

we have the posterior mean and variance by completing the square:

$$\sigma_{\text{post}}^2 = \frac{s^2 \sigma^2}{\sigma^2 + N s^2},$$

$$\mu_{\text{post}} = \left(\frac{m}{s^2} + \frac{\sum_{n=1}^N x_n}{\sigma^2} \right) \cdot \sigma_{\text{post}}^2.$$

This finishes question (a).

For question (b), we already knew that the MLE is:

$$\mu_{\text{MLE}} = \frac{\sum_{n=1}^N x_n}{N}.$$

As N increases, we have:

$$\begin{aligned} \lim_{N \rightarrow \infty} \mu_{\text{post}} &= \lim_{N \rightarrow \infty} \frac{\frac{\sigma^2}{s^2} \cdot m + \sum_{n=1}^N x_n}{\frac{\sigma^2}{s^2} + N} \\ &= \frac{\sum_{n=1}^N x_n}{N}. \end{aligned}$$

For question (c), when $s^2 \rightarrow \infty$, μ_{post} also converges to μ_{MLE} since $\frac{\sigma^2}{s^2} \rightarrow 0$.

For question (d), when $s^2 \rightarrow 0$, then $\frac{\sigma^2}{s^2} \rightarrow \infty$ and μ_{post} converges to m .

Both (c) and (d) are very intuitive. $s^2 \rightarrow \infty$ means a non-informative prior has been introduced so MAP is the same as MLE. $s^2 \rightarrow 0$ means that the knowledge that μ is close to m is very strong so that finite observations cannot modify this belief.

4.15 Sequential(recursive) updating of covariance matrix

For question (a), note that:

$$n\mathbf{C}_{n+1} - (n-1)\mathbf{C}_n = \sum_{i=1}^{n+1} (\mathbf{x}_i - \mathbf{m}_{n+1})(\mathbf{x}_i - \mathbf{m}_{n+1})^T - \sum_{i=1}^n (\mathbf{x}_i - \mathbf{m}_n)(\mathbf{x}_i - \mathbf{m}_n)^T.$$

Making use of:

$$\mathbf{m}_{n+1} = \frac{n\mathbf{m}_n + \mathbf{x}_{n+1}}{n+1},$$

we have:

$$\begin{aligned} n\mathbf{C}_{n+1} - (n-1)\mathbf{C}_n &= \mathbf{x}_{n+1}\mathbf{x}_{n+1}^T - (n+1)\mathbf{m}_{n+1}\mathbf{m}_{n+1}^T + n\mathbf{m}_n\mathbf{m}_n^T \\ &= \frac{n}{n+1}(\mathbf{x}_{n+1} - \mathbf{m}_n)(\mathbf{x}_{n+1} - \mathbf{m}_n)^T. \end{aligned}$$

For question (b), the complexity is $\mathcal{O}(d^2)$.

For question (c), plugging (4.281) directly into (4.279) yields (4.280).

For question (d), the complexity remains $\mathcal{O}(d^2)$.

4.16 Likelihood ratio for Gaussians

Consider a classifier for two classes, the generative distribution for them are two normal distributions $p(x|y = C_i) = \mathcal{N}(x|\mu_i, \Sigma_i)$, by the Bayes rule:

$$\log \frac{p(y=1|x)}{p(y=0|x)} = \log \frac{p(x|y=1)}{p(x|y=0)} + \log \frac{p(y=1)}{p(y=0)}.$$

The first term on r.h.s. is the ratio of likelihood probability.

When we have arbitrary covariance matrices:

$$\frac{p(x|y=1)}{p(x|y=0)} = \sqrt{\frac{|\Sigma_0|}{|\Sigma_1|}} \cdot \exp \left\{ -\frac{1}{2}(x - \mu_1)^T \Sigma_1^{-1} (x - \mu_1) + \frac{1}{2}(x - \mu_0)^T \Sigma_0^{-1} (x - \mu_0) \right\}.$$

As Σ_0, Σ_1 are arbitrary matrices, this formulation cannot be reduced further:

$$\log \frac{p(x|y=1)}{p(x|y=0)} = \frac{1}{2} \log \frac{|\Sigma_0|}{|\Sigma_1|} - \frac{1}{2} (x - \mu_1)^T \Sigma_1^{-1} (x - \mu_1) + \frac{1}{2} (x - \mu_0)^T \Sigma_0^{-1} (x - \mu_0).$$

Note that the decision boundary ($\log \frac{p(x|y=1)}{p(x|y=0)} = 0$) is a quadratic surface in D -dimension space.

When both covariance matrixes are given by Σ :

$$\frac{p(x|y=1)}{p(x|y=0)} = \exp \left\{ -\frac{1}{2} (x - \mu_1)^T \Sigma^{-1} (x - \mu_1) + \frac{1}{2} (x - \mu_0)^T \Sigma^{-1} (x - \mu_0) \right\},$$

so:

$$\begin{aligned} \log \frac{p(x|y=1)}{p(x|y=0)} &= -\frac{1}{2} (x - \mu_1)^T \Sigma^{-1} (x - \mu_1) + \frac{1}{2} (x - \mu_0)^T \Sigma^{-1} (x - \mu_0) \\ &= \frac{1}{2} \text{tr} \left(\Sigma^{-1} [(x - \mu_1)(x - \mu_1)^T - (x - \mu_0)(x - \mu_0)^T] \right). \end{aligned}$$

When Σ is a diagonal matrix, we have:

$$\begin{aligned} \log \frac{p(x|y=1)}{p(x|y=0)} &= -\frac{1}{2} (x - \mu_1)^T \Sigma^{-1} (x - \mu_1) + \frac{1}{2} (x - \mu_0)^T \Sigma^{-1} (x - \mu_0) \\ &= \frac{1}{2} \text{tr} \left(\Sigma^{-1} [(x - \mu_1)(x - \mu_1)^T - (x - \mu_0)(x - \mu_0)^T] \right) \\ &= \frac{1}{2} \text{tr} (\Lambda^{-1} \Phi) \\ &= \frac{1}{2} \sum_{i=1}^d \lambda_i^{-1} \Phi_{i,i}. \end{aligned}$$

where:

$$\Phi = (x - \mu_1)(x - \mu_1)^T - (x - \mu_0)(x - \mu_0)^T.$$

Finally, if $\Sigma = \sigma^2 I$ then:

$$\log \frac{p(x|y=1)}{p(x|y=0)} = \frac{1}{2\sigma^2} \text{tr}(\Phi).$$

Note that for the last three cases, a decision boundary is a linear plane in the space, since the quadratic term on x has been canceled in Φ .

4.17 LDA/QDA on height/weight data

...

4.18 Naive Bayes with mixed features

For question (a):

$$\begin{aligned} p(y = 1|x_1 = 0, x_2 = 0) &= \frac{p(y = 1) \cdot p(x_1|y = 1) \cdot p(x_2 = 0|y = 1)}{p(x_1 = 0, x_2 = 0)} \\ &= \frac{0.5 \cdot 0.5 \cdot \frac{\exp(-0.5)}{\sqrt{2\pi}}}{p(x_1 = 0, x_2 = 0)}. \end{aligned}$$

Similarly,

$$\begin{aligned} p(y = 2|x_1 = 0, x_2 = 0) &= \frac{0.25 * 0.5 * \frac{1}{\sqrt{2\pi}}}{p(x_1 = 0, x_2 = 0)}, \\ p(y = 3|x_1 = 0, x_2 = 0) &= \frac{0.25 * 0.5 * \frac{\exp(-0.5)}{\sqrt{2\pi}}}{p(x_1 = 0, x_2 = 0)}. \end{aligned}$$

A normalization yields the final result by eliminating $p(x_1 = 0, x_2 = 0)$.

For question (b), we have:

$$p(y = 1|x_1 = 0) = 0.5,$$

$$p(y = 2|x_1 = 0) = 0.25,$$

$$p(y = 3|x_1 = 0) = 0.25,$$

since x_1 yields no more information for the classification label.

For question (c), we have:

$$\begin{aligned} p(y = 1|x_2 = 0) &\propto 0.5 * \frac{\exp(-0.5)}{\sqrt{2\pi}}, \\ p(y = 2|x_2 = 0) &\propto 0.25 * \frac{1}{\sqrt{2\pi}}, \\ p(y = 3|x_2 = 0) &\propto 0.25 * \frac{\exp(-0.5)}{\sqrt{2\pi}}. \end{aligned}$$

One can observe that unlike $p(y|x_1 = 0)$, $p(y|x_2 = 0)$ is different from the prior on labels.

4.19 Decision boundary for LDA with semi tied covariances

We begin from the Bayes rule:

$$p(y = 1|\mathbf{x}, \theta) \propto p(y = 1|\theta) \cdot p(\mathbf{x}|y = 1, \theta),$$

where we have omitted the terms independent of y . With a uniform prior on two classes:

$$\begin{aligned} p(y = 1|\mathbf{x}, \theta) &\propto \mathcal{N}(\mathbf{x}|\mu_1, k\Sigma_0), \\ p(y = 0|\mathbf{x}, \theta) &\propto \mathcal{N}(\mathbf{x}|\mu_0, \Sigma_0). \end{aligned}$$

The decision boundary, in which we are interested, is a curve depicted by $f(\mathbf{x}) = 0$, where:

$$f(x) = \log \frac{p(y = 1|\mathbf{x}, \theta)}{p(y = 0|\mathbf{x}, \theta)}.$$

Therefore the decision boundary is:

$$\text{tr}(\Sigma_0^{-1}[\Phi(\mathbf{x})]) = d \cdot \ln k,$$

where:

$$\Phi(\mathbf{x}) = \frac{(\mathbf{x} - \mu_1)(\mathbf{x} - \mu_1)^T}{k} - (\mathbf{x} - \mu_0)(\mathbf{x} - \mu_0)^T.$$

The decision boundary is a quadratic curve unless $k = 1$, which is geometrically very intuitive.

Let us consider a 2d case, focusing on the transformed coordinates where Σ_0 is the identity matrix. With out loss of generality, let $\mu_0 = (0, 0)$, $\mu_1 = (z, 0)$, denote the distance between a point p in this place from μ_0 and μ_1 by $a(p)$ and $b(p)$. The decision boundary is exactly:

$$\left\{ p : \frac{b(p)^2}{k} - a(p)^2 = -\frac{1}{2} \ln k \right\}.$$

Plugging in the Cartesian representation $p(x, y)$, we ends up with a conical curve. The linear transform of the space would not change its conical nature.

4.20 Logistic regression vs LDA/QDA

The underlying assumptions for all four classifiers are as follows:

- GaussI assumes a covariance matrix as an identity matrix;
- GaussX has no prior assumption on the covariance matrix;
- LinLog assumes that different classes share the same covariance matrix;

- QuadLog has no prior assumption on covariance matrix, yet it assumes that all data from one class are subject to a normal distribution;

From the perspective of complexity we have the following order:

QuadLog=GaussX>LinLog>GaussL.

The MLE likelihood should follow the same order, this answers the question (a)-(d).

For question (e), the argument is untrue in general. For example, model M predicts two samples belonging to the first class with probability vectors $(0.49, 0.51)$ and $(0.99, 0.1)$. While M' outputs $(0.51, 0.49)$ and $(0.51, 0.49)$. Now M' is correct on both samples so $R(M) > R(M')$, but:

$$\frac{\log(0.49) + \log(0.99)}{2} > \frac{\log(0.51) + \log(0.51)}{2},$$

so $L(M) > L(M')$, this is sufficient for disproving the argument.

4.21 Gaussian decision boundaries

We have:

$$p(x|\mu_i, \sigma_i^2) = \frac{1}{\sqrt{2\pi\sigma_i^2}} \cdot \exp\left(-\frac{(x - \mu_i)^2}{2\sigma_i^2}\right),$$

so:

$$p(x|\mu_1, \sigma_1^2) = \frac{1}{\sqrt{2\pi}} \cdot \exp\left(-\frac{x^2}{2}\right),$$

$$p(x|\mu_2, \sigma_2^2) = \frac{1}{10^3 \cdot \sqrt{2\pi}} \cdot \exp\left(-\frac{(x-1)^2}{2 \times 10^6}\right).$$

The decision region satisfies:

$$\frac{p(x|\mu_1, \sigma_1^2)}{p(x|\mu_2, \sigma_2^2)} \geq 1,$$

this is tantamount to:

$$\frac{(x-1)^2}{10^6} - x^2 \geq -6 \cdot \ln 10.$$

Denote x_1 and x_2 as the zeros of this quadratic form, then

$$R_1 = [x_1, x_2].$$

When $\sigma^2 = \sigma_1^2$, R_1 is exactly $(-\infty, \frac{1}{2})$.

One can solve for (a) and (b) by plugging in (4.289).

4.22 QDA with 3 classes

Solve (a) and (b) numerically:

```

1 import math
2 import numpy as np
3 mu1=np.array([0,0])
4 mu2=np.array([1,1])
5 mu3=np.array([1,-1])
6 s1=np.array([[0.7,0],[0,0.7]])
7 s2=np.array([[0.8,0.2],[0.2,0.8]])
8 s3=np.array([[0.8,0.2],[0.2,0.8]])
9 def p(x):
10     f1=np.linalg.det(s1)**(-0.5)*math.exp(-0.5*(mu1-x)@np.
        linalg.inv(s1)@(mu1-x))
11     f2=np.linalg.det(s2)**(-0.5)*math.exp(-0.5*(mu2-x)@np.
        linalg.inv(s2)@(mu2-x))
12     f3=np.linalg.det(s3)**(-0.5)*math.exp(-0.5*(mu3-x)@np.
        linalg.inv(s3)@(mu3-x))
13     return f1,f2,f3
14 x1=np.array([-0.5,0.5])
15 x2=np.array([0.5,0.5])
16 print(p(x1))
17 print(p(x2))

```

The outcome is:

```

1 (0.9995321962501862, 0.31309336105606445, 0.030361279346887253)
2 (0.9995321962501862, 1.005427487616282, 0.18990072283298057)

```

So the predicted labels are 1 and 2 respectively.

4.23 Scalar QDA

```

1 import math
2 hm=[67,79,71]
3 hf=[68,67,60]
4 def mu(h):
5     return (h[0]+h[1]+h[2])/3

```

```
6 def sigma2(h):
7     m=mu(h)
8     return ((h[0]-m)**2+(h[1]-m)**2+(h[2]-m)**2)/3
9 # For question (a):
10 mu_m=mu(hm)
11 mu_f=mu(hf)
12 sigma2_m=sigma2(hm)
13 sigma2_f=sigma2(hf)
14 print(mu_m)
15 print(sigma2_m)
16 print(mu_f)
17 print(sigma2_f)
18 # \pi_{m}=\pi_{f}=0.5.
19 # For question (b):
20 temp_m=(2*math.pi*sigma2_m)**(-0.5)*math.exp(-(72-mu_m)**2/2/
21         sigma2_m)
22 temp_f=(2*math.pi*sigma2_f)**(-0.5)*math.exp(-(72-mu_f)**2/2/
23         sigma2_f)
24 print(temp_m/(temp_m+temp_f))
```

For question (c) we can use a naive Bayes model, which is tantamount to adopting diagonal covariance matrices for both classes.

5 Bayesian statistics

This Chapter is a continuation of Chapter 2. Section 5.7.1.5. sheds light on PAC theory, the fundamental theory underlying machine learning, elegant but somewhat impractical. Section 5.7.3.1. introduces some basic elements from reinforcement learning. Both branches require solid and involved theories, especially those from probability.

5.1 Proof that a mixture of conjugate priors is indeed conjugate

The mixed conjugate prior takes the form (5.68):

$$p(\theta) = \sum_k p(z = k) \cdot p(\theta|z = k),$$

where k is the index for mixed components. Each $p(\theta|z = k)$ is a conjugate prior for the model, i.e., $p(\theta|z = k)$ and $p(\theta|z = k, \mathcal{D})$ takes the same form.

For the posterior in this case:

$$\begin{aligned} p(\theta|\mathcal{D}) &= \sum_k p(\theta, z = k|\mathcal{D}) \\ &= \sum_k p(z = k|\mathcal{D}) \cdot p(\theta|z = k, \mathcal{D}). \end{aligned}$$

So the posterior is still a mixture of conjugate priors. This is exactly (5.69), finally:

$$\begin{aligned} p(z = k|\mathcal{D}) &= \frac{p(z = k, \mathcal{D})}{p(\mathcal{D})} \\ &= \frac{p(z = k) \cdot p(\mathcal{D}|p(z = k))}{\sum_{k'} p(z = k') \cdot p(\mathcal{D}|p(z = k'))}, \end{aligned}$$

from Bayes rules. The computation bottleneck in computing $p(z = k|\mathcal{D})$ is:

$$p(\mathcal{D}|z = k) = \int p(\mathcal{D}|\theta) \cdot p(\theta|z = k) d\theta,$$

which is not a easy task as well.

5.2 Optimal threshold on classification probability

For question (a), the posterior loss expectation for choosing action \hat{y} , given x , is:

$$\begin{aligned}\rho(\hat{y}|x) &= \mathbb{E}_{p(y|x)}[L(y, \hat{y})] \\ &= p_0 \cdot L(\hat{y}, 0) + p_1 \cdot L(\hat{y}, 1) \\ &= p_0 \cdot L(\hat{y}, 0) + (1 - p_0) \cdot L(\hat{y}, 1) \\ &= L(\hat{y}, 1) + p_0(L(\hat{y}, 0) - L(\hat{y}, 1)).\end{aligned}$$

Thus:

$$\begin{aligned}\rho(0|x) &= \lambda_{01} - p_0 \cdot \lambda_{01}, \\ \rho(1|x) &= p_0 \cdot \lambda_{10}.\end{aligned}$$

As both $\rho(0|x)$ and $\rho(1|x)$ are linear functions of p_0 , hence the unique optimal threshold is where $\rho(0|x) = \rho(1|x)$, i.e.,

$$\begin{aligned}p_0 &= \frac{\lambda_{01}}{\lambda_{01} + \lambda_{10}}, \\ p_1 &= \frac{\lambda_{10}}{\lambda_{01} + \lambda_{10}}.\end{aligned}$$

For question (b), let $\lambda_{10} = 1$, $\lambda_{01} = 9$ will do.

5.3 Reject option in classifiers

For question (a), the posterior expected loss for choosing an non-reject action \hat{y} given data x is:

$$\begin{aligned}\rho(\hat{y}|x) &= \sum_{y=1}^C p(y|x) \cdot L(\hat{y}, y) \\ &= \mathbf{p}^T \mathbf{l}(\hat{y}),\end{aligned}$$

where \mathbf{p} is the column vector encoding $p(y|x)$ and $\mathbf{l}(\hat{y})$ is a column vector whose elements are λ_s except for the \hat{y} -th one. Thus the expected loss is $(1 - p(\hat{y}|x)) \cdot \lambda_s$ in this case, whose minimum is obtained by let

$$\hat{y} = \arg \max_y (p(y|x)).$$

For the reject option, the loss is uniform λ_s .

Thus one should choose reject or \hat{y} by minimizing:

$$\min(\lambda_r, (1 - p(\hat{y}|x)) \cdot \lambda_s).$$

If

$$\lambda_r \leq (1 - p(\hat{y}|x)) \cdot \lambda_s,$$

then we readily adopt the reject option. This condition is tantamount to what is required to be prove:

$$p(\hat{y}|x) \leq 1 - \frac{\lambda_r}{\lambda_s}.$$

For question (b), the minimum of the expected loss is:

$$\lambda_s \cdot \min \left\{ \frac{\lambda_r}{\lambda_s}, 1 - p(\hat{y}|x) \right\},$$

where \hat{y} is the most probable class. When $\frac{\lambda_r}{\lambda_s}$ is negligible, the reject option would always be chosen. When $\frac{\lambda_r}{\lambda_s}$, the reject option would never be chosen.

5.4 More reject options

For question (a), we need to choose from:

$$(3, 10 \times (1 - 0.8) = 2),$$

hence the optimal decision is class 0.

For question (b), we need to choose from:

$$(3, 10 \times (1 - 0.6) = 4),$$

hence the optimal decision is reject.

For question (c), the threshold is where:

$$\lambda_r = (1 - p(\hat{y}|x)) \cdot \lambda_s.$$

Plugging in figures we have:

$$p(\hat{y}|x) = 0.7.$$

Thus $\theta_1 = 0.7$, $\theta_0 = 0.3$ by symmetry.

5.5 Newsvendor problem

We have:

$$\mathbb{E}(\pi|Q) = P \int_0^Q Df(D)dD - CQ \int_0^Q f(D)dD + (P - C)Q \int_Q^{+\infty} f(D)dD.$$

Take derivative w.r.t. Q :

$$\frac{\partial}{\partial Q} \mathbb{E}(\pi|Q) = PQf(Q) - C \int_0^Q f(D)dD - CQf(Q) + (P - C) \int_Q^{+\infty} f(D)dD - (P - C)Qf(Q).$$

Setting it to zero by making use of $\int_0^Q f(D)dD + \int_Q^{+\infty} f(D)dD = 1$, we arrive in:

$$F(Q^*) = \frac{P - C}{P},$$

where:

$$F(Q) = \int_0^Q f(D)dD$$

is the c.d.f. of D .

5.6 Bayes factors and ROC curves

If the prior on models are identical then both methods yield the same decision boundary.

5.7 Bayes model averaging helps predictive accuracy

Suppose the variable Δ is generated from a mixture of models, so:

$$p(\Delta|\mathcal{D}) = \sum_{m \in \mathcal{M}} p(\Delta|m, \mathcal{D}) \cdot p(m|\mathcal{D}).$$

The Bayes model averaging (BMA) result is just:

$$p^{\text{BMA}}(\Delta) = \sum_{m \in \mathcal{M}} p(\Delta|m, \mathcal{D}) \cdot p(m|\mathcal{D}).$$

While that from an individual model $m \in \mathcal{M}$ is:

$$p^m(\Delta) = p(\Delta|m, \mathcal{D}).$$

The expected loss of the BMA result is:

$$\mathbb{E}_{p(\Delta)}[-\log p^{\text{BMA}}(\Delta)],$$

while that of model m is:

$$\mathbb{E}_{p(\Delta)}[-\log p^m(\Delta)].$$

Now it is easy to see that:

$$\mathbb{E}_{p(\Delta)}[-\log p^m(\Delta)] - \mathbb{E}_{p(\Delta)}[-\log p^{\text{BMA}}(\Delta)] = \mathbb{KL}(p^{\text{BMA}}(\Delta) || p^m(\Delta)),$$

since the distribution on which the expectation is computed is just $p^{\text{BMA}}(\Delta)$. Therefore the non-negativity of the KL divergence yields (5.127).

The conclusion from this exercise is of hardly any practical significance. Since the underlying distribution is usually intractable, even with the mixture Bayes model. Once the form of the latent distribution is revealed, it is obvious that other distributions result in higher loss.

5.8 MLE and model selection for a 2d discrete distribution

For question (a), the joint distribution $p(x, y | \theta_1, \theta_2)$ is given by:

$$\begin{aligned} p(x=0, y=0) &= (1 - \theta_1) \cdot \theta_2, \\ p(x=0, y=1) &= (1 - \theta_1) \cdot (1 - \theta_2), \\ p(x=1, y=0) &= \theta_1 \cdot (1 - \theta_2), \\ p(x=1, y=1) &= \theta_1 \cdot \theta_2. \end{aligned}$$

This can be compactly written as:

$$p(x, y | \theta_1, \theta_2) = \theta_1^x (1 - \theta_1)^{(1-x)} \theta_2^{x \odot y} (1 - \theta_2)^{(1-x \odot y)},$$

where \odot is the Exclusive NOR operator.

For question (b), the MLE for θ_1 is $\frac{4}{7}$, while that for θ_2 is $\frac{4}{7}$. Since we assumed the independency between θ_1 and θ_2 , both MLE can be arrived at by simply counting. The evidence is given by:

$$p(\mathcal{D} | \theta_{\text{MLE}}) = \left(\frac{4}{7}\right)^4 \cdot \left(\frac{3}{7}\right)^3 \cdot \left(\frac{3}{7}\right)^3 \cdot \left(\frac{4}{7}\right)^4.$$

For question (c), the MLE for θ is computed by normalizing the counting vector: $(2, 1, 2, 2)$, so:

$$\theta_{\text{MLE}} = \left(\frac{2}{7}, \frac{2}{7}, \frac{1}{7}, \frac{2}{7}\right).$$

The evidence is:

$$\left(\frac{2}{7}\right)^2 \cdot \left(\frac{2}{7}\right)^2 \cdot \left(\frac{1}{7}\right) \cdot \left(\frac{2}{7}\right)^2 \cdot$$

For question (d):

```

1  import math
2  x=[1,1,0,1,1,0,0]
3  y=[1,0,0,0,1,0,1]
4  l2=0
5  l4=0
6  e=1/10**5
7  for shadow in range(7):
8      temp1=0
9      temp2=0
10     temp00=0
11     temp10=0
12     temp01=0
13     temp11=0
14     for i in range(len(x)):
15         if i==shadow:
16             continue
17         if x[i]==1:
18             temp1=temp1+1
19         if x[i]==y[i]:
20             temp2=temp2+1
21         if x[i]==0 and y[i]==0:
22             temp00=temp00+1
23         if x[i]==0 and y[i]==1:
24             temp01=temp01+1
25         if x[i]==1 and y[i]==0:
26             temp10=temp10+1
27         if x[i]==1 and y[i]==1:
28             temp11=temp11+1
29     theta_1=temp1/(len(x)-1)
30     theta_2=temp2/(len(x)-1)
31     s=temp00+temp01+temp10+temp11
32     theta_00=temp00/s
33     theta_01=temp01/s
34     theta_10=temp10/s

```

```

35     theta_11=temp11/s
36     p2=theta_1**(x[shadow])*(1-theta_1)**(1-x[shadow])*theta_2
        *(1-x[shadow]^y[shadow])*(1-theta_2)**(x[shadow]^y[
        shadow])
37     p4=theta_00**(x[shadow]==0 and y[shadow]==0)*theta_01**(x[
        shadow]==0 and y[shadow]==1)*theta_10**(x[shadow]==1
        and y[shadow]==0)*theta_11**(x[shadow]==1 and y[shadow
        ==1])
38     l2=l2+math.log(p2+e)
39     l4=l4+math.log(p4+e)
40     print(l2)
41     print(l4)

```

The result is:

```

1     -12.136441189337646
2     -28.04302596169576

```

Hence the CV will pick M_2 . The reason behind is that M_4 assumes zero probability for $(0, 1)$ during the cross-validation, which significantly declines the confidence.

For question (e), the BICs for M_2 and M_4 are respectively:

$$\text{BIC}(M_2, \mathcal{D}) = -11.51,$$

$$\text{BIC}(M_4, \mathcal{D}) = -12.38.$$

Hence the BIC prefers M_2 as well.

5.9 Posterior median is optimal estimate under L1 loss

The posterior loss expectation is (where we have omitted \mathcal{D} w.l.o.g.):

$$\begin{aligned}
 \rho(a) &= \int |y - a|p(y)dy = \int_{-\infty}^a (a - y)p(y)dy + \int_a^{+\infty} (y - a)p(y)dy \\
 &= a \left\{ \int_{-\infty}^a p(y)dy - \int_a^{+\infty} p(y)dy \right\} - \int_{-\infty}^a yp(y)dy + \int_a^{+\infty} yp(y)dy.
 \end{aligned}$$

Differentiating the loss w.r.t. a ends up with:

$$\frac{\partial}{\partial a} \rho(a) = \left\{ \int_{-\infty}^a p(y)dy - \int_a^{+\infty} p(y)dy \right\} + a \cdot 2p(a) - a \cdot 2p(a).$$

Set it to zero so:

$$\int_{-\infty}^a p(y)dy = \int_a^{+\infty} p(y),$$

whose summation is unity, hence a satisfies:

$$\int_{-\infty}^a p(y)dy = \int_a^{+\infty} p(y) = \frac{1}{2},$$

which identifies the median.

5.10 Decision rule for trading off FPs and FNs

This is tantamount to exercise 5.2. Replacing λ_{01} by $c \cdot \lambda_{10}$, the rest follows the corresponding results.

6 Frequentist statistics

For one who has delved into Bayesian statistics before acknowledging the frequentist statistics, the latter is somewhat awkward and intimidating. However, in most scenarios regarding science and engineering, orthodox statistics remains the optimal candidate. Since when the size of the dataset grows, the privilege of using Bayes is easily overwhelmed by the asymptotic performance of classical statistics. Moreover, although Bayesians claim to have eliminated the overfitting problem with prior distributions, the subjectiveness in selecting hyperparameters also increases. This problem cannot be solved with empirical Bayesian or more complex graphical structures, since none of them cut down the degree of freedom for parameters.

For the reasons above, the frequentist statistics should receive equal attention as, if not more than, the Bayesian version.

6.1 Pessimism of LOOCV

Under the setting in this problem, the best classification method can achieve is the accuracy of 50%, since the features are independent of the labels. This is, however, the idealistic situation where we assume that

$$N \rightarrow \infty,$$

where $N = N_1 = N_2$. When N is finite, the classification accuracy usually fluctuates around 50%. For the label to be *randomly* assigned, we should have that for an arbitrary (probabilistic) algorithm/machine that terminates within a time complexity a polynomial of N_i , its classification accuracy should not deviate from 50% larger than $\frac{1}{p(N)}$, in which p is an arbitrary polynomial function. The machine in this case is referred to as a probabilistic polynomial time (PPT) machine and is of remarkable interest to cryptologists. To generate cheap pseudo-random numbers efficiently is one of the most important tasks for cryptography.

For LOOCV, if a sample for class one is omitted, then the classification probability is:

$$\left(\frac{N-1}{2N-1}, \frac{N}{2N-1}\right).$$

Hence the probability that the omitted sample is correctly labelled is:

$$\frac{N-1}{2N-1} \leq \frac{1}{2},$$

justifies its pessimism in this setting.

The dependency between features and labels paves the way for decoupling features or security/forensics in machine learning. However, it is impossible to justify that some dataset *cannot be learned* by any machines, whose converse side is extremely simple.

6.2 James Stein estimator for Gaussian means

The prior for θ_i is:

$$\mathcal{N}(\theta_i | m_0, \tau_0^2),$$

and the likelihood is given by:

$$\mathcal{N}(y_i | \theta_i, \sigma^2).$$

For question (a), we begin by integrating out θ_i and establishing the dependency of $\mathcal{D} = \{y_i\}_{i=1}^6$ on m_0 and τ_0^2 :

$$\begin{aligned} p(y|m_0, \tau_0^2) &= \int p(y, \theta | m_0, \tau_0^2) d\theta \\ &= \int p(y | \theta, \sigma^2) \cdot p(\theta | m_0, \tau_0^2) d\theta \\ &= \frac{1}{2\pi\sqrt{\sigma^2\tau_0^2}} \cdot \int \exp \left\{ -\frac{1}{2} \left(\frac{(y-\theta)^2}{\sigma^2} + \frac{(\theta-m_0)^2}{\tau_0^2} \right) \right\} d\theta \\ &\propto \exp \left\{ -\frac{1}{2} \left(-\frac{(\frac{y}{\sigma^2} + \frac{m_0}{\tau_0^2})^2}{\frac{1}{\sigma^2} + \frac{1}{\tau_0^2}} + \frac{y^2}{\sigma^2} \right) \right\} \int \exp \{ (a \cdot \theta - b(y))^2 \} d\theta \\ &\propto \exp \left\{ -\frac{1}{2} \left(-\frac{(y + \frac{m_0\sigma^2}{\tau_0^2})^2}{\sigma^2 + \frac{\sigma^4}{\tau_0^2}} + \frac{y^2}{\sigma^2} \right) \right\} \\ &= \mathcal{N}(y | m_0, \sigma^2 + \tau_0^2), \end{aligned}$$

where we have canceled terms independent from y and completing the square in the final step of deduction. Given $\sigma^2 = 500$, we have:

$$\hat{m}_0 = \bar{x} = 1527.5,$$

$$\hat{\tau}_0^2 = 1878.58 - \sigma^2 = 1378.58,$$

For question (b), the posterior distribution of θ_i given y_i and other hyperparameters is:

$$\begin{aligned} p(\theta|y) &\propto p(\theta) \cdot p(y|\theta) \\ &\propto \exp \left\{ -\frac{(\theta - m_0)^2}{2\tau_0^2} - \frac{(\theta - y)^2}{2\sigma^2} \right\} \\ &= \mathcal{N}(\theta | \frac{m_0\sigma^2 + y\tau_0^2}{\sigma^2 + \tau_0^2}, \frac{\sigma^2\tau_0^2}{\sigma^2 + \tau_0^2}). \end{aligned}$$

For question (c), the interval is $(\mu' - 1.96\sigma', \mu' + 1.96\sigma')$, where μ' and σ' are from question (b).

For question (d), a smaller σ^2 would reduce the ML-II into the ordinary posterior analysis. The parameter σ^2 can be understood as the *noise* on the observations. The less noise we assumed, the more precise the observations are, and the intermedia θ becomes less necessary.

6.3 $\hat{\sigma}_{\text{MLE}}^2$ is biased

Consider:

$$\mathbb{E}[(x_n - \bar{x})^2] = \mathbb{E}[x_n^2] + \frac{1}{N^2} \sum_{j,k=1}^N \mathbb{E}[x_j x_k] - \frac{2}{N} \sum_{j=1}^N \mathbb{E}[x_n x_j].$$

The second term on the r.h.s. contains N components that equal to $\mathbb{E}[x^2] = \mu^2 + \sigma^2$ and $N^2 - N$ components that equal to $\mathbb{E}[x_l x_m] = \mu^2$, where $l \neq m$. For the third term, the numbers are 1 and $N - 1$, thus:

$$\mathbb{E}[(x_n - \bar{x})^2] = \frac{N-1}{N} \sigma^2,$$

so is $\hat{\sigma}_{\text{MLE}}^2$. The unbiased estimation, of which one who has taken several elementary courses on probability and statistics should have been very familiar, is:

$$\frac{N}{N-1} \cdot \hat{\sigma}_{\text{MLE}}^2.$$

6.4 Estimation of σ^2 when μ is known

The likelihood is:

$$p(x|\sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp \left(-\frac{(x - \mu)^2}{2\sigma^2} \right).$$

Taking logarithm and gradient:

$$\frac{d}{d\sigma^2} \ln p(\mathcal{D}|\sigma^2) = -\frac{N}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{n=1}^N (x_n - \mu)^2.$$

Setting it to zero yields:

$$\sigma_{\text{MLE}}^2 = \frac{\sum_{n=1}^N (x_n - \mu)^2}{N}.$$

In this case the MLE is unbiased since:

$$\mathbb{E}(x_n - \mu^2) = \mathbb{E}(x_n^2) + \mu^2 - 2\mu\mathbb{E}(x_n) = \sigma^2.$$

7 Linear regression

7.1 Behavior of training set error with increasing sample size

The statement **the error on the test will always decrease as we get more training data since the model will be better estimated** is not a gold standard. The prerequisites behind this statement are at least:

- The test data shares an identical distribution with the training data. Which, although plausible in general, fails in cases such as adversarial training (where an adversary schemes for samples that hinders the model) or defense against zero-day attack (where the model is required to detect an unknown threat from benign samples solely).
- The model is complex enough to learn the knowledge embedded in the data. For a basic example, a tabular predictor that only saves the latest finite samples and predicts by table looking-up. In this case, the size of the dataset would hardly help. Moreover, the break of the i.i.d. assumption could worsen its performance.

We now reduce the discussion in cases where the assumptions for PAC learning holds. When the training set is small, the trained model is usually over-fitted to the current data set (since the model is very complex), so the accuracy can be relatively high. As the training set increases, the model has to learn to adapt to more general-purpose parameters, thus reducing the overfitting effect laterally, resulting in lower accuracy. As pointed out in Section 7.5.4, increasing the training set is an important method of countering over-fitting besides adding regularizers. Increasing dataset is fundamentally the only solution to models' robustness since one cannot exhaust all possible data augmentation strategies.

7.2 Multi-output linear regression

For multi-output linear regression, if the outputs are independent, then for each output dimension subscripted by j , we have:

$$p(y_j|\mathbf{x}, \mathbf{w}_j) = \mathcal{N}(y_j|\mathbf{x}^T \mathbf{w}_j, \sigma_j^2),$$

then:

$$p(\mathbf{y}|\mathbf{x}, \mathbf{W}) = \prod_{j=1}^M \mathcal{N}(y_j|\mathbf{x}^T \mathbf{w}_j, \sigma_j^2).$$

The independence implies:

$$p(\mathbf{Y}|\mathbf{X}, \mathbf{W}) = \prod_{n=1}^N \prod_{j=1}^M \mathcal{N}(y_{n,j}|\mathbf{x}_n^T \mathbf{w}_j, \sigma_j^2).$$

Taking its logarithm and saving terms dependent on \mathbf{W} , we have:

$$\mathcal{L}(\mathbf{W}) = \sum_{n=1}^N \sum_{j=1}^M \frac{1}{\sigma_j^2} (y_{n,j} - \mathbf{x}_n^T \mathbf{w}_j)^2.$$

Interchanging the order of summation helps to decompose the loss into:

$$\sum_{j=1}^M \mathcal{L}_j(\mathbf{w}_j) = \sum_{j=1}^M \frac{1}{\sigma_j^2} \sum_{n=1}^N (y_{n,j} - \mathbf{x}_n^T \mathbf{w}_j)^2,$$

where each:

$$\mathcal{L}_j(\mathbf{w}_j) = \frac{1}{\sigma_j^2} \sum_{n=1}^N (y_{n,j} - \mathbf{x}_n^T \mathbf{w}_j)^2,$$

is but the MLE loss for a 1D linear regression. Thus the columns of \mathbf{W} can be estimated independently by:

$$\mathbf{w}_j^{\text{MLE}} = (\mathbf{X}\mathbf{X}^T)^{-1} \mathbf{X}\mathbf{Y}_j,$$

where \mathbf{X} is the $D \times N$ design matrix and \mathbf{Y} is a column matrix with length N that embeds the j -th component of the output. This is a little different from the symbols from the textbook but simple α -reductions would eliminate such difference. Equation (7.90) is incorrect by missing one design matrix.

As a compact way of writing \mathbf{W}^{MLE} , we would have:

$$\mathbf{W}^{\text{MLE}} = (\mathbf{X}\mathbf{X}^T)^{-1} \mathbf{X}\mathbf{Y},$$

where \mathbf{Y} is a $N * M$ matrix.

For the case in this exercise, we have $D = 2$, $N = 6$, $M = 2$:

$$\mathbf{X} = \begin{pmatrix} 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 \end{pmatrix},$$

$$\mathbf{Y} = \begin{pmatrix} -1 & -1 \\ -1 & -2 \\ -2 & -1 \\ 1 & 1 \\ 1 & 2 \\ 2 & 1 \end{pmatrix}.$$

Thus:

$$\mathbf{W} = \begin{pmatrix} -\frac{4}{3} & -\frac{4}{3} \\ \frac{4}{3} & \frac{4}{3} \end{pmatrix}.$$

One can observe that two columns for \mathbf{W} are identical, which is obvious by examining the two columns from \mathbf{Y} .

7.3 Centering and ridge regression

The loss is given by:

$$\mathcal{L}(\mathbf{w}, w_0) = (\mathbf{y} - \mathbf{X}^T \mathbf{w} - w_0 \mathbf{1})^T (\mathbf{y} - \mathbf{X}^T \mathbf{w} - w_0 \mathbf{1}).$$

Taking partial gradient w.r.t. w_0 yields:

$$\frac{\partial}{\partial w_0} \mathcal{L} = \frac{\mathbf{1}^T (\mathbf{y} - \mathbf{X}^T \mathbf{w})}{N}.$$

Setting it to zero gives (7.94), where we have made use of $\bar{x} = 0$.

Taking partial gradient w.r.t. \mathbf{w} yields:

$$\mathbf{w}^{\text{MLE}} = (\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I})^{-1} \mathbf{X}(\mathbf{y} - w_0 \mathbf{1}).$$

In the equation (7.95), the observed data has to be centralized.

7.4 MLE for σ^2 for linear regression

The likelihood is given by:

$$\begin{aligned}
 p(\mathcal{D}|\mathbf{w}, \sigma^2) &= p(\mathbf{Y}|\mathbf{w}, \sigma^2, \mathbf{X}) \\
 &= \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}, \sigma^2) \\
 &= \prod_{n=1}^N \mathcal{N}(y_n|\mathbf{w}^T \mathbf{x}_n, \sigma^2) \\
 &= \frac{1}{(2\pi\sigma^2)^{\frac{N}{2}}} \cdot \exp \left\{ -\frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{w}^T \mathbf{x}_n)^2 \right\}.
 \end{aligned}$$

As for σ^2 :

$$\frac{\partial}{\partial \sigma^2} \log p(\mathcal{D}|\mathbf{w}, \sigma^2) = -\frac{N}{2\sigma^2} + \frac{1}{2(\sigma^2)^2} \sum_{n=1}^N (y_n - \mathbf{w}^T \mathbf{x}_n)^2.$$

Setting it to zero yields:

$$\sigma_{\text{MLE}}^2 = \frac{1}{N} \sum_{n=1}^N (y_n - \mathbf{w}^T \mathbf{x}_n)^2.$$

7.5 MLE for the offset term in linear regression

For (7.97), we only have to use the second equation derived in exercise 7.3.

For (7.98), recall the third equation that we have derived in exercise 7.3.

7.6 MLE for simple linear regression

For (7.99), we plug the $D = 1$ assumption into (7.98) so:

$$w_1^{\text{MLE}} = \frac{\sum_{n=1}^N (y_n - \bar{y})(x_n - \bar{x})}{\sum_{n=1}^N (x_n - \bar{x})^2}.$$

Recall that $\bar{x} = \frac{1}{N} \sum_{n=1}^N x_n$ and that for \bar{y} , so (7.99) is proven.

For (7.100), use (7.97) directly.

7.7 Sufficient statistics for online linear regression

For question (a), according to (7.99), w_1 can be estimated from $C_{xy}^{(n)}$ and $C_{xx}^{(n)}$ solely.

For question (b), according to (7.100), w_0 can be estimated from $\bar{x}^{(n)}$, $\bar{y}^{(n)}$ and w_1 . Hence $\bar{x}^{(n)}$, $\bar{y}^{(n)}$, $C_{xy}^{(n)}$ and $C_{xx}^{(n)}$ are necessary.

For question (c), the solution has been given by (7.103)-(7.104). For y , the matter is simply an α -reduction.

For question (d), we need to prove:

$$(n+1)C_{xy}^{(n+1)} = nC_{xy}^{(n)} + x_{n+1}y_{n+1} + n\bar{x}^{(n)}\bar{y}^{(n)} - (n+1)\bar{x}^{(n+1)}\bar{y}^{(n+1)}.$$

Expand the C_{xy} in both two sides then the l.h.s. becomes:

$$\sum_{i=1}^{n+1} (x_i - \bar{x}^{(n+1)})(y_i - \bar{y}^{(n+1)}),$$

the r.h.s. becomes:

$$\sum_{i=1}^n (x_i - \bar{x}^{(n)})(y_i - \bar{y}^{(n)}) + x_{n+1}y_{n+1} + n\bar{x}^{(n)}\bar{y}^{(n)} - (n+1)\bar{x}^{(n+1)}\bar{y}^{(n+1)}.$$

Finally, substituting the average at the $(n+1)$ -th iteration in the l.h.s. and r.h.s. by (7.104), this is sufficient for arriving in (7.105).

For question (e) and (f):

```

1  import math
2  import matplotlib.pyplot as plt
3  x=[94,96,94,95,104,106,108,113,115,121,131]
4  y=[0.47,0.75,0.83,0.98,1.18,1.29,1.40,1.60,1.75,1.90,2.23]
5  bx=(x[0]+x[1])/2
6  by=(y[0]+y[1])/2
7  Cxx=((x[0]-bx)**2+(x[1]-bx)**2)/2
8  Cxy=((y[0]-by)*(x[0]-bx)+(y[1]-by)*(x[1]-bx))/2
9  w1=[]
10 w0=[]
11 for n in list(range(2,11)):
12     bx_=bx+(x[n]-bx)/(n+1)
13     by_=by+(y[n]-by)/(n+1)
14     Cxx=(x[n]*x[n]+n*Cxx+n*bx*bx-(n+1)*bx_*bx_)/(n+1)

```



```

15     Cxy=(x[n]*y[n]+n*Cxy+n*bx*by-(n+1)*bx_*by_)/(n+1)
16     bx=bx_
17     by=by_
18     w1.append(Cxy/Cxx)
19     w0.append(by-w1[n-2]*bx)

```

With: Where we borrow data from exercise 7.8 for illustration. The data




Figure. 4. Exercise 7.7. P1




Figure. 5. Exercise 7.7. P2

in x and y appear in sequential order for Figure. 7. 14. That is to say, if we shuffle x and y accordingly, the estimation for weights is expected to converge faster.

7.8 Bayesian linear regression in 1d with known σ^2

For question (a), the estimation for σ^2 is 0.3173.

For question (b), we have:

$$p(\mathbf{w}) \propto \mathcal{N}(w_1|0, 1) \propto \exp\left\{-\frac{1}{2}w_1^2\right\}.$$

To simplify the algebra, we observe that w_1 and w_0 are independent in this prior. Thus the prior distribution can be reduced to: $\mathcal{N}(w_1|0, 1) \cdot \mathcal{N}(w_0|\gamma, \infty)$, where γ can be an arbitrary finite number, thus:

$$\mathbf{m}_0 = \begin{pmatrix} 0 \\ \gamma \end{pmatrix},$$

$$\mathbf{V}_0 = \begin{pmatrix} 1 & 0 \\ 0 & 0 \end{pmatrix}.$$

For question (c) and (d), we consider the posterior distribution for parameters:

$$\begin{aligned} p(\mathbf{w}|\mathcal{D}, \sigma^2) &= \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{V}_0) \cdot \prod_{n=1}^N \mathcal{N}(y_n|w_0 + w_1x_n, \sigma^2) \\ &\propto \exp\left\{-\frac{w_1^2}{2}\right\} \cdot \prod_{n=1}^N \exp\left\{-\frac{(y_n - w_0 - w_1x_n)^2}{2\sigma^2}\right\}, \end{aligned}$$

where we only maintain terms dependent on w_0 and w_1 . To marginalize out w_0 , we have:

$$\begin{aligned} p(w_1|\mathcal{D}, \sigma^2) &= \int p(w_1, w_0)dw_0 \\ &\propto \exp\left\{-\frac{w_1^2}{2}\right\} \cdot \int \exp\{Aw_1^2 + Bw_0^2 + Cw_0w_1 + Dw_1 + Ew_0 + F\}dw_0 \\ &= \exp\left\{-\frac{w_1^2}{2} + Aw_1^2 + Dw_1 + F\right\} \cdot \int \exp\{Bw_0^2 + Cw_0w_1 + Ew_0\}dw_0 \\ &= \exp\left\{-\frac{w_1^2}{2} + Aw_1^2 + Dw_1 + F - \frac{(Cw_1 + E)^2}{4B}\right\} \\ &\quad \int \exp\left\{Bw_0^2 + (Cw_1 + E)w_0 + \frac{(Cw_1 + E)^2}{4B}\right\}dw_0 \\ &\propto \exp\left\{-\frac{w_1^2}{2} + Aw_1^2 + Dw_1 + F - \frac{(Cw_1 + E)^2}{4B}\right\}. \end{aligned}$$

Hence the posterior distribution over w_1 is a normal distribution. The coefficients for w_1^2 and w_1 in the exponential are respectively:

$$-\frac{1}{2} + A - \frac{C^2}{4B},$$

$$D - \frac{CE}{2B}.$$

Thence its posterior variance is:

$$\frac{1}{1 - 2A + \frac{C^2}{2B}},$$

its mean is:

$$\frac{2BD - CE}{2B - 4AB + C^2}.$$

Finally, let us plug A to E with statistics of \mathcal{D} :

$$A = -\frac{\sum_{n=1}^N x_n^2}{2\sigma^2},$$

$$B = -\frac{N}{2\sigma^2},$$

$$C = -\frac{\sum_{n=1}^N x_n}{\sigma^2},$$

$$D = \frac{\sum_{n=1}^N x_n y_n}{\sigma^2},$$

$$E = \frac{\sum_{n=1}^N y_n}{\sigma^2}.$$

The posterior variance is:

$$\frac{\sigma^2}{\sigma^2 + \sum x^2 - \frac{(\sum x)^2}{N}},$$

from which we observe that, with N grows, the denominator increases as the bound from the Cauchy inequality:

$$\sum_{n=1}^N x_n^2 - \frac{(\sum_{n=1}^N x_n)^2}{N} \geq 0.$$

To put in other words, The larger the Cauchy difference is, the more confidence we have for the estimation of w_1 . Such difference is determined by the variance of the distribution on x . The posterior variance can be reduced to:

$$\frac{\sigma^2}{\sigma^2 + N \text{var}(x)}.$$

Therefore for any fixed generative distribution on x , the uncertainty on w_1 declines as $\mathcal{O}\left\{\frac{1}{N}\right\}$.

7.9 Generative model for linear regression

For question (a), assume that \mathbf{x} and y are jointly subject to a Gaussian:

$$\mathcal{N}\left(\begin{pmatrix} \mathbf{x} \\ y \end{pmatrix} \middle| \begin{pmatrix} \mu_{\mathbf{x}} \\ \mu_y \end{pmatrix}, \begin{pmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{xy}^T & \Sigma_{yy} \end{pmatrix}\right).$$

We know from (4.68) that the marginal distribution of \mathbf{x} and y are $\mathcal{N}(\mathbf{x}|\mu_{\mathbf{x}}, \Sigma_{xx})$ and $\mathcal{N}(y|\mu_y, \Sigma_{yy})$ respectively, this is sufficient for MLE $\mu_{\mathbf{x}}$, μ_y , Σ_{xx} and Σ_{yy} :

$$\begin{aligned} \mu_{\mathbf{x}} &= \frac{1}{N} \cdot \sum_{n=1}^N \mathbf{x}_n, \\ \mu_y &= \frac{1}{N} \cdot \sum_{n=1}^N y_n, \\ \Sigma_{xx} &= \frac{1}{N} \cdot \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T = \frac{1}{N} \cdot \bar{\mathbf{X}}\bar{\mathbf{X}}^T, \\ \Sigma_{yy} &= \frac{1}{N} \cdot \sum_{n=1}^N (y_n - \bar{y})^2 = \frac{1}{N} \bar{\mathbf{Y}}^T \bar{\mathbf{Y}}. \end{aligned}$$

For an estimation of Σ_{xy} , connecting \mathbf{x} and y together then estimating this vector's covariance matrix (this is an ordinary Gaussian model), picking only the first D components from its last column:

$$\Sigma_{xy} = \frac{1}{N} \bar{\mathbf{X}} \bar{\mathbf{Y}}.$$

Finally, plugging our observations into (4.69):

$$\mu_{y|\mathbf{x}} = \mu_y + \Sigma_{yx} \Sigma_{xx}^{-1} (\mathbf{x} - \mu_{\mathbf{x}}),$$

which equals:

$$\bar{y} + \frac{1}{N} \cdot \bar{\mathbf{Y}}^T \bar{\mathbf{X}}^T \cdot N \cdot (\bar{\mathbf{X}} \bar{\mathbf{X}}^T)^{-1} \cdot (\mathbf{x} - \bar{\mathbf{x}}).$$

This is identical to what (7.109)-(7.111) implies, hence the proof is completed.

For question (b), there is no significant difference except that we are now equipped with a distribution of x . This enables us the ability of generating more samples and sheds light on active querying strategies, in which we can query an oracle for samples in order to accelerate the convergence of the learning process.

7.10 Bayesian linear regression using the g-prior

For Bayesian linear regression model, the likelihood is as always:

$$p(\mathcal{D}|\mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(y_n|\mathbf{w}^T \mathbf{x}_n, \sigma^2).$$

The prior distribution is Gaussian-Inverse Gamma distribution:

$$\begin{aligned} p(\mathbf{w}, \sigma^2) &= \text{NIG}(\mathbf{w}, \sigma^2 | \mathbf{w}_0, \mathbf{V}_0, a_0, b_0) \\ &= \mathcal{N}(\mathbf{w} | \mathbf{w}_0, \sigma^2 \mathbf{V}_0) \cdot \text{IG}(\sigma^2 | a_0, b_0) \\ &= \frac{1}{(2\pi)^{\frac{D}{2}}} \frac{1}{|\sigma^2 \mathbf{V}_0|^{\frac{1}{2}}} \exp \left\{ -\frac{1}{2} (\mathbf{w} - \mathbf{w}_0)^T (\sigma^2 \mathbf{V}_0)^{-1} (\mathbf{w} - \mathbf{w}_0) \right\} \cdot \\ &\quad \frac{b_0^{a_0}}{\Gamma(a_0)} (\sigma^2)^{-(a_0+1)} \exp \left\{ -\frac{b_0}{\sigma^2} \right\} \\ &= \frac{b_0^{a_0}}{(2\pi)^{\frac{D}{2}} |\mathbf{V}_0|^{\frac{1}{2}} \Gamma(a_0)} (\sigma^2)^{-(a_0+\frac{D}{2}+1)} \cdot \exp \left\{ -\frac{(\mathbf{w} - \mathbf{w}_0)^T \mathbf{V}_0^{-1} (\mathbf{w} - \mathbf{w}_0) + 2b_0}{2\sigma^2} \right\}. \end{aligned}$$

The posterior distribution takes the form:

$$\begin{aligned} p(\mathbf{w}, \sigma^2 | \mathcal{D}) &\propto p(\mathbf{w}, \sigma^2) p(\mathcal{D} | \mathbf{w}, \sigma^2) \\ &\propto \frac{b_0^{a_0}}{(2\pi)^{\frac{D}{2}} |\mathbf{V}_0|^{\frac{1}{2}} \Gamma(a_0)} (\sigma^2)^{-(a_0+\frac{D}{2}+1)} \cdot \\ &\quad \exp \left\{ -\frac{(\mathbf{w} - \mathbf{w}_0)^T \mathbf{V}_0^{-1} (\mathbf{w} - \mathbf{w}_0) + 2b_0}{2\sigma^2} \right\} \cdot \\ &\quad (\sigma^2)^{-\frac{N}{2}} \cdot \exp \left\{ -\frac{\sum_{n=1}^N (y_n - \mathbf{w}^T \mathbf{x}_n)^2}{2\sigma^2} \right\}. \end{aligned}$$

To decompose the updated hyperparameters from this form, we have to find:

- The exponential of σ^2 .
- The squared term within the exponential.

The exponential of σ^2 in the posterior is:

$$-\left(a_0 + \frac{D}{2} + 1\right) - \frac{N}{2},$$

thus we have:

$$a_N = a_0 + \frac{N}{2} = \frac{N}{2}.$$

The coefficient of $\mathbf{w}^T \mathbf{w}$ in the exponential (as the introducer matrix of the inner product) is:

$$-\frac{\mathbf{V}_0^{-1}}{2\sigma^2} - \frac{\mathbf{X}\mathbf{X}^T}{2\sigma^2},$$

therefore:

$$\mathbf{V}_N^{-1} = \mathbf{V}_0^{-1} + \mathbf{X}\mathbf{X}^T,$$

with $\mathbf{V}_0 = g(\mathbf{X}\mathbf{X}^T)^{-1}$:

$$\mathbf{V}_N = \frac{g}{g+1} (\mathbf{X}\mathbf{X}^T)^{-1}.$$

The coefficient of \mathbf{w}^T in the exponential is:

$$\frac{\mathbf{V}_0^{-1} \mathbf{w}_0}{\sigma^2} + \frac{\sum_{n=1}^N y_n \mathbf{x}_n}{\sigma^2} = \frac{\mathbf{V}_0^{-1} \mathbf{w}_0}{\sigma^2} + \frac{\mathbf{X}\mathbf{Y}}{\sigma^2}.$$

So we have:

$$\mathbf{V}_0^{-1} \mathbf{w}_0 + \mathbf{X}\mathbf{Y} = \mathbf{V}_N^{-1} \mathbf{w}_N.$$

This yields:

$$\begin{aligned} \mathbf{w}_N &= \mathbf{V}_N (\mathbf{V}_0^{-1} \mathbf{w}_0 + \mathbf{X}\mathbf{Y}) \\ &= \frac{g}{g+1} (\mathbf{X}\mathbf{X}^T)^{-1} \left(\frac{\mathbf{X}\mathbf{X}^T}{g} \mathbf{0} + \mathbf{X}\mathbf{Y} \right) \\ &= \frac{g}{g+1} (\mathbf{X}\mathbf{X}^T)^{-1} \mathbf{X}\mathbf{Y}. \end{aligned}$$

Finally, completing the square within the exponential yields:

$$(\mathbf{w} - \mathbf{w}_0)^T \mathbf{V}_0^{-1} (\mathbf{w} - \mathbf{w}_0) + 2b_0 + (\mathbf{w}^T \mathbf{X} - \mathbf{Y})^2 = (\mathbf{w} - \mathbf{w}_N)^T \mathbf{V}_N^{-1} (\mathbf{w} - \mathbf{w}_N) + 2b_N.$$

Plugging in what we have already known about \mathbf{w}_N and \mathbf{V}_N results in:

$$b_N = \frac{\mathbf{Y}^T \mathbf{Y}}{2} + \frac{g}{2(g+1)} \mathbf{w}_{\text{MLE}}^T \mathbf{X}\mathbf{X}^T \mathbf{w}_{\text{MLE}}.$$

Now we have (7.113)-(7.116) established.

8 Logistic regression

Logistic regression is the basic form of classification, which is the fundamental task of machine intelligence. Even for classifying complex data structures, the last layers of the model are usually logistic structures.

8.1 Spam classification using logistic regression

Practice by yourself. I know nothing about MATLAB.

8.2 Spam classification using naive Bayes

Practice by yourself.

8.3 Gradient and Hessian of log-likelihood for logistic regression

For question (a),

$$\frac{d}{da} \sigma(a) = \frac{\exp(-a)}{(1 + \exp(-a))^2} = \frac{1}{1 + e^{-a}} \frac{e^{-a}}{1 + e^{-a}} = \sigma(a) \cdot (1 - \sigma(a)).$$

For question (b),

$$\begin{aligned} g(\mathbf{w}) &= \frac{\partial}{\partial \mathbf{w}} \text{NLL}(\mathbf{w}) \\ &= - \sum_{n=1}^N \frac{\partial}{\partial \mathbf{w}} [y_i \log \mu_i + (1 - y_i) \log(1 - \mu_i)] \\ &= - \sum_{n=1}^N y_i \frac{1}{\sigma_i} \sigma_i(1 - \sigma_i) \cdot \mathbf{x}_i + (1 - y_i) \frac{-1}{1 - \sigma_i} \sigma(1 - \sigma_i) \cdot \mathbf{x}_i \\ &= \sum_{n=1}^N (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i, \end{aligned}$$

where $\sigma_i = \sigma(\mathbf{w}^T \mathbf{x}_i)$.

For question (c), the result is obvious. For an arbitrary vector \mathbf{u} :

$$\begin{aligned}\mathbf{u}^T \mathbf{H} \mathbf{u} &= (\mathbf{X} \mathbf{u})^T \mathbf{S} (\mathbf{X} \mathbf{u}) \\ &= \mathbf{v}^T \mathbf{S} \mathbf{v} \\ &= \sum_{d=1}^D v_d^2 \cdot \mu_d \cdot (1 - \mu_d) \geq 0.\end{aligned}$$

Hence \mathbf{H} is positive definite if all μ_i are within $(0, 1)$, otherwise it is semi-positive definite.

8.4 Gradient and Hessian of log-likelihood for multinomial logistic regression

For question (a), given a sample indexed i we have,

$$\begin{aligned}\mu_{ik} &= \frac{\exp(\mathbf{w}_k^T \mathbf{x}_i)}{\sum_c \exp(\mathbf{w}_c^T \mathbf{x}_i)}, \\ \eta_{ij} &= \mathbf{w}_j^T \mathbf{x}_i.\end{aligned}$$

Now we have:

$$\begin{aligned}\frac{\partial \mu_{ik}}{\partial \eta_{ij}} &= \frac{\frac{\partial \exp(\eta_{ik})}{\partial \eta_{ij}} \cdot \sum_c \exp(\eta_{ic}) - \frac{\partial \sum_c \exp(\eta_{ic})}{\partial \eta_{ij}} \cdot \exp(\eta_{ik})}{(\sum_c \exp(\eta_{ic}))^2} \\ &= \frac{\exp(\eta_{ik}) \cdot \delta_{kj} \cdot \sum_c \exp(\eta_{ic}) - \exp(\eta_{ij}) \cdot \exp(\eta_{ik})}{(\sum_c \exp(\eta_{ic}))^2} \\ &= \mu_{ik} \cdot \delta_{kj} - \mu_{ij} \cdot \mu_{ik},\end{aligned}$$

what dominates is but the elementary calculus.

For question (b), recall that:

$$l(\mathbf{W}) = \sum_{i=1}^N \sum_c y_{ic} \cdot \log \mu_{ic}.$$

Let $l_i(\mathbf{W}) = \sum_c y_{ic} \cdot \log \mu_{ic}$, we are now ready for reduction:

$$\begin{aligned}
\frac{\partial l_i}{\partial \mathbf{w}_j} &= \frac{\partial}{\partial \mathbf{w}_j} \sum_c y_{ic} \cdot \log \mu_{ic} \\
&= \sum_c \frac{y_{ic}}{\mu_{ic}} \cdot \frac{\partial \mu_{ic}}{\partial \eta_{ij}} \cdot \frac{\partial \eta_{ij}}{\partial \mathbf{w}_j} \\
&= \sum_c \frac{y_{ic}}{\mu_{ic}} \cdot \mu_{ic} \cdot (\delta_{cj} - \mu_{ij}) \cdot \mathbf{x}_i \\
&= \sum_c y_{ic} \cdot (1 - \mu_{ij}) \cdot \mathbf{x}_i \\
&= y_{ij} \cdot (1 - \mu_{ij}) \cdot \mathbf{x}_i - \sum_{c \neq j} y_{ic} \cdot \mu_{ij} \cdot \mathbf{x}_i \\
&= y_{ij} \cdot (1 - \mu_{ij}) \cdot \mathbf{x}_i + (y_{ij} - 1) \cdot \mu_{ij} \cdot \mathbf{x}_i \\
&= (y_{ij} - \mu_{ij}) \cdot \mathbf{x}_i.
\end{aligned}$$

Summarizing over i yields (8.126).

For question (c), we have by definition:

$$\mathbf{H}_{c,c'} = \nabla_{\mathbf{w}_{c'}} \nabla_{\mathbf{w}_c} l(\mathbf{W}).$$

Hence we begin with the result from question (b):

$$\begin{aligned}
\nabla_{\mathbf{w}_{c'}} \nabla_{\mathbf{w}_c} l_i &= \frac{\partial}{\partial \mathbf{w}_{c'}} (y_{ic} - \mu_{ic}) \cdot \mathbf{x}_i \\
&= - \frac{\partial}{\partial \mathbf{w}_{c'}} \mu_{ic} \cdot \mathbf{x}_i \\
&= - \frac{\partial \mu_{ic}}{\partial \eta_{ic'}} \frac{\partial \eta_{ic'}}{\partial \mathbf{w}_{c'}} \cdot \mathbf{x}_i \\
&= - \mu_{ic} (\delta_{cc'} - \mu_{ic'}) \cdot \mathbf{x}_i \mathbf{x}_i^T,
\end{aligned}$$

where in the last step we have to adopt the outer product to span the Hessian. Summarizing over i yields the desired result (8.127).

8.5 Symmetric version of l2 regularized multinomial logistic regression

We borrow the results from exercise 8.5. Once the l_2 prior is introduced, the likelihood becomes:

$$l_2 = l - \sum_c \lambda \mathbf{w}_c^T \mathbf{w}_c.$$

Therefore (8.126) becomes:

$$\nabla_{\mathbf{w}_c} l_2 = \sum_i (y_{ic} - \mu_{ic}) \cdot \mathbf{x}_i - \lambda \mathbf{w}_c.$$

At the unique optimum we have $\forall c$,

$$\nabla_{\mathbf{w}_c} l_2 = 0,$$

which is identical to:

$$\hat{w}_{c,j} = \frac{1}{\lambda} \sum_i (y_{ic} - \mu_{ic}) \cdot x_{ij}.$$

Therefore:

$$\sum_c \hat{w}_{c,j} = \frac{1}{\lambda} \sum_i \left[\sum_c (y_{ic} - \mu_{ic}) \right] x_{ij},$$

whose value is zero since $\sum_c y_{ic} = \sum_c \mu_{ic} = 1$.

8.6 Elementary properties of l2 regularized logistic regression

For question (a), the Hessian of l_2 regularized negative log likelihood is:

$$\mathbf{H} + \lambda \mathbf{I},$$

where \mathbf{H} , following the derivation in exercise 8.3, is at least semi-positive definite. So the Hessian for this model is strictly (for non-trivial λ) positive definite, there is a unique optimal solution. The answer is False.

For question (b), the result is not necessarily true. For a sparse optimum, one should resort to the LASSO model, where a Laplace prior is exerted on weights.

For question (c), if $\lambda = 0$ then the model reduces to ordinary logistic regression. If the dataset is linearly separable then there exists \mathbf{w}' such that $\forall i$:

$$\mathbf{w}'^T \mathbf{x}_i \cdot y_i \geq 0.$$

Now for an arbitrary number $\alpha > 0$, the weights $\alpha \cdot \mathbf{w}'$ also meets the separation condition. Let $\alpha \rightarrow \infty$ justifies that the statement in (c) is True.

For question (d), the statement is False since the model now has to trade-off between fitness and prior knowledge. Concretely, we can prove the other statement: *as we increase λ , the likelihood of the training dataset monotonically decreases*.

Assume that $\hat{\mathbf{w}}_1$ minimize the (8.131) with λ_1 , denoted by J_1 . Now increase λ_1 to λ_2 , we are now optimizing the loss:

$$J_2(\mathbf{w}) = -l(\mathbf{w}, \mathcal{D}_{\text{train}}) + \lambda_2 \mathbf{w}^T \mathbf{w},$$

whose optimal solution is denoted by $\hat{\mathbf{w}}_2$.

If $\hat{\mathbf{w}}_1 = \hat{\mathbf{w}}_2$ then we already have:

$$l(\hat{\mathbf{w}}_1, \mathcal{D}_{\text{train}}) = l(\hat{\mathbf{w}}_2, \mathcal{D}_{\text{train}}).$$

Otherwise, we would have:

$$-l(\hat{\mathbf{w}}_1, \mathcal{D}_{\text{train}}) + \lambda_2 \hat{\mathbf{w}}_1^T \hat{\mathbf{w}}_1 > -l(\hat{\mathbf{w}}_2, \mathcal{D}_{\text{train}}) + \lambda_2 \hat{\mathbf{w}}_2^T \hat{\mathbf{w}}_2.$$

If $l(\hat{\mathbf{w}}_2, \mathcal{D}_{\text{train}}) > l(\hat{\mathbf{w}}_1, \mathcal{D}_{\text{train}})$ then we would have:

$$\Delta = l(\hat{\mathbf{w}}_2, \mathcal{D}_{\text{train}}) - l(\hat{\mathbf{w}}_1, \mathcal{D}_{\text{train}}) > \lambda_2 (\hat{\mathbf{w}}_2^T \hat{\mathbf{w}}_2 - \hat{\mathbf{w}}_1^T \hat{\mathbf{w}}_1).$$

Finally, consider:

$$J_1(\hat{\mathbf{w}}_1) - J_1(\hat{\mathbf{w}}_2) = \Delta + \lambda_1 (\hat{\mathbf{w}}_1^T \hat{\mathbf{w}}_1 - \hat{\mathbf{w}}_2^T \hat{\mathbf{w}}_2).$$

If $\hat{\mathbf{w}}_1^T \hat{\mathbf{w}}_1 > \hat{\mathbf{w}}_2^T \hat{\mathbf{w}}_2$ then $\hat{\mathbf{w}}_1$ is not the optimum of J_1 and we arrive in a contradiction. Otherwise:

$$\lambda_1 (\hat{\mathbf{w}}_2^T \hat{\mathbf{w}}_2 - \hat{\mathbf{w}}_1^T \hat{\mathbf{w}}_1) < \lambda_2 (\hat{\mathbf{w}}_2^T \hat{\mathbf{w}}_2 - \hat{\mathbf{w}}_1^T \hat{\mathbf{w}}_1) < \Delta.$$

Hence we still have the optimality of $\hat{\mathbf{w}}_1$ fail. This finishes the proof.

For question (e), the statement is False. This can be easily shown by imagine $\lambda \rightarrow \infty$.

8.7 Regularizing separate terms in 2d logistic regression

The decision boundary is given by:

$$w_0 + w_1 \cdot x_1 + w_2 \cdot x_2 = 0,$$

hence a straight line in the $X_1 - X_2$ plane.

For question (a), the decision boundary is an arbitrary line.

For question (b), the decision boundary is a line that passes the origin.

For question (c), the decision boundary is a line parallel to the X_1 axis.

For question (e), the decision boundary is a line parallel to the X_2 axis.

9 Generalized linear models and the exponential family

9.1 Conjugate prior for univariate Gaussian in exponential family form

Recall that the 1d Gaussian distribution is:

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp \left\{ -\frac{1}{2\sigma^2}(x - \mu)^2 \right\}.$$

Rewrite it into the standard exponential family form:

$$p(x|\mu, \sigma^2) = \exp \left\{ -\frac{1}{2\sigma^2}x^2 + \frac{\mu}{\sigma^2}x - \left\{ \frac{\mu^2}{2\sigma^2} + \frac{\ln(2\pi\sigma^2)}{2} \right\} \right\}.$$

With $\lambda = \frac{1}{\sigma^2}$, denote:

$$\begin{aligned} \theta &= \left(-\frac{\lambda}{2}, \lambda\mu\right)^T, \\ \phi(x) &= (x^2, x)^T, \\ A(\theta) &= \frac{\lambda\mu^2}{2} + \frac{\ln(2\pi)}{2} - \frac{\ln \lambda}{2}. \end{aligned}$$

Now consider the likelihood w.r.t. a dataset $\mathcal{D} = \{x_n\}_{n=1}^N$:

$$\log p(\mathcal{D}|\theta) = \theta^T \left(\sum_{n=1}^N \phi(x_n) \right) - N \cdot A(\theta).$$

The prior distribution of θ should satisfy the following variational form:

$$\begin{aligned} p(\theta|\mathbf{v}, M) &= \exp \{ \theta^T \mathbf{v} + M \cdot A(\theta) \} \\ &= \exp \{ \theta_1 \cdot v_1 + \theta_2 \cdot v_2 + M \cdot A(\theta) \} \\ &= \exp \left\{ -\frac{\lambda v_1}{2} + \lambda \mu v_2 + M \cdot \left(\frac{\lambda \mu^2}{2} + \frac{\ln(2\pi)}{2} - \frac{\ln \lambda}{2} \right) \right\} \\ &= \exp \left\{ -\frac{\lambda v_1}{2} - \frac{M \ln \lambda}{2} + \lambda \mu v_2 + \frac{M \lambda \mu^2}{2} \right\} \cdot \exp \left\{ \frac{\ln 2\pi}{2} \right\} \\ &\propto \exp \left\{ -\frac{\lambda v_1}{2} - \frac{M \ln \lambda}{2} \right\} \cdot \exp \left\{ \lambda \mu v_2 + \frac{M \lambda \mu^2}{2} \right\}. \end{aligned}$$

The first term, in which λ is the target variable, takes the form of a Gamma distribution since:

$$\text{Ga}(\lambda|\alpha, \beta) \propto \exp \{ -\beta\lambda + (\alpha - 1) \ln \lambda \}.$$

The second term is just another Gaussian distribution since the sufficient statistics are μ and μ^2 . Combine these two observations together, we have:

$$p(\theta|\mathbf{v}, M) = \text{Ga}(\lambda|\alpha, \beta) \cdot \mathcal{N}(\mu|\gamma, \tau^2).$$

The transformations between variables are:

$$\begin{aligned}\beta &= \frac{v_1}{2}, \\ \alpha &= 1 - \frac{M}{2}, \\ \tau^2 &= -\frac{1}{M\lambda}, \\ \gamma &= -\frac{v_2}{M}.\end{aligned}$$

In case the variance of the prior for μ is written in the information form, the precision can be written by: $\lambda \cdot (2\alpha - 2)$.

9.2 The MVN is in the exponential family

The sufficient statistics for \mathbf{x} is:

$$\phi(\mathbf{x}) = (x_1, x_2, \dots, x_M, x_1x_1, \dots, x_Mx_M)^T,$$

including altogether $M + M^2$ components.

Expand the information form for a Gaussian, the linear terms within the exponential are:

$$-\frac{1}{2} \sum_{m,m'} \Lambda_{m,m'} x_m x_{m'} + \sum_m x_m \sum_{m'} \Lambda_{m,m'} \mu_{m'},$$

hence the parameters can be collected by:

$$\theta = \left(\sum_{m'} \Lambda_{1,m'} \mu_{m'}, \dots, \sum_{m'} \Lambda_{M,m'} \mu_{m'}, \Lambda_{1,1}, \dots, \Lambda_{M,M} \right).$$

θ contains exactly the same information as μ, Σ . Finally, we have:

$$A(\theta) = \frac{1}{2} \mu^T \Lambda \mu + \frac{M}{2} \ln(2\pi) + \frac{1}{2} \ln |\Sigma|.$$

10 Directed graphical models(Bayes nets)

Graphical models are at the central position for knowledge-based machine learning. Since with graph, most practical relationships among structured variables can be described, learned and crystallized. Although graphical models are not influential regarding more complex scenarios, it remains the most important element in this textbook.

10.1 Marginalizing a node in a DGM

Since Figure 10.14.(a) implies:

$$\begin{aligned}
 p(A, B, C, D, E, F, X) &= \sum_x p(A, B, C, D, E, F, x) \\
 &= p(A) \cdot p(B) \cdot p(C) \cdot p(D) \sum_x p(x|A, B) \cdot p(E|C, x) \cdot p(F|x, D) \\
 &= p(A) \cdot p(B) \cdot p(C) \cdot p(D) \cdot f(A, B, C, D, E, F).
 \end{aligned}$$

Thus the marginalized graph has to decompose in such a way with clique $\{A, B, C, D, E, F\}$. One option is as follows: The edge from E to F is neces-




Figure. 6. Exercise 10.1.

sary since the joint probability cannot be decomposed into $f(A, B, C, D, E) \cdot g(A, B, C, D, F)$.

10.2 Bayes Ball

For question (a), for node C , there is a path:

$$C \rightarrow A,$$

so C is not independent from A given B .

For node D , there is a path:

$$D \rightarrow G \rightarrow E \rightarrow C \rightarrow A,$$

so D is not independent from A given B .

For node E , there is a path:

$$E \rightarrow C \rightarrow A,$$

so E is not independent from A given B .

For node F , there is a path:

$$F \rightarrow C \rightarrow A,$$

so F is not independent from A given B .

For node G , there is a path:

$$G \rightarrow E \rightarrow C \rightarrow A,$$

so G is not independent from A given B .

For node H , there is a path:

$$H \rightarrow F \rightarrow C \rightarrow A,$$

so H is not independent from A given B .

For node I , all pathes from I to A have to go through a collider whose intermedium is not shadowed (G and H), hence I is independent from A given B .

For question (b), for node B , there exist a path to A ,

$$B \rightarrow D \rightarrow G \rightarrow A,$$

which contains a collider structure that is not blocked since G 's child, J , is shadowed. Hence B is not independent from A given J .

For node C , each of its pathes to A must bypass $F \rightarrow I \rightarrow E$, which is a blocked collider, hence C is conditionally independent from A .

For node D , the result is the same as B .

For E , pathes:

$$E \rightarrow B \rightarrow D \rightarrow G \rightarrow A,$$

is unblocked, hence E is not independent from A given J .

For F , the result is the same as C .

For G , it is obviously not independent from A given J .

For H , path:

$$H \rightarrow D \rightarrow G \rightarrow A,$$

is unblocked, hence H is not independent from A given J .

For I , the path:

$$I \rightarrow E \rightarrow B \rightarrow D \rightarrow G \rightarrow A,$$

is unblocked, hence I is not independent from A given J .

10.3 Markov blanket for a DGM

The trick in the required reduction is to partition all variables into:

$$p(\mathcal{X}) = p(\mathcal{A}) \cdot p(X_i | \text{Pa}(X_i)) \prod_{Y_j \in \text{ch}(X_i)} p(Y_j | \text{Pa}(Y_j)) \cdot p(\mathcal{B} | \mathcal{A}, \mathcal{Y}),$$

where $\mathcal{Y} = \text{ch}(X_i)$, \mathcal{A} contains the collection of all variable that are topological smaller than X_i , hence $\text{Pa}(X_i) \subset \mathcal{A}$. Moreover, we have \mathcal{A} incorporate:

$$\cup_{Y_j \in \text{ch}(X_i)} \text{Pa}(Y_j) - \{X_i\} - \cup_{Y_j \in \text{ch}(X_i)} \text{Pa}\{Y_j\}.$$

So all elements of \mathcal{A} can be safely evaluated before X_i . Then all X_i 's children can be computed according to a topological order. Finally, the rest variables \mathcal{B} can be evaluated.

We have:

$$\begin{aligned} p(X_i | X_{-i}) &= \frac{p(\mathcal{X})}{p(X_{-i})} \\ &= \frac{p(\mathcal{X})}{\sum_{x_i} p(\mathcal{X}, X_i = x_i)}. \end{aligned}$$

Eliminating the terms $p(\mathcal{A})$ and $p(\mathcal{B} | \mathcal{A}, \mathcal{Y})$ from both the numerator and the denominator we have (10.58).

10.4 Hidden variables in DGMs

For question (a), $p(X_i)$ where $i = 1, 2, 3$ include three parameters. The term $p(H = h|X_{1:3})$ include $2^3 = 8$ free parameters. Each $p(X_i|H = h)$, where $i = 4, 5, 6$ includes two parameters. Thus we have:

$$3 + 8 + 3 * 2 = 17,$$

free parameters.

For question (b), $p(X_i)$ where $i = 1, 2, 3$ include three parameters. $p(X_4|X_{1:3})$ contains 2^3 free parameters. $p(X_5|X_{1:4})$ contains 2^4 free parameters. $p(X_6|X_{1:4})$ contains 2^5 free parameters. Thus we have:

$$3 + 2^3 + 2^4 + 2^5 = 59,$$

free parameters.

For question (c), the second model is easier since we only have to estimate six terms independently by counting.

10.5 Bayes nets for a rainy day

We first write down the entire probability:

$$p(V, G, R, S) = p(V) \cdot p(G) \cdot p(R|V, G) \cdot p(S|G).$$

For question (a), we have:

$$\begin{aligned} p(S = 0|V = 1) &= \frac{p(S = 0, V = 1)}{p(V = 1)} \\ &= \frac{1}{p(V = 1)} \sum_{R, G \in \{0,1\}} p(V = 1, G, R, S = 0) \\ &= 0.2\alpha\gamma + 0.8\alpha\gamma + 0.1(1 - \alpha)\beta + 0.9(1 - \alpha)\beta \\ &= \alpha\gamma + (1 - \alpha)\beta. \end{aligned}$$

For question (b), the answer is exactly the same. Since given G , S is independent from V .

For question (c), we independently estimate δ , α , γ and β :

$$\begin{aligned}\delta &= 1, \\ \alpha &= \frac{1}{3}, \\ \gamma &= 1, \\ \beta &= 0.\end{aligned}$$

from counting.

10.6 Fishing nets

For question (a), we are to compute the posterior:

$$p(X_2|X_4 = \text{thin}, X_1 \sim \pi),$$

where:

$$\pi = (0.5, 0, 0, 0.5)^T.$$

This shall be done straightforwardly:

$$\begin{aligned}p(X_2 = \text{salmon}|X_4 = \text{thin}, X_1 \sim \pi) &= \sum_{x_1, x_3} p(X_2 = \text{salmon}, X_1 = x_1, X_3 = x_3|X_4 = \text{thin}) \\ &= \frac{p(X_1 = \text{winter}, X_2 = \text{salmon}, X_3 = \text{light}, X_4 = \text{thin})}{p(X_4 = \text{thin})} \\ &+ \frac{p(X_1 = \text{winter}, X_2 = \text{salmon}, X_3 = \text{medium}, X_4 = \text{thin})}{p(X_4 = \text{thin})} \\ &+ \frac{p(X_1 = \text{winter}, X_2 = \text{salmon}, X_3 = \text{dark}, X_4 = \text{thin})}{p(X_4 = \text{thin})} \\ &+ \frac{p(X_1 = \text{autumn}, X_2 = \text{salmon}, X_3 = \text{light}, X_4 = \text{thin})}{p(X_4 = \text{thin})} \\ &+ \frac{p(X_1 = \text{autumn}, X_2 = \text{salmon}, X_3 = \text{medium}, X_4 = \text{thin})}{p(X_4 = \text{thin})} \\ &+ \frac{p(X_1 = \text{autumn}, X_2 = \text{salmon}, X_3 = \text{dark}, X_4 = \text{thin})}{p(X_4 = \text{thin})} \\ &= \frac{0.5 * 0.6 * 1.7}{p(X_4 = \text{thin})}.\end{aligned}$$

While that for sea bass is:

$$p(X_2 = \text{salmon}|X_4 = \text{thin}, X_1 \sim \pi) = \frac{0.5 * 0.05 * 0.3}{p(X_4 = \text{thin})},$$

hence the fish is likely to be salmon.

For question (b),

$$\begin{aligned}
 p(X_1 = \text{winter} | X_3 = \text{medium}, X_4 = \text{thin}) &= \sum_{x_2} p(X_1 = \text{winter}, X_2 = x_2 | X_3 = \text{medium}, X_4 = \text{thin}) \\
 &= \frac{0.25 * 0.9 * 0.33 * 0.6 + 0.25 * 0.1 * 0.1 * 0.05}{p(X_3 = \text{medium}, X_4 = \text{thin})} \\
 &= \frac{0.04468}{p(X_3 = \text{medium}, X_4 = \text{thin})}.
 \end{aligned}$$

Analogously, we have:

$$\begin{aligned}
 p(X_1 = \text{spring} | X_3 = \text{medium}, X_4 = \text{thin}) &= \frac{0.01573}{p(X_3 = \text{medium}, X_4 = \text{thin})}, \\
 p(X_1 = \text{summer} | X_3 = \text{medium}, X_4 = \text{thin}) &= \frac{0.02055}{p(X_3 = \text{medium}, X_4 = \text{thin})}, \\
 p(X_1 = \text{autumn} | X_3 = \text{medium}, X_4 = \text{thin}) &= \frac{0.03985}{p(X_3 = \text{medium}, X_4 = \text{thin})}.
 \end{aligned}$$

Therefore the season is likely to be winter. Intuitively, a fish that is thin and medium lightness is likely to be a salmon. And the season in which a salmon is caught is most likely to be winter.

10.7 Removing leaves in BN20 networks

For question (a), note that given \mathbf{z} , all components from \mathbf{x} are mutually independent, hence:

$$p(\mathbf{Z} | X_1, X_2, X_4) = \frac{1}{p(X_1, X_2, X_4)} \cdot p(\mathbf{Z}) \cdot p(X_1, X_2, X_4 | \mathbf{Z}),$$

according to Figure 10.16.(b), while that computed w.r.t. Figure 10.16.(a) is:

$$\begin{aligned}
 p(\mathbf{Z} | X_1, X_2, X_4) &= \sum_{x_3, x_5} p(\mathbf{Z}, X_3 = x_3, X_5 = x_5 | X_1, X_2, X_4) \\
 &= \frac{1}{p(X_1, X_2, X_4)} \cdot \sum_{x_3, x_5} p(\mathbf{Z}) \cdot p(X_1, X_2, X_4 | \mathbf{Z}) p(X_3 = x_3, X_5 = x_5 | \mathbf{Z}) \\
 &= \frac{1}{p(X_1, X_2, X_4)} \cdot p(\mathbf{Z}) \cdot p(X_1, X_2, X_4 | \mathbf{Z}) \sum_{x_3, x_5} p(X_3 = x_3, X_5 = x_5 | \mathbf{Z}) \\
 &= \frac{1}{p(X_1, X_2, X_4)} \cdot p(\mathbf{Z}) \cdot p(X_1, X_2, X_4 | \mathbf{Z}).
 \end{aligned}$$

This completes the proof.

For question (b), consider the following decomposition of the joint probability:

$$\begin{aligned} p(\mathbf{X}_{\text{on}}, \mathbf{X}_{\text{off}}, \mathbf{Z}) &= p(\mathbf{Z}) \cdot \prod_{X \in \mathbf{X}_{\text{on}}} p(X|\mathbf{Z}) \cdot \prod_{Y \in \mathbf{X}_{\text{off}}} p(Y|\mathbf{Z}) \\ &= \prod_{Z \in \mathbf{Z}} p(Z) \cdot \prod_{X \in \mathbf{X}_{\text{on}}} p(X|\mathbf{Z}) \cdot \prod_{Y \in \mathbf{X}_{\text{off}}} \prod_{Z \in \mathbf{Z}} \theta_{Z,Y}, \end{aligned}$$

where $\theta_{Z,Y}$ is defined from the correlation between disease Z and symptom Y . If Z and Y are independent then $\theta_{Z,Y} = 1$.

Now:

$$p(\mathbf{X}_{\text{on}}, \mathbf{X}_{\text{off}}, \mathbf{Z}) = \prod_{Z \in \mathbf{Z}} p(Z) \left[\prod_{Y \in \mathbf{X}_{\text{off}}} \theta_{Z,Y} \right] \cdot \prod_{X \in \mathbf{X}_{\text{on}}} p(X|\mathbf{Z}).$$

By changing the prior over Z from $p(Z)$ into $p(Z) [\prod_{Y \in \mathbf{X}_{\text{off}}} \theta_{Z,Y}]$ we complete the proof.

10.8 Handling negative findings in the QMR network

Counting and conducting MLE on $p(\mathbf{d})$ is of complexity $\mathcal{O}(|\mathbf{d}|)$. For each disease, it costs $\mathcal{O}(|\mathbf{f}^-|)$ to MLE all corresponding symptoms. Hitherto we have completed the proof.

10.9 Moralization does not introduce new independence statements

To prove that:

$$\text{CI}(M) \subset \text{CI}(G),$$

it is sufficient to prove that:

$$\neg \text{CI}(G) \subset \neg \text{CI}(M),$$

where $\neg \text{CI}(G)$ is the collection of denials of conditional independence from G .

A statement in $\neg \text{CI}(G)$ taking form:

$$\neg \mathcal{A} \perp \mathcal{B} | \mathcal{C},$$

is tantamount to that: *there exists one unblocked path, p , from \mathcal{A} to \mathcal{B} .* Consider this p in M , where the directions on the edges have vanished. The only possibility that this path in M is blocked is that p encounters a collider structure in G and the intermedium node m , or some of its descents is in \mathcal{E} . In this case there exists an extra edge between the parents of m in M , hence we replace the two edges connecting m and m 's parents with the extra edge. This procedure ends up with another unblocked edge, hence we have the statement $\neg\text{CI}(G) \subset \neg\text{CI}(M)$ completed. This finishes the proof of the original proposition.

11 Mixture models and the EM algorithm

EM is perhaps the most important and powerful that probabilistic machine learning theory has provided. Its intuitive conciseness and analytic elegance make it the most popular method for general graph models. Apart from the properties covered in the textbook, the speed of convergence for EM is also a variable of significance. The speed of convergence for EM can be easily derived from the Hessian of the likelihood, combined with elementary matrix algebra, which has been discussed in Dempster's pioneering papers. Without an involved understanding of EM, one can hardly proceed to absorb variational inference or further stochastic learning algorithms.

Despite formal complexity, the intuition behind EM is simple and straightforward. It originates from the conjugate iterative optimization problem, which has been discussed by the optimization community for centuries. Instead of remembering all the symbolic deductions, it is this intuition that should be engraved into one's mind.

11.1 Student T as infinite mixture of Gaussian

Recall that the 1d Student-t distribution takes the form:

$$\text{St}(x|\mu, \sigma^2, v) = \frac{\Gamma(\frac{v}{2} + \frac{1}{2})}{\Gamma(\frac{v}{2})} \left(\frac{1}{\pi v \sigma^2} \right)^{\frac{1}{2}} \left(1 + \frac{(x - \mu)^2}{v \sigma^2} \right)^{-\frac{v+1}{2}}.$$

While the r.h.s. of (11.61) is:

$$\begin{aligned} & \int \mathcal{N}(x|\mu, \frac{\sigma^2}{z}) \cdot \text{Ga}\left(z|\frac{v}{2}, \frac{v}{2}\right) dz \\ &= \int \frac{\sqrt{z}}{\sqrt{2\pi\sigma^2}} \cdot \exp\left\{-\frac{z}{2\sigma^2}(x - \mu)^2\right\} \cdot \frac{(\frac{v}{2})^{\frac{v}{2}}}{\Gamma(\frac{v}{2})} \cdot z^{\frac{v}{2}-1} \cdot \exp\left\{-\frac{vz}{2}\right\} dz \\ &= \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \frac{(\frac{v}{2})^{\frac{v}{2}}}{\Gamma(\frac{v}{2})} \int z^{\frac{v-1}{2}} \cdot \exp\left\{-\left(\frac{v}{2} + \frac{(x - \mu)^2}{2\sigma^2}\right)z\right\} dz. \end{aligned}$$

Instead of integrating analytically, we observe that the term being integrated takes the form of a Gamma density:

$$\text{Ga}\left(z|\frac{v+1}{2}, \frac{v}{2} + \frac{(x - \mu)^2}{2\sigma^2}\right) = \frac{\left(\frac{v}{2} + \frac{(x - \mu)^2}{2\sigma^2}\right)^{\frac{v+1}{2}}}{\Gamma\left(\frac{v+1}{2}\right)} \cdot z^{\frac{v-1}{2}} \cdot \exp\left\{-\left(\frac{v}{2} + \frac{(x - \mu)^2}{2\sigma^2}\right)z\right\}.$$

Therefore the r.h.s. of (11.61) becomes:

$$\begin{aligned} \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \frac{\left(\frac{v}{2}\right)^{\frac{v}{2}}}{\Gamma\left(\frac{v}{2}\right)} \cdot \frac{\Gamma\left(\frac{v+1}{2}\right)}{\left(\frac{v}{2} + \frac{(x-\mu)^2}{2\sigma^2}\right)^{\frac{v+1}{2}}} &= \frac{\Gamma\left(\frac{v+1}{2}\right)}{\Gamma\left(\frac{v}{2}\right)} \cdot \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \frac{\left(\frac{v}{2}\right)^{\frac{v}{2}}}{\left(\frac{v}{2} + \frac{(x-\mu)^2}{2\sigma^2}\right)^{\frac{v+1}{2}}} \\ &= \frac{\Gamma\left(\frac{v+1}{2}\right)}{\Gamma\left(\frac{v}{2}\right)} \cdot \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \left(\frac{v}{2}\right)^{-\frac{1}{2}} \left(1 + \frac{(x-\mu)^2}{2v\sigma^2}\right)^{-\frac{v+1}{2}}. \end{aligned}$$

This completes the proof.

11.2 EM for mixture of Gaussians

We are to optimize the following target w.r.t. θ :

$$\begin{aligned} Q(\theta, \theta^{\text{old}}) &= \mathbb{E}_{p(\mathbf{z}|\mathcal{D}, \theta^{\text{old}})} \left[\sum_{n=1}^N \log p(\mathbf{x}_n, \mathbf{z}_n | \theta) \right] \\ &= \sum_{n=1}^N \mathbb{E}_{p(\mathbf{z}|\mathcal{D}, \theta^{\text{old}})} \left[\log \prod_{k=1}^K (\pi_k \cdot p(\mathbf{x}_n | z_k, \theta))^{z_{nk}} \right] \\ &= \sum_{n=1}^N \sum_{k=1}^K \mathbb{E}_{p(\mathbf{z}|\mathcal{D}, \theta^{\text{old}})} [z_{nk} \cdot \log(\pi_k \cdot p(\mathbf{x}_n | z_k, \theta))] \\ &= \sum_{n=1}^N \sum_{k=1}^K \mathbb{E}_{p(\mathbf{z}|\mathcal{D}, \theta^{\text{old}})} [z_{nk}] \cdot \log(\pi_k \cdot p(\mathbf{x}_n | z_k, \theta)) \\ &= \sum_{n=1}^N \sum_{k=1}^K r_{nk} \cdot \log(\pi_k \cdot p(\mathbf{x}_n | z_k, \theta)), \end{aligned}$$

where:

$$r_{nk} = p(z_{nk} = 1 | \mathbf{x}_n, \theta^{\text{old}}).$$

(Recall the graphical structure of GMM model. \mathbf{z}_n is the one-hot variable that encodes the belonging of sample \mathbf{x}_n to the centroids.) When the base distribution $p(\mathbf{x}|\mathbf{z}, \theta)$ is Gaussian, consider the terms involving μ_k and Σ_k in $Q(\theta, \theta^{\text{old}})$ first (adopting non-information prior):

$$\begin{aligned} \sum_{n=1}^N r_{nk} \cdot \log p(\mathbf{x}_n | z_k = 1, \theta) &= \sum_{n=1}^N r_{nk} \cdot \log \mathcal{N}(\mathbf{x}_n | \mu_k, \sigma_k^2) \\ &= \sum_{n=1}^N r_{nk} \cdot \left(C - \frac{1}{2} \log |\Sigma_k| - \frac{1}{2} (\mathbf{x}_n - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_n - \mu_k) \right) \\ &= \mathcal{L}(\mu_k, \Sigma_k). \end{aligned}$$

Optimizing this target w.r.t. μ_k and Σ_k is tantamount to optimizing the mean and covariance of a weighted Gaussian model, hence:

$$\frac{\partial \mathcal{L}}{\partial \mu_k} = \sum_{n=1}^n r_{nk} \cdot \Sigma^{-1}(\mathbf{x}_n - \mu_k).$$

Setting it to zero yields:

$$\mu_k = \frac{\sum_{n=1}^N r_{nk} \cdot \mathbf{x}_n}{\sum_{n=1}^N r_{nk}}.$$

Finally:

$$\frac{\partial \mathcal{L}}{\partial \Lambda_k} = \sum_{n=1}^N r_{nk} \left(\frac{1}{2} \Lambda_k^{-1} - \frac{1}{2} (\mathbf{x} - \mu_k)(\mathbf{x} - \mu_k)^T \right),$$

where $\Lambda_k = \Sigma_k^{-1}$. Setting it to zero yields:

$$\Sigma_k = \Lambda_k^{-1} = \frac{\sum_{n=1}^N r_{nk} (\mathbf{x} - \mu_k)(\mathbf{x} - \mu_k)^T}{\sum_{n=1}^N r_{nk}}.$$

So far we have proven (11.114) and (11.115).

11.3 EM for mixtures of Bernoullis

For the mixture of Bernoullis model, consider K bases, from which each is a Bernoulli distribution:

$$\text{Ber}(x|\theta_k) = \theta_k^{\mathbb{I}(x=1)} \cdot (1 - \theta_k)^{\mathbb{I}(x=0)}.$$

The auxiliary function, whom we are to optimize w.r.t. θ is:

$$\begin{aligned} Q(\theta, \theta^{\text{old}}) &= \mathbb{E}_{p(\mathbf{z}|\mathcal{D}, \theta^{\text{old}})} \left[\sum_{n=1}^N \log p(x_n, \mathbf{z}_n | \theta) \right] \\ &= \sum_{n=1}^N \sum_{k=1}^K \mathbb{E}[z_{nk}] \cdot (\log \pi_k + \mathbb{I}(x_n = 1) \log \theta_k + \mathbb{I}(x_n = 0) \log(1 - \theta_k)) \\ &= \sum_{n=1}^N \sum_{k=1}^K r_{nk} \cdot (\log \pi_k + \mathbb{I}(x_n = 1) \theta_k + \mathbb{I}(x_n = 0) (1 - \theta_k)). \end{aligned}$$

Taking differential w.r.t. θ_k

$$\frac{\partial Q}{\partial \theta_k} = \sum_{n=1}^N r_{nk} \cdot \left(\mathbb{I}(x_n = 1) \frac{1}{\theta_k} - \mathbb{I}(x_n = 0) \frac{1}{1 - \theta_k} \right),$$

set it to zero:

$$\theta_k = \frac{\sum_{n=1}^N r_{nk} \mathbb{I}(x_n = 1)}{\sum_{n=1}^N r_{nk}}.$$

This is exactly (11.116) modules α -reduction.

If a Beta(α_k, β_k) prior is introduced for each base then we introduce $\alpha_k - 1$ positive samples and $\beta_k - 1$ negative samples into the computation, this is tantamount to setting $r_{nk} = 1$ for $n = N + 1, \dots, N + \alpha_k + \beta_k - 2$, so:

$$\theta_k = \frac{\sum_{n=1}^N r_{nk} \mathbb{I}(x_n = 1) + \alpha_k - 1}{\sum_{n=1}^N r_{nk} + \alpha_k + \beta_k - 2}.$$

At this point one might wonder the necessity of introducing a mixture of Bernoullis. Unlike the mixture of Gaussians, that of Bernoullis seems less convincing. Let θ denotes the weighted average of base models:

$$\theta = \sum_k \pi_k \theta_k,$$

then the variance of the mixture model remains $\theta - \theta^2$. There is no need of using a mixture of Bernoullis (regardin prediction) unless we have to explicitly model a scenario in which there has to be a mixture structure. For example, if we were told that a binary string is generated from a set of unbalanced coins where each coin has different dynamics and we are asked to tell which coin generates some specific toss. But even this scenario might lead to abnormality, considering a coin that always yields head and another that always yields tail.

11.4 EM for mixture of Student distributions

The log-likelihood for the complete data set for the mixture of Student distribution is:

$$l_c(\mathbf{x}, z_k = 1) = \pi_k \cdot \frac{\Gamma(\frac{D}{2} + \frac{v_k}{2})}{\Gamma(\frac{v_k}{2})} \cdot \frac{|\Sigma_k|^{-\frac{1}{2}}}{v_k^{\frac{D}{2}} \pi^{\frac{D}{2}}} \times \left(1 + \frac{1}{v_k} (\mathbf{x} - \mu) \Sigma_k^{-1} (\mathbf{x} - \mu) \right)^{-\frac{v_k + D}{2}}.$$

Following the similar path as exercise 11.2 and exercise 11.3, we have:

$$Q(\theta, \theta^{\text{old}}) = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \cdot \left(\log \pi_k + \log \left(\frac{\Gamma(\frac{D}{2} + \frac{v_k}{2})}{\Gamma(\frac{v_k}{2})} \cdot \frac{|\Sigma_k|^{-\frac{1}{2}}}{v_k^{\frac{D}{2}} \pi^{\frac{D}{2}}} \times \left(1 + \frac{1}{v_k} (\mathbf{x} - \mu) \Sigma_k^{-1} (\mathbf{x} - \mu) \right)^{-\frac{v_k + D}{2}} \right) \right).$$

The rest is ordinary EM, whose detail symbolic forms is too complex to write down explicitly.

11.5 Gradient descent for fitting GMM

From the given (11.118) and (11.119):

$$\begin{aligned} p(\mathbf{x}|\theta) &= \sum_k \pi_k \cdot \mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k) \\ l(\theta) &= \sum_{n=1}^N \log p(\mathbf{x}_n|\theta). \end{aligned}$$

While $r_{nk} = p(z_{nk} = 1|\mathbf{x}_n, \theta)$ is defined by (11.120).

For question (a), recall that:

$$\begin{aligned} l(\theta) &= \sum_{n=1}^N \log \left[\sum_{\mathbf{z}_n} p(\mathbf{x}_n, \mathbf{z}_n|\theta) \right] \\ &= \sum_{n=1}^N \log \left[\sum_{k=1}^K p(\mathbf{x}_n, z_{nk} = 1|\theta) \right] \\ &= \sum_{n=1}^N \log \left[\sum_{k=1}^K \pi_k \cdot \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k) \right]. \end{aligned}$$

We are now ready to taking partial gradient of l w.r.t. μ_k , which yields:

$$\begin{aligned} \frac{\partial l}{\partial \mu_k} &= \sum_{n=1}^N \frac{\partial}{\partial \mu_k} \log \left[\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n|\mu_{k'}, \Sigma_{k'}) \right] \\ &= \sum_{n=1}^N \frac{\pi_k}{\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n|\mu_{k'}, \Sigma_{k'})} \cdot \frac{\partial \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k)}{\partial \mu_k} \\ &= \sum_{n=1}^N \frac{\pi_k \cdot \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k)}{\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n|\mu_{k'}, \Sigma_{k'})} \cdot \Sigma_k^{-1}(\mathbf{x}_n - \mu_k), \end{aligned}$$

using (4.10) for the last step. Now we have arrived in (11.121).

For question (b):

$$\begin{aligned} \frac{\partial l}{\partial \pi_k} &= \sum_{n=1}^N \frac{\partial}{\partial \pi_k} \log \left[\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n|\mu_{k'}, \Sigma_{k'}) \right] \\ &= \sum_{n=1}^N \frac{\mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k)}{\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n|\mu_{k'}, \Sigma_{k'})}, \end{aligned}$$

For question (c), with:

$$\pi_k = \frac{\exp(w_k)}{\sum_{k'=1}^K \exp(w_{k'})},$$

we have:

$$\begin{aligned} \frac{\partial l}{\partial w_k} &= \sum_j \frac{\partial l}{\partial \pi_j} \frac{\partial \pi_j}{\partial w_k} \\ &= \sum_j \sum_{n=1}^N \frac{\mathcal{N}(\mathbf{x}_n | \mu_j, \Sigma_j)}{\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n | \mu_{k'}, \Sigma_{k'})} \cdot \pi_j (1 - \pi_j)^{\mathbb{I}(j=k)} (-\pi_k)^{\mathbb{I}(j \neq k)} \\ &= \sum_{n=1}^N \sum_j r_{nj} (1 - \pi_j)^{\mathbb{I}(j=k)} (-\pi_k)^{\mathbb{I}(j \neq k)} \\ &= \sum_{n=1}^N r_{nk} \cdot (1 - \pi_k) + (1 - r_{nk}) \cdot (-\pi_k) \\ &= \sum_{n=1}^N r_{nk} - \pi_k, \end{aligned}$$

where no constant factor is missed.

For question (d), we have:

$$\begin{aligned} \frac{\partial l}{\partial \Sigma_k} &= \sum_{n=1}^N \frac{\partial}{\partial \Sigma_k} \log \left[\sum_{k'=1}^K \pi_{k'} \cdot \mathcal{N}(\mathbf{x}_n | \mu_{k'}, \Sigma_{k'}) \right] \\ &= \sum_{n=1}^N r_{nk} \left(-\frac{1}{2} \right) \left((\text{textbf{x}_n} \mu_k)(\text{textbf{x}_n} \mu_k)^T - \Sigma_k \right), \end{aligned}$$

during which process we need to use (4.10) and the fact:

$$\frac{\partial |\mathbf{A}|}{\partial \mathbf{A}} = |\mathbf{A}| \mathbf{A}^{-1},$$

for a symmetric matrix \mathbf{A} . Thus the optimal Σ_k takes the same form as what has been derived in exercise 11.2.

For question (e), this process is redundant since the MLE for Σ is already a positive definite matrix.

11.6 EM for a finite scale mixture of Gaussians

For question (a), we advance straightforwardly:

$$\begin{aligned} p(J_n = j, K_n = k | x_n, \theta) &= \frac{p(x_n, J_n = j, K_n = k | \theta)}{p(x_n, \theta)} \\ &= \frac{1}{A} p_j \cdot q_k \cdot \mathcal{N}(x_n | \mu_j, \sigma_k^2) \\ &= \frac{p_j \cdot q_k \cdot \mathcal{N}(x_n | \mu_j, \sigma_k^2)}{\sum_{j', k'} p_{j'} \cdot q_{k'} \cdot \mathcal{N}(x_n | \mu_{j'}, \sigma_{k'}^2)}. \end{aligned}$$

For question (b):

$$\begin{aligned} Q(\theta^{\text{new}}, \theta^{\text{old}}) &= \mathbb{E}_{p(J_n, K_n | x_n, \theta^{\text{old}})} \left[\sum_{n=1}^N \log p(x_n, J_n, K_n | x_n, \theta^{\text{new}}) \right] \\ &= \sum_{n=1}^N \mathbb{E}_{p(J_n, K_n | x_n, \theta^{\text{old}})} \left[\sum_{j=1}^m \sum_{k=1}^l z_{jk} \log (p_j^{\text{new}} q_k^{\text{new}} \cdot \mathcal{N}(x_n | \mu_j^{\text{new}}, \sigma_k^{2, \text{new}})) \right] \\ &= \sum_n \sum_{j, k} \mathbb{E}_{p(J_n, K_n | x_n, \theta^{\text{old}})} [z_{jk}] (\log p_j^{\text{new}} + \log q_k^{\text{new}} + \log \mathcal{N}(x_n | \mu_j^{\text{new}}, \sigma_k^{2, \text{new}})). \end{aligned}$$

What left is trivial calculus. Define:

$$r_{njk} = \mathbb{E}_{p(J_n, K_n | x_n, \theta^{\text{old}})} [z_{jk}] = p(J_n = j, K_n = k | x_n, \theta)$$

For question (c), we have:

$$\begin{aligned} \frac{\partial Q}{\partial \mu_{j'}^{\text{new}}} &= \sum_n \sum_{j, k} r_{njk} \cdot \frac{\partial}{\partial \mu_{j'}^{\text{new}}} \log \mathcal{N}(x_n | \mu_j^{\text{new}}, \sigma_j^{2, \text{new}}) \\ &= \sum_n \sum_k r_{nj'k} \cdot \frac{(x_n - \mu_{j'}^{\text{new}})}{\sigma_k^{2, \text{new}}}. \end{aligned}$$

Setting it to zero yields:

$$\mu_{j'}^{\text{new}} = \frac{\sum_n \sum_k \frac{r_{nj'k} \cdot x_n}{\sigma_k^{2, \text{new}}}}{\sum_n \sum_k \frac{r_{nj'k}}{\sigma_k^{2, \text{new}}}}.$$

11.7 Manual calculation of the M step for a GMM

For question (a), we are to optimize:

$$\sum_{n=1}^3 \sum_{k=1}^2 r_{nk} (\log \pi_i + \log \mathcal{N}(x_n | \mu_k, \sigma_k^2)).$$

For question (b), the new optimal assignment for π_1 , π_2 is derived by differentiating the auxiliary function added with a regularizer to ensure $\pi_1 + \pi_2 = 1$:

$$\frac{\partial Q + \lambda(\pi_1 + \pi_2 - 1)}{\partial \pi_1} = \frac{\sum_{n=1}^3 r_{n,1}}{\pi_1} + \lambda = \frac{1.4}{\pi_1} + \lambda,$$

$$\frac{\partial Q + \lambda(\pi_1 + \pi_2 - 1)}{\partial \pi_2} = \frac{\sum_{n=1}^3 r_{n,2}}{\pi_2} + \lambda = \frac{1.6}{\pi_2} + \lambda.$$

Setting both gradients to zero ends up with:

$$\pi_1 = \frac{7}{15},$$

$$\pi_2 = \frac{8}{15},$$

$$\lambda = -3.$$

For question (c), we use the results from exercise 11.2:

$$\mu_1 = \frac{\sum_{n=1}^3 r_{n1} \cdot x_n}{\sum_{n=1}^3 r_{n1}} = \frac{25}{7},$$

$$\mu_2 = \frac{\sum_{n=1}^3 r_{n2} \cdot x_n}{\sum_{n=1}^3 r_{n2}} = \frac{65}{4}.$$

11.8 Moments of a mixture of Gaussians

For question (a), the expectation of a mixture Gaussian distribution is:

$$\begin{aligned} \mathbb{E}(\mathbf{x}) &= \int \mathbf{x} \sum_k \pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k) d\mathbf{x} \\ &= \sum_k \pi_k \left(\int \mathbf{x} \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k) d\mathbf{x} \right) \\ &= \sum_k \pi_k \mu_k. \end{aligned}$$

For question (b), recall that $\text{cov}(\mathbf{x}) = \mathbb{E}(\mathbf{x}\mathbf{x}^T) - \mathbb{E}(\mathbf{x})\mathbb{E}(\mathbf{x})^T$. We have:

$$\begin{aligned} \mathbb{E}(\mathbf{x}\mathbf{x}^T) &= \int \mathbf{x}\mathbf{x}^T \sum_k \pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k) d\mathbf{x} \\ &= \sum_k \pi_k \int \mathbf{x}\mathbf{x}^T \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k) d\mathbf{x}, \end{aligned}$$

in which:

$$\begin{aligned}\int \mathbf{x}\mathbf{x}^T \mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k) d\mathbf{x} &= \mathbb{E}(\mathbf{x}\mathbf{x}^T) \\ &= \text{cov}(\mathbf{x}) + \mathbb{E}(\mathbf{x})\mathbb{E}(\mathbf{x})^T \\ &= \Sigma_k + \mu_k\mu_k^T.\end{aligned}$$

Therefore:

$$\text{cov}(\mathbf{x}) = \sum_k \pi_k (\Sigma_k + \mu_k\mu_k^T) - \mathbb{E}(\mathbf{x})\mathbb{E}(\mathbf{x})^T.$$

11.9 K-means clustering by hand

```

1  import math
2  def distance(a,b):
3      return math.sqrt((a[0]-b[0])**2+(a[1]-b[1])**2)
4  data
5      =[ [1,0], [3,0], [5,0], [7,0], [9,0], [11,0], [13,0], [15,0], [17,0],
6          [1,3], [3,3], [5,3], [7,3], [9,3], [11,3], [13,3], [15,3], [17,3]]
7
8  c1=[9,3]
9  c2=[11,3]
10 membership=[1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1]
11 for i in range(10):
12     cli=0
13     cli=0
14     cli=0
15     cli=0
16     cli=0
17     cli=0
18     for j in range(len(data)):
19         if (distance(data[j],c1)<=distance(data[j],c2)):
20             membership[i]=1
21             cli=cli+1
22             cli=cli+data[j][0]
23             cli=cli+data[j][1]
24         else:
25             membership[i]=2

```

```

24         c2i=c2i+1
25         c2x=c2x+data[j][0]
26         c2y=c2x+data[j][1]
27     c1[0]=c1x/c1i
28     c1[1]=c1y/c1i
29     c2[0]=c2x/c2i
30     c2[1]=c2y/c2i
31     print(c1)
32     print(c2)

```

The result is:

```

1  [5.0, 1.5]
2  [15.5, 0.0]

```

11.10 Deriving the K-means cost function

For every term for a given k , note that:

$$\begin{aligned}
 \sum_{i:z_i=k} \sum_{i':z_{i'}=k} (x_i - x_{i'})^2 &= \sum_i \sum_{i'} [(x_i - \bar{x}) - (x_{i'} - \bar{x})]^2 \\
 &= \sum_i \sum_{i'} (x_i - \bar{x})^2 + (x_{i'} - \bar{x})^2 - 2(x_i - \bar{x})(x_{i'} - \bar{x}) \\
 &= 2 \cdot n_k \cdot \sum_i (x_i - \bar{x})^2 - 2 \left[\sum_i (x_i - \bar{x}) \right] \left[\sum_{i'} (x_{i'} - \bar{x}) \right] \\
 &= 2 \cdot n_k \cdot \sum_i (x_i - \bar{x})^2.
 \end{aligned}$$

This finishes the proof. One can also adopt (11.132)-(11.135) for the desired result.

11.11 Visible mixtures of Gaussians are in exponential family

Encode latent variable as binary code:

$$z_k = \mathbb{I}(x \text{ is generated from the } k\text{-th base distribution}),$$

then

$$p(\mathbf{z}|\theta) = \prod_{k=1}^K \pi_k^{z_k},$$

$$p(x|\mathbf{z}, \theta) = \prod_{k=1}^K \left(\frac{1}{\sqrt{2\pi\sigma_k^2}} \cdot \exp \left\{ -\frac{1}{2\sigma_k^2} (x - \mu_k)^2 \right\} \right)^{z_k},$$

where $\theta = (\pi, \mu, \sigma^2)$.

The logarithm for the joint distribution is:

$$\begin{aligned} \log p(x, \mathbf{z}|\theta) &= \log \prod_{k=1}^K \left(\frac{\pi_k}{\sqrt{2\pi\sigma_k^2}} \cdot \exp \left\{ -\frac{1}{2\sigma_k^2} (x - \mu_k)^2 \right\} \right)^{z_k} \\ &= \sum_{k=1}^K z_k \cdot \left(\log \pi_k - \frac{1}{2} \log 2\pi\sigma_k^2 - \frac{1}{2\sigma_k^2} (x - \mu_k)^2 \right). \end{aligned}$$

To reduce $p(x, \mathbf{z}|\theta)$ into the exponential family, note that $\log p(x, \mathbf{z}|\theta)$ is linearly dependent on \mathbf{z} and x , hence we can rewrite:

$$\phi(x, \mathbf{z}) = (\mathbf{z}^T, x\mathbf{z}^T, x^2\mathbf{z}^T)^T,$$

the parameters for this form are:

$$(\log \pi_{\dots} - \frac{1}{2} \log 2\pi\sigma^2, \mu_{\dots} \odot \sigma_{\dots}^{-2}, -\frac{1}{2\sigma_{\dots}^2})^T$$

in vector form.

For the mixture of MVN, since:

$$\log p(\mathbf{x}, \mathbf{z}|\theta) = \sum_{k=1}^K z_k \left(\log \pi_k - \frac{D}{2} \log 2\pi - \frac{1}{2} \log |\Sigma_k| - \frac{1}{2} (\mathbf{x} - \mu_k)^T \Sigma_k^{-1} (\mathbf{x} - \mu_k) \right),$$

so the distribution is still an exponential family member, with sufficient statistics:

$$\phi(\mathbf{x}, \mathbf{z}) = (\mathbf{z}, \mathbf{x} \odot \mathbf{z}, \mathbf{x} \odot \mathbf{x} \odot \mathbf{z}),$$

rearranged as a vector. (\odot denotes the tensor/outer product.)

11.12 EM for robust linear regression with a Student t likelihood

Assuming that v and σ^2 are fixed, then the likelihood takes the form:

$$p(y|\mathbf{x}, \mathbf{w}, \sigma^2, v) = f(\sigma^2, v) \times \left(1 + \frac{(y - \mathbf{w}^T \mathbf{x})^2}{2\sigma^2} \right)^{-\frac{v+1}{2}}.$$

We only derive the M-step for this exercise with fixed σ^2 and v . The negative logarithm likelihood is:

$$\mathcal{L}(\mathbf{w}) \propto \sum_{i=1}^N \cdot \log \left(1 + \frac{(y_i - \mathbf{w}^T \mathbf{x}_i)^2}{2\sigma^2} \right).$$

Taking gradient:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} \propto \sum_{i=1}^N \frac{(y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i}{2\sigma^2 + (y_i - \mathbf{w}^T \mathbf{x}_i)^2}.$$

The dependency of the denominator on \mathbf{w} makes the optimum analytically intractable. One possible solution is to randomly initialize \mathbf{w} with \mathbf{w}_0 and conduct iteration:

$$\mathbf{w}_{t+1} = \left(\sum_{i=1}^n \frac{\mathbf{x}_i \mathbf{x}_i^T}{2\sigma^2 + (y_i - \mathbf{w}_t^T \mathbf{x}_i)^2} \right)^{-1} \left(\sum_{i=1}^n \frac{y_i \mathbf{x}_i}{2\sigma^2 + (y_i - \mathbf{w}_t^T \mathbf{x}_i)^2} \right).$$

11.13 EM for EB estimation of Gaussian shrinkage model

This is an example of non-mixture latent graphical model. In this case the latent variable is no longer the one-hot type, making it different from the EM forms that we have developed.

Recall that the complete likelihood for Gaussian shrinkage model is:

$$\begin{aligned} p(\theta, \mathcal{D} | \mu, \tau^2, \{\sigma_j^2\}_{j=1}^D) &= p(\theta | \mu, \tau^2) \cdot p(\mathcal{D} | \theta, \{\sigma_j^2\}_{j=1}^D) \\ &= \prod_{j=1}^D \left[\mathcal{N}(\theta_j | \mu, \tau^2) \prod_{i=1}^{N_j} \mathcal{N}(x_{ij} | \theta_j, \sigma_j^2) \right]. \end{aligned}$$

Taking logarithm yields:

$$\begin{aligned}
\log p\left(\theta, \mathcal{D}|\mu, \tau^2, \{\sigma_j^2\}_{j=1}^D\right) &= \sum_{j=1}^D \left[\log \mathcal{N}(\theta_j|\mu, \tau^2) + \sum_{i=1}^{N_j} \log \mathcal{N}(x_{ij}|\theta_j, \sigma_j^2) \right] \\
&= \sum_{j=1}^D \left[-\frac{1}{2} \log 2\pi\tau^2 - \frac{1}{2\tau^2}(\theta_j - \mu)^2 \right] \\
&\quad + \sum_{j=1}^D \sum_{i=1}^{N_j} \left[-\frac{1}{2} \log 2\pi\sigma_j^2 - \frac{1}{2\sigma_j^2}(x_{ij} - \theta_j)^2 \right] \\
&= -\frac{D}{2} \log 2\pi\tau^2 - \frac{\sum_{j=1}^D (\theta_j - \mu)^2}{2\tau^2} + \sum_{j=1}^D \left[-\frac{N_j}{2} \log 2\pi\sigma_j^2 \right] \\
&\quad - \sum_{j=1}^D \sum_{i=1}^{N_j} \frac{(x_{ij} - \theta_j)^2}{2\sigma_j^2}.
\end{aligned}$$

Note that $p(\theta, \mathcal{D}|\mu, \tau^2, \sigma_j^2\text{s})$ is essentially Gaussian, hence the posterior over θ can be analytically written down with (4.125) (though tedious). Hence all terms that dependent on θ in the logarithm of the complete likelihood can be estimated as moments of their posterior. This is possible since all such terms taking the form θ_j or θ_j^2 . This completes the E-step.

For the M-step, this model is not different from others we have developed so far. Taking partial gradient w.r.t. μ and τ^2 and setting them to zero would yield the update rules.

11.14 EM for censored linear regression

The model for censored linear regression (non-Bayesian version) is:

$$\begin{aligned}
\epsilon_i &\leftarrow \mathcal{N}(\epsilon|0, \sigma^2), \\
z_i &= \mathbf{w}^T \mathbf{x}_i + \epsilon_i, \\
y_i &= \min(z_i, c_i).
\end{aligned}$$

The observed variables are y_i , c_i and \mathbf{x}_i and we are to estimate \mathbf{w} and σ^2 . The latent variable in this model is z_i . The complete likelihood is:

$$p(y_i, z_i|c_i, \mathbf{x}_i, \mathbf{w}, \sigma^2) = \mathcal{N}(y_i|\mathbf{w}^T \mathbf{x}_i, \sigma^2),$$

if $y_i \neq c_i$, and is:

$$p(y_i, z_i | c_i, \mathbf{x}_i, \mathbf{w}, \sigma^2) = \int_{c_i}^{\infty} \mathcal{N}(z | \mathbf{w}^T \mathbf{x}_i, \sigma^2) dz,$$

if $y_i = c_i$, which implies $z_i \geq c_i$. We observe that the integral is going to appear inside the logarithm operator, so it is better to approximate this value from its moments. One possible approximation is to use (11.137) and (11.138), so when $y_i = c_i$, the first and the second moment of z_i are:

$$\begin{aligned} & \mathbf{w}^T \mathbf{x}_i + \sigma \cdot H\left(\frac{c_i - \mathbf{w}^T \mathbf{x}_i}{\sigma}\right), \\ & (\mathbf{w}^T \mathbf{x}_i)^2 + \sigma^2 + \sigma(c_i + \mathbf{w}^T \mathbf{x}_i) \cdot H\left(\frac{c_i - \mathbf{w}^T \mathbf{x}_i}{\sigma}\right), \end{aligned}$$

respectively. Note that this is a variant version of the E-step.

The log likelihood for the entire dataset now becomes:

$$\begin{aligned} \log p(\mathbf{Y}, \mathbf{Z} | c_i, \mathbf{X}, \mathbf{w}, \sigma^2) &= \sum_{y_i \neq c_i} \log p(y_i, z_i | c_i, \mathbf{x}_i, \mathbf{w}, \sigma^2) + \sum_{y_i = c_i} \log p(y_i, z_i | c_i, \mathbf{x}_i, \mathbf{w}, \sigma^2) \\ &= \sum_{y_i \neq c_i} \left[-\frac{1}{2} \log 2\pi\sigma^2 - \frac{(y_i - \mathbf{w}^T \mathbf{x}_i)^2}{2\sigma^2} \right] \\ &\quad + \sum_{y_i = c_i} \left[-\frac{1}{2} \log 2\pi\sigma_i'^2 - \frac{\left(y_i - \mathbf{w}^T \mathbf{x}_i - \sigma \cdot H\left(\frac{c_i - \mathbf{w}^T \mathbf{x}_i}{\sigma}\right)\right)^2}{2\sigma_i'^2} \right], \end{aligned}$$

where $\sigma_i'^2 = \mathbb{E}[z_i^2 | z_i \geq c_i] - \mathbb{E}[z_i | z_i \geq c_i]^2$.

Finally, in the M-step, we take the partial gradient of the log likelihood w.r.t. \mathbf{w} and σ^2 , set them to zero in order to update them. This step is straightforward given H 's gradient can be automatically solved efficiently.

11.15 Posterior mean and variance of a truncated Gaussian

Denote $A = \frac{c_i - \mu_i}{\sigma}$, for the conditional mean, by linearity:

$$\mathbb{E}[z_i | z_i \geq c_i] = \mu_i + \sigma \cdot \mathbb{E}[\epsilon_i | \epsilon_i \geq A].$$

And we have:

$$\mathbb{E}[\epsilon_i | \epsilon_i \geq A] = \frac{1}{p(\epsilon_i \geq A)} \cdot \int_A^{+\infty} \epsilon_i \cdot \mathcal{N}(\epsilon_i | 0, 1) d\epsilon_i = \frac{\phi(A)}{1 - \Phi(A)} = H(A),$$

where H is defined by (11.139). (Recall the definition of the conditional expectation!) Therefore we have:

$$\mathbb{E}[z_i | z_i \geq c_i] = \mu_i + \sigma \cdot H(A).$$

Now we proceed to calculate the expectation for the squared term:

$$\mathbb{E}[z_i^2 | z_i \geq c_i] = \mu_i^2 + 2\mu_i\sigma\mathbb{E}[\epsilon_i | \epsilon_i \geq A] + \sigma^2\mathbb{E}[\epsilon_i^2 | \epsilon_i \geq A].$$

To evaluate $\mathbb{E}[\epsilon_i^2 | \epsilon_i \geq A]$, we make use of the hint of this exercise:

$$\frac{d}{dw}(w \cdot \mathcal{N}(w|0, 1)) = \mathcal{N}(w|0, 1) - w^2 \cdot \mathcal{N}(w|0, 1).$$

Hence we can solve for the following integral by parts:

$$\int_b^c w^2 \cdot \mathcal{N}(w|0, 1)dw = \Phi(c) - \Phi(b) - c \cdot \mathcal{N}(c|0, 1) + b \cdot \mathcal{N}(b|0, 1).$$

Thence:

$$\mathbb{E}[\epsilon_i^2 | \epsilon_i \geq A] = \frac{1}{p(\epsilon_i \geq A)} \cdot \int_A^{+\infty} w^2 \cdot \mathcal{N}(w|0, 1)dw = \frac{1 - \Phi(A) + A \cdot \phi(A)}{1 - \Phi(A)}.$$

Plug it into the previous formula:

$$\begin{aligned} \mathbb{E}[z_i^2 | z_i \geq c_i] &= \mu_i^2 + 2\mu_i\sigma \cdot H(A) + \sigma^2 \frac{1 - \Phi(A) + A\phi(A)}{1 - \Phi(A)} \\ &= \mu_i^2 + \sigma^2 + H(A)(\sigma c_i + \sigma \mu_i). \end{aligned}$$

12 Latent linear models

PPCA is an elegant bridge between probabilistic machine learning models and classical data processing algorithms.

12.1 M-step for FA

The EM for FA, as *a useful exercise if you want to become proficient at the math*, is presented in detail as follows. As for the mixture of FAs, you can refer to *The EM Algorithm for Mixtures of Factor Analyzers*, Zoubin Ghahramani, Geoffrey E. Hinton, 1996.

We begin with: (centralize \mathbf{X} to cancel μ w.l.o.g):

$$p(\mathbf{z}_n) = \mathcal{N}(\mathbf{z}|0, \mathbf{I}),$$

$$p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \Psi) = \mathcal{N}(\mathbf{x}|\mathbf{W}\mathbf{z}, \Psi),$$

where we have centralized \mathbf{X} to simplify the deduction. Now we apply (4.124) and (4.125) to the two equations before, this ends up with:

$$p(\mathbf{z}_n|\mathbf{x}_n, \mathbf{W}, \Psi) = \mathcal{N}(\mathbf{z}_n|\mathbf{m}, \Sigma),$$

$$\Sigma = (\mathbf{I} + \mathbf{W}^T \Psi^{-1} \mathbf{W})^{-1},$$

$$\mathbf{m} = \Sigma \mathbf{W}^T \Psi^{-1} \mathbf{x}_n.$$

The log-likelihood for the complete data set $\{\mathbf{x}, \mathbf{z}\}$ is:

$$\begin{aligned} \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n|\mathbf{W}, \Psi) &= \sum_{n=1}^N [\log p(\mathbf{z}_n) + \log p(\mathbf{x}_n|\mathbf{z}_n, \mathbf{W}, \Psi)] \\ &= \sum_{n=1}^N \left[-\frac{\mathbf{z}_n^T \mathbf{z}_n}{2} - \frac{N}{2} \log |\Psi| + \sum_{n=1}^N -\frac{1}{2} (\mathbf{x}_n - \mathbf{W}\mathbf{z}_n)^T \Psi^{-1} (\mathbf{x}_n - \mathbf{W}\mathbf{z}_n) \right] \\ &= -\frac{1}{2} \left(\sum_{n=1}^N \mathbf{z}_n^T \mathbf{z}_n \right) - \frac{N}{2} \log |\Psi| \\ &\quad - \frac{1}{2} \left(\sum_{n=1}^N \mathbf{x}_n^T \Psi^{-1} \mathbf{x}_n + \mathbf{z}_n^T \mathbf{W}^T \Psi^{-1} \mathbf{W} \mathbf{z}_n - 2 \mathbf{z}_n^T \mathbf{W}^T \Psi^{-1} \mathbf{x}_n \right). \end{aligned}$$

We are now ready to formulate the auxiliary function, let $\theta = (\mathbf{W}, \Psi)$:

$$\begin{aligned} Q(\theta, \theta^{\text{old}}) &= \mathbb{E}_{p(\mathbf{z}|\mathbf{x}, \theta^{\text{old}})} \left[\sum_{n=1}^N \log p(\mathbf{x}_n, \mathbf{z}_n | \theta) \right] \\ &= -\frac{1}{2} \sum_{n=1}^N \mathbb{E}[\mathbf{z}_n^T \mathbf{z}_n] - \frac{N}{2} \log |\Psi| \\ &\quad - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n^T \Psi^{-1} \mathbf{x}_n + \mathbb{E}[\mathbf{z}_n^T \mathbf{W}^T \Psi^{-1} \mathbf{W} \mathbf{z}_n] - 2\mathbb{E}[\mathbf{z}_n^T] \mathbf{W}^T \Psi^{-1} \mathbf{x}_n), \end{aligned}$$

where the conditional first and second moments are:

$$\begin{aligned} \mathbb{E}[\mathbf{z}_n] &= (\mathbf{I} + \mathbf{W}^T \Psi^{-1} \mathbf{W})^{-1} \mathbf{W}^T \Psi^{-1} \mathbf{x}_n = \mathbf{m}_n, \\ \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^T] &= (\mathbf{I} + \mathbf{W}^T \Psi^{-1} \mathbf{W})^{-1} + \mathbf{m}_n \mathbf{m}_n^T = \Sigma + \mathbf{m}_n \mathbf{m}_n^T. \end{aligned}$$

Therefore:

$$\begin{aligned} Q(\theta, \theta^{\text{old}}) &= -\frac{1}{2} \sum_{n=1}^N \text{tr} [\Sigma + \mathbf{m}_n \mathbf{m}_n^T] - \frac{N}{2} \log |\Psi| \\ &\quad - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n^T \Psi^{-1} \mathbf{x}_n + \text{tr} [\mathbf{W}^T \Psi^{-1} \mathbf{W} (\Sigma + \mathbf{m}_n \mathbf{m}_n^T)] - 2\mathbf{m}_n^T \mathbf{W}^T \Psi^{-1} \mathbf{x}_n). \end{aligned}$$

Finally, let us take partial gradient of the auxiliary function w.r.t. \mathbf{W} and Ψ . We start with \mathbf{W} , note that Σ and \mathbf{m}_n only depends on \mathbf{W}^{old} , hence the first term is a constant for θ :

$$\begin{aligned} \frac{\partial Q}{\partial \mathbf{W}} &= -\frac{1}{2} \sum_{n=1}^N \frac{\partial}{\partial \mathbf{W}} \text{tr} [\mathbf{W}^T \Psi^{-1} \mathbf{W} \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^T]] - 2 \cdot \frac{\partial}{\partial \mathbf{W}} \text{tr} [\mathbf{x}_n \mathbf{m}_n^T \Psi^{-1} \mathbf{W}] \\ &= -\frac{1}{2} \sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^T] \cdot 2\Psi^{-1} \mathbf{W} - 2 \cdot \mathbf{x}_n \mathbf{m}_n^T \Psi^{-1}. \end{aligned}$$

Setting it to zero yields:

$$\mathbf{W} = \left(\sum_{n=1}^N \mathbf{x}_n \mathbf{m}_n^T \right) \left(\sum_{n=1}^N \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^T] \right)^{-1}.$$

Meanwhile, let $\Lambda = \Psi^{-1}$:

$$\frac{\partial Q}{\partial \Lambda} = \frac{N}{2} \Lambda^{-1} - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n \mathbf{x}_n^T + \mathbf{W} \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^T] \mathbf{W}^T - 2\mathbf{x}_n \mathbf{m}_n^T \mathbf{W}^T).$$

Hence:

$$\Psi = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n \mathbf{x}_n^T + \mathbf{W} \mathbb{E}[\mathbf{z}_n \mathbf{z}_n^T] \mathbf{W}^T - 2 \mathbf{x}_n \mathbf{m}_n^T \mathbf{W}^T).$$

This completes the proof.

12.2 MAP estimation for the FA model

Assume that the model is exerted the prior distribution $p(\mathbf{W})$ and $p(\Psi)$. Compare with the question before, the M-step needs to be modified into:

$$\begin{aligned} \frac{\partial}{\partial \mathbf{W}} (Q + \log p(\mathbf{W})) &= 0, \\ \frac{\partial}{\partial \Psi} (Q + \log p(\Psi)) &= 0, \end{aligned}$$

whose results dependent on the concrete form of the prior distributions.

12.3 Heuristic for assessing applicability of PCA

We derive this heuristics from an information theory's perspective. Recall that the differential entropy for a MVN is (with $\sigma_i^2 = \lambda_i^{-1}$):

$$h(\{\lambda_i\}_{i=1}^d) = \frac{1}{2} \sum_{i=1}^d d \log_2(2\pi e) + \log_2(\sigma_i^2).$$

After PCA, the covariance for this MVN model is obtained by replacing the smallest d' variances into $\sigma^2 \rightarrow 0$, hence the difference in entropy is:

$$\Delta h(d') = \frac{1}{2} \sum_{i=1}^{d'} \log_2 \frac{\sigma_i^2}{\sigma^2} = \frac{d'}{2} \log_2 \lambda - \frac{1}{2} \sum_{i=1}^{d'} \log_2 \lambda_i.$$

For two eigen series with the same mean $\bar{\lambda}$, it is plausible to expect that the product of the largest d' values in the series with a larger variance is larger, hence the information loss is smaller, making the PCA better regarding information compression.

12.4 Deriving the second principal component

For this exercise, minimizing J makes \mathbf{v}_1 and \mathbf{v}_2 the first and second principal component for the dataset. By definition, $z_{i,j}$ (where $j = 1, 2$) is the projection of \mathbf{x}_i onto \mathbf{v}_j .

For question (a), $J(\mathbf{v}_1, \mathbf{v}_2)$ is the reconstruction loss measured by l_2 norm, hence we would have:

$$z_{i,2} = \mathbf{v}_2^T \mathbf{x}_i,$$

from the physics of projection. On the other hand, we could use a more mathematical way for deduction, with:

$$J(\mathbf{v}_2, \mathbf{z}_2) = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - z_{n,1} \mathbf{v}_1 - z_{n,2} \mathbf{v}_2)^T (\mathbf{x}_n - z_{n,1} \mathbf{v}_1 - z_{n,2} \mathbf{v}_2),$$

we have:

$$\frac{\partial J}{\partial z_{m,2}} = \frac{1}{N} (2 \cdot z_{m,2} \mathbf{v}_2^T \mathbf{v}_2 - 2 \mathbf{v}_2^T (\mathbf{x}_m - z_{m,1} \mathbf{v}_1)) = 0.$$

Using the fact that $\mathbf{v}_2^T \mathbf{v}_2 = 1$ and $\mathbf{v}_2^T \mathbf{v}_1 = 0$, we arrive at:

$$z_{m,2} = \mathbf{v}_2^T \mathbf{x}_m.$$

For question (b), with:

$$\tilde{J}(\mathbf{v}_2) = -\mathbf{v}_2^T \mathbf{C} \mathbf{v}_2 + \lambda_2 (\mathbf{v}_2^T \mathbf{v}_2) + \lambda_{12} (\mathbf{v}_2^T \mathbf{v}_1 - 0),$$

we adopt straightforward matrix algebra:

$$\frac{\partial \tilde{J}}{\partial \mathbf{v}_2} = -2\mathbf{C} \mathbf{v}_2 + 2\lambda_2 \mathbf{v}_2 + \lambda_{12} \mathbf{v}_1,$$

where we have assumed that \mathbf{C} is symmetric and semi-positive definite. The next step is to decompose \mathbf{v}_2 along the eigenvectors of \mathbf{C} as:

$$\mathbf{v}_2 = \sum_{d=1}^D f_d \cdot \mathbf{u}_d,$$

where d is arranged in the inverse order of the eigenvalues y_d , i.e., $\mathbf{u}_1 = \mathbf{v}_1$, $f_1 = \lambda_1$. Taking partial gradient w.r.t. λ_2 and λ_{12} implies:

$$\begin{aligned} \sum_{d=1}^D f_d^2 &= 1, \\ f_1 &= 0. \end{aligned}$$

Hence, taking the gradient w.r.t. \mathbf{v}_2 as zero is tantamount to write:

$$\lambda_{12}\mathbf{v}_1 = \sum_{d=1}^D 2f_d \cdot (y_d - \lambda_2)\mathbf{u}_d.$$

Recall that the eigenvectors for \mathbf{C} are orthogonal to each other, hence we have:

$$\begin{aligned} f_1 \cdot (y_1 - \lambda_2) &= \frac{\lambda_{12}}{2}, \\ \forall d \neq 1, f_d \cdot (y_d - \lambda_2) &= 0. \end{aligned}$$

These equations tell that $\lambda_{12} = 0$, and λ_2 equals one eigenvalue of \mathbf{C} , y_d , whose corresponding eigenvector \mathbf{u}_d is the optimal value for \mathbf{v}_2 . (We ignore the degenerate case here, but the generalization is straightforward.) For such \mathbf{v}_2 , the value of \tilde{J} is $-y_d^2$, whose minimum is reached when \mathbf{v}_2 is the eigenvector corresponding to the second largest eigenvalue.

12.5 Deriving the residual error for PCA

For question (a), we have:

$$\begin{aligned} \left\| \mathbf{x}_n - \sum_{j=1}^K z_{nj} \mathbf{v}_j \right\|^2 &= \left(\mathbf{x}_n - \sum_{j=1}^K z_{nj} \mathbf{v}_j \right)^T \left(\mathbf{x}_n - \sum_{j=1}^K z_{nj} \mathbf{v}_j \right) \\ &= \mathbf{x}_n^T \mathbf{x}_n - 2\mathbf{x}_n^T \sum_{j=1}^K z_{nj} \mathbf{v}_j + \sum_{j=1}^K \sum_{j'=1}^K z_{nj} z_{nj'} \mathbf{v}_j^T \mathbf{v}_{j'} \\ &= \mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^K z_{nj}^2, \end{aligned}$$

which is tantamount to (12.127).

For question (b), we have:

$$\begin{aligned}
J_K &= \frac{1}{N} \sum_{n=1}^N \left(\mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^K z_{nj}^2 \right) \\
&= \frac{1}{N} \sum_{n=1}^T \mathbf{x}_n^T \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N \sum_{j=1}^K \mathbf{v}_j^T \mathbf{x}_n \mathbf{x}_n^T \mathbf{v}_j \\
&= \frac{1}{N} \sum_{n=1}^T \mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^K \mathbf{v}_j^T \left(\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^T \right) \mathbf{v}_j \\
&= \frac{1}{N} \sum_{n=1}^T \mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^K \mathbf{v}_j^T \mathbf{C} \mathbf{v}_j \\
&= \frac{1}{N} \sum_{n=1}^T \mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^K \lambda_j.
\end{aligned}$$

For question (c), we have:

$$\begin{aligned}
J_K &= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^K \lambda_j \\
&= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n^T \mathbf{x}_n - \left(\sum_{j=1}^d \lambda_j - \sum_{j=K+1}^d \lambda_j \right) \\
&= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n^T \mathbf{x}_n - \sum_{j=1}^d \lambda_j + \sum_{j=K+1}^d \lambda_j \\
&= J_d + \sum_{j=K+1}^d \lambda_j = \sum_{j=K+1}^d \lambda_j.
\end{aligned}$$

12.6 Derivation of Fisher's linear discriminant

This is but straightforward algebra:

$$\begin{aligned}
\frac{\partial J}{\partial \mathbf{w}} &= \frac{\mathbf{S}_B \mathbf{w} \mathbf{w}^T \mathbf{S}_W \mathbf{w} - \mathbf{w}^T \mathbf{S}_B \mathbf{w} \mathbf{S}_W \mathbf{w}}{(\mathbf{w}^T \mathbf{S}_W \mathbf{w})^2} \\
&= \frac{1}{(\mathbf{w}^T \mathbf{S}_W \mathbf{w})^2} \mathbf{w}^T \mathbf{S}_W \mathbf{w} \left(\mathbf{S}_B \mathbf{w} - \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}} \mathbf{S}_W \mathbf{w} \right).
\end{aligned}$$

Note that minimizing the Fisher loss is but one option in maximizing $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$ while minimizing $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$. Modern numerical toolkits offer more choices to be explored.

12.7 PCA via successive deflation

The matrix:

$$\mathbf{I} - \mathbf{v}\mathbf{v}^T$$

is unknown as the *projection matrix* that maps an arbitrary vector into the subspace that is orthogonal to \mathbf{v} , since:

$$\mathbf{v}^T(\mathbf{I} - \mathbf{v}\mathbf{v}^T)\mathbf{x} = \mathbf{v}^T\mathbf{x} - \mathbf{v}^T\mathbf{x}.$$

The only requirement is that \mathbf{v} is a unit vector: $\mathbf{v}^T\mathbf{v} = 1$.

For generalization to the projection matrix to the supplementary of a multi-dimensional space spanned by a set of orthogonal and unit bases $\{\mathbf{v}_1, \dots, \mathbf{v}_M\}$ (otherwise using the Schmidt procedure first), we only have to use:

$$\mathbf{I} - \sum_{m=1}^M \mathbf{v}_m \mathbf{v}_m^T.$$

For question (a), we have:

$$\begin{aligned} \frac{1}{N} \tilde{\mathbf{X}} \tilde{\mathbf{X}}^T &= \frac{1}{N} (\mathbf{I} - \mathbf{v}_1 \mathbf{v}_1^T) \mathbf{X} \mathbf{X}^T (\mathbf{I} - \mathbf{v}_1 \mathbf{v}_1^T)^T \\ &= \frac{1}{N} \mathbf{X} \mathbf{X}^T - \frac{1}{N} \mathbf{v}_1 \mathbf{v}_1^T \mathbf{X} \mathbf{X}^T \mathbf{v}_1 \mathbf{v}_1^T \\ &= \frac{1}{N} \mathbf{X} \mathbf{X}^T - \mathbf{v}_1 \mathbf{v}_1^T \lambda_1 \mathbf{v}_1 \mathbf{v}_1^T \\ &= \frac{1}{N} \mathbf{X} \mathbf{X}^T - \lambda_1 \mathbf{v}_1 \mathbf{v}_1^T. \end{aligned}$$

For question (b), let $\tilde{\mathbf{C}} = \frac{1}{N} \tilde{\mathbf{X}} \tilde{\mathbf{X}}^T$ then its principal eigenvector \mathbf{u} satisfies:

$$\tilde{\mathbf{C}}\mathbf{u} = \lambda\mathbf{u}.$$

To solve this equation, expand \mathbf{u} onto the eigenvectors of \mathbf{C} :

$$\mathbf{u} = \sum_{d=1}^D f_d \cdot \mathbf{v}_d,$$

then:

$$\begin{aligned} \tilde{\mathbf{C}}\mathbf{u} &= \sum_{d=1}^D f_d \cdot (\mathbf{C} - \lambda_1 \mathbf{v}_1 \mathbf{v}_1^T) \mathbf{v}_d \\ &= \sum_{d=2}^D f_d \cdot \lambda_d \cdot \mathbf{v}_d. \end{aligned}$$

Then the eigenequation becomes:

$$\begin{aligned} f_1 \lambda &= 0, \\ \forall d = 2, \dots, D, f_d \lambda_d &= f_d \lambda, \end{aligned}$$

s.t.:

$$\sum_{d=1}^D f_d^2 = 1.$$

To solve for this system, we have:

$$\begin{aligned} f_1 &= 0, \\ \lambda &= \lambda_{d'}, d' \geq 2 \\ f_{d'} &= 1, \\ f_d &= 1, d \neq d'. \end{aligned}$$

Hence the solution is $d' = 2$, this finishes the proof.

For question (c), the procedure is roughly:

1. $[\lambda, \mathbf{u}] = f(\mathbf{C})$;
2. Save λ and \mathbf{u} ;
3. $\mathbf{C}- = \lambda \mathbf{u} \mathbf{u}^T$;
4. Repeat.

12.8 Latent semantic indexing

Practice by yourself.

12.9 Imputation in a FA model

In this context, \mathbf{x}_h and \mathbf{x}_v are complementary components of the data space. Recall that:

$$p(\mathbf{x}_v, \mathbf{x}_h | \theta) = \mathcal{N}(\mathbf{x}_v, \mathbf{x}_h | \mathbf{0}, \Psi + \mathbf{W} \mathbf{W}^T).$$

Now we use (4.69):

$$p(\mathbf{x}_h | \mathbf{x}_v, \theta) = \mathcal{N}(\mu, \Sigma),$$

where:

$$\begin{aligned}\mu &= \Sigma_{hv} \Sigma_{vv}^{-1} \mathbf{x}_v, \\ \Sigma &= \Sigma_{hh} - \Sigma_{hv} \Sigma_{vv}^{-1} \Sigma_{vh},\end{aligned}$$

in which Σ_{hv} corresponding to the submatrix in $\Psi + \mathbf{W}\mathbf{W}^T$.

12.10 Efficiently evaluating the PPCA density

We firstly complete the proof of Theorem 12.2.2, this can be done by plugging (12.61) back into (12.60). Recall that the MLE is where \mathbf{W} satisfies (using $\frac{\partial \mathbf{A}^{-1}}{\partial \mathbf{A}} = -\mathbf{A}^{-1} \mathbf{A}^{-1}$):

$$(\mathbf{I} - \mathbf{S}\mathbf{C}^{-1})\mathbf{C}^{-1}\mathbf{W} = 0.$$

With:

$$\mathbf{W} = \mathbf{Q} \begin{pmatrix} \lambda_1 - \sigma^2 & \cdots \\ \cdots & \cdots \\ \cdots & \lambda_L - \sigma^2 \\ 0 & \cdots \\ \cdots & 0 \end{pmatrix},$$

we have:

$$\mathbf{C}^{-1} = \mathbf{Q} \begin{pmatrix} \lambda_1^{-1} & 0 & \cdots & \cdots & \cdots & 0 \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ 0 & \cdots & \lambda_L^{-1} & 0 & \cdots & 0 \\ 0 & \cdots & 0 & \frac{1}{\sigma^2} & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots & \cdots & \cdots \\ 0 & \cdots & \cdots & \cdots & 0 & \frac{1}{\sigma^2} \end{pmatrix} \mathbf{Q}^T.$$

In which \mathbf{Q} is the singular vectors/eigen vectors for \mathbf{S} . This would reduce $(\mathbf{I} - \mathbf{S}\mathbf{C}^{-1})\mathbf{C}^{-1}\mathbf{W}$ to zero, hence complete the proof. For MLE of the σ^2 , using the trace trick.

Now for $p(\mathbf{x}|\tilde{\mathbf{W}}, \tilde{\sigma}^2)$, it is a normal distribution with zero mean and covariance whose last $L - D$ components are uniformly $\tilde{\sigma}^2$.

12.11 PPCA vs FA

```

1  import math
2  import numpy as np
3  from numpy.linalg import eig
4  import matplotlib.pyplot as plt
5  from mpl_toolkits.mplot3d.axes3d import Axes3D
6  np.random.seed(520)
7  z1=np.random.normal(0,1,size=200)
8  z2=np.random.normal(0,1,size=200)
9  z3=np.random.normal(0,1,size=200)
10 x1=z1
11 x2=z1+0.001*z2
12 x3=10*z3
13 fig=plt.figure()
14 axes3d=Axes3D(fig)
15 axes3d.view_init(elev=20., azim=30)
16 axes3d.scatter(z1,z2,z3,color="cyan")
17 axes3d.scatter(x1,x2,x3,color="red")
18 Z=np.vstack((z1,z2,z3))
19 z1m=np.sum(z1)/200
20 z2m=np.sum(z2)/200
21 z3m=np.sum(z3)/200
22 zm=np.array([z1m,z2m,z3m])
23 for j in range(3):
24     for i in range(200):
25         Z[j][i]=Z[j][i]-zm[j]
26
27 X=np.vstack((x1,x2,x3))
28 x1m=np.sum(x1)/200
29 x2m=np.sum(x2)/200
30 x3m=np.sum(x3)/200
31 xm=np.array([x1m,x2m,x3m])
32 for j in range(3):
33     for i in range(200):
34         X[j][i]=X[j][i]-xm[j]
35 SZ=Z@Z.T/200
36 SX=X@X.T/200
37 vals,vecs=eig(SX)
38 print(vals)

```

The scatters are shown in the following figure: Where the red scatters are

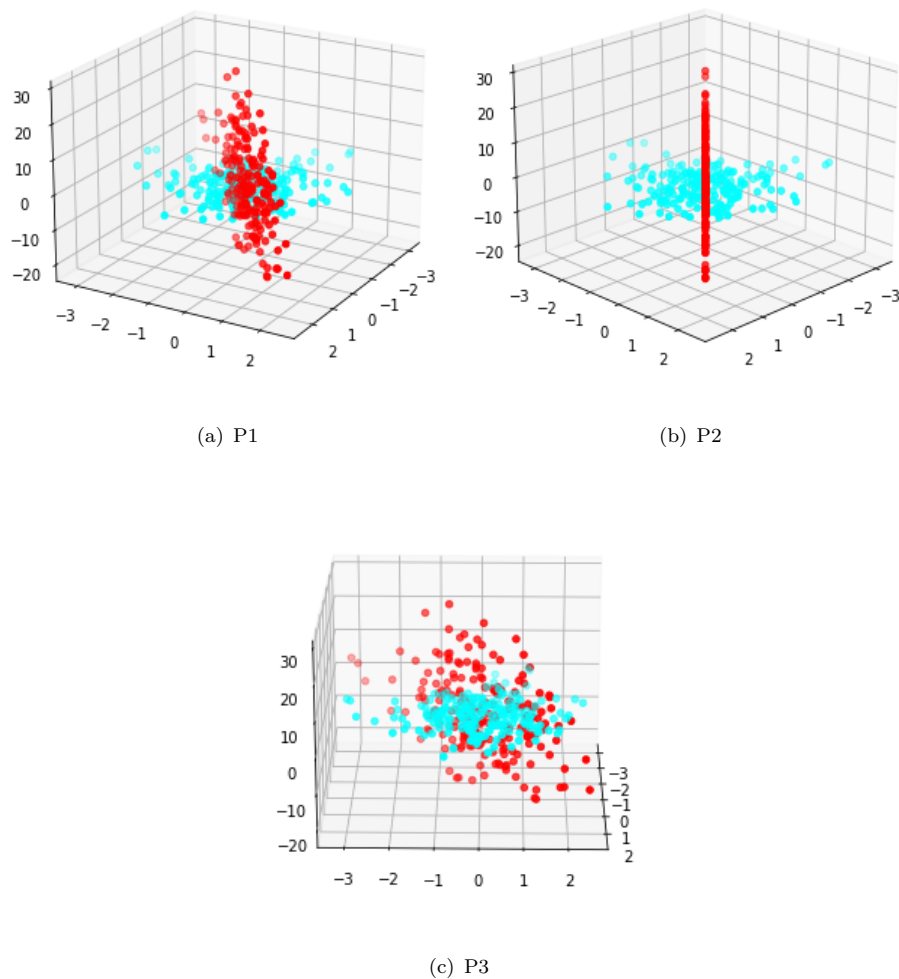


Figure. 7. Exercise 12.11.

\mathbf{X} and the cyan ones are \mathbf{Z} . PCA would select dimension 3 as the principal component since it has the largest variance. PPCA would select the same if σ^2 is small enough. Otherwise, σ^2 would be estimated as approximately $\frac{100}{3}$, hence the reduced variance is larger for the first and the second dimension.

13 Sparse linear models

The techniques covered in this chapter are examples of elementary sparse models, especially linear models. In cases where we are to explore sparsity in modern machine learning models, especially deep models, we are equipped with another set of tools such as dropout, pruning and Bayesian neural networks. When the type of the hyperparameters becomes more complex, e.g., we are to decide the optimal minimal number of layers within a neural network, the sparsity remains a challenge.

13.1 Partial derivative of the RSS

For question (a), define:

$$\text{RSS}(\mathbf{w}) = \sum_{n=1}^N (y_n - \mathbf{w}^T \mathbf{x}_n)^2.$$

Then we have straightforwardly:

$$\begin{aligned} \frac{\partial}{\partial w_j} \text{RSS}(\mathbf{w}) &= \sum_{n=1}^N 2 \cdot (y_n - \mathbf{w}^T \mathbf{x}_n) (-x_{nj}) \\ &= - \sum_{n=1}^N 2 \cdot (x_{nj} y_n - x_{nj} \sum_{i=1}^D w_i x_{ni}) \\ &= - \sum_{n=1}^N 2 (x_{nj} y_n - x_{nj} \sum_{i \neq j}^D w_i x_{ni} - x_{nj}^2 w_j). \end{aligned}$$

From which we observe that w_j 's coefficient is:

$$a_j = 2 \sum_{n=1}^N x_{nj}^2,$$

while the rest irrelevant terms can be absorbed into:

$$c_j = 2 \sum_{n=1}^N x_{nj} (y_n - \mathbf{w}_{-j}^T \mathbf{x}_{n,-j}).$$

The optimal value for w_j is:

$$\hat{w}_j = \frac{c_j}{a_j}.$$

For question (b), (13.184) is obvious by plugging the definition of \mathbf{r}_k into (13.182)-(13.183) and the expression for \hat{w}_j .

13.2 Derivation of M-step for EB for linear regression

We give the EM for Automatic Relevance Determination(ARD) for the linear regression scene, the model is defined by:

$$\begin{aligned} p(\mathbf{Y}|\mathbf{X}, \mathbf{w}, \beta) &= \mathcal{N}(\mathbf{Y}|\mathbf{X}^T \mathbf{w}, \beta^{-1} \mathbf{I}_N), \\ p(\mathbf{w}) &= \mathcal{N}(\mathbf{w}|\mathbf{0}, \mathbf{A}^{-1}), \\ \mathbf{A} &= \begin{pmatrix} \alpha_1 & \cdots & \cdots \\ \cdots & \cdots & \cdots \\ \cdots & \cdots & \alpha_D \end{pmatrix}. \end{aligned}$$

In which the latent variables are \mathbf{w} , and it is the hyperparameters \mathbf{A} that requires estimation.

During the E-step, we are to estimate the expectation of \mathbf{w} , conditioned on \mathbf{X} and \mathbf{Y} . Recall (4.125):

$$p(\mathbf{w}|\mathbf{Y}, \mathbf{X}, \beta, \mathbf{A}) = \mathcal{N}(\mathbf{w}|\mu_E, \Sigma_E),$$

where:

$$\begin{aligned} \Sigma_E &= (\mathbf{A} + \beta \mathbf{X} \mathbf{X}^T)^{-1}, \\ \mu_E &= \beta \Sigma_E \mathbf{X} \mathbf{Y}. \end{aligned}$$

We are now ready to write down the logarithm of the complete likelihood:

$$\begin{aligned} p(\mathbf{Y}, \mathbf{w}|\mathbf{X}, \beta, \mathbf{A}) &= \mathcal{N}(\mathbf{w}|\mathbf{0}, \mathbf{A}) \cdot \mathcal{N}(\mathbf{Y}|\mathbf{X}^T \mathbf{w}, \beta^{-1} \mathbf{I}_N) \\ &\propto |\mathbf{A}|^{\frac{1}{2}} \beta^{\frac{N}{2}} \exp \left\{ -\frac{\beta}{2} (\mathbf{Y} - \mathbf{X}^T \mathbf{w})^T (\mathbf{Y} - \mathbf{X}^T \mathbf{w}) - \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w} \right\}. \end{aligned}$$

Hence the auxiliary function is (let $\theta = (\beta, \mathbf{A})$):

$$\begin{aligned} Q(\theta, \theta^{\text{old}}) &= \mathbb{E} \left[-\frac{1}{2} \log |\mathbf{A}| + \frac{N}{2} \log \beta - \frac{\beta}{2} (\mathbf{Y} - \mathbf{X}^T \mathbf{w})^T (\mathbf{Y} - \mathbf{X}^T \mathbf{w}) - \frac{1}{2} \mathbf{w}^T \mathbf{A} \mathbf{w} \right] \\ &= -\frac{1}{2} \log |\mathbf{A}| + \frac{N}{2} \log \beta - \frac{\beta}{2} \mathbb{E} \left[(\mathbf{Y} - \mathbf{X}^T \mathbf{w})^T (\mathbf{Y} - \mathbf{X}^T \mathbf{w}) \right] - \frac{1}{2} \mathbb{E} [\mathbf{w}^T \mathbf{A} \mathbf{w}]. \end{aligned}$$

The dependence of Q on \mathbf{A} is through:

$$-\frac{1}{2} \log |\mathbf{A}| - \frac{1}{2} \text{tr} (\mathbf{A} \mathbb{E}[\mathbf{w} \mathbf{w}^T]) = -\frac{1}{2} \log |\mathbf{A}| - \frac{1}{2} \text{tr} (\mathbf{A} (\Sigma_E + \mu_E \mu_E^T)).$$

Hence:

$$\frac{\partial Q}{\partial \alpha_j} = \frac{1}{2\alpha_j} - \frac{1}{2} (\Sigma_E + \mu_E \mu_E^T)_{jj}.$$

Using a conjugate onto α_j yields (13.166).

Finally, the dependence of Q on β is through:

$$\frac{N}{2} \log \beta - \frac{\beta}{2} \mathbb{E} \left[(\mathbf{Y} - \mathbf{X}^T \mathbf{w})^T (\mathbf{Y} - \mathbf{X}^T \mathbf{w}) \right],$$

where:

$$\begin{aligned} \mathbb{E} \left[(\mathbf{Y} - \mathbf{X}^T \mathbf{w})^T (\mathbf{Y} - \mathbf{X}^T \mathbf{w}) \right] &= \mathbf{Y}^T \mathbf{Y} - 2\mathbb{E}[\mathbf{w}]^T \mathbf{X} \mathbf{Y} + \text{tr}(\mathbf{X} \mathbf{X}^T \mathbb{E}[\mathbf{w} \mathbf{w}^T]) \\ &= \mathbf{Y}^T \mathbf{Y} - 2\mu_E \mathbf{X} \mathbf{Y} + \text{tr}(\mathbf{X} \mathbf{X}^T [\Sigma_E + \mu_E \mu_E^T]). \end{aligned}$$

Hence we have:

$$\beta = \frac{N}{\mathbf{Y}^T \mathbf{Y} - 2\mu_E \mathbf{X} \mathbf{Y} + \text{tr}(\mathbf{X} \mathbf{X}^T [\Sigma_E + \mu_E \mu_E^T])}.$$

Adding a Gamma prior results in (13.168).

13.3 Derivation of fixed point updates for EB for linear regression

Instead of EM, this method directly optimize the posterior probability, whose negative logarithm is:

$$l(\alpha, \beta) = -\log p(\mathbf{Y}|\mathbf{X}, \alpha, \beta) + \sum_j (a \cdot \log \alpha_j - b \cdot \alpha_j) + c \cdot \log \beta - d \cdot \beta.$$

By (4.126), marginalizing out \mathbf{w} yields:

$$p(\mathbf{Y}|\mathbf{X}, \alpha, \beta) = \mathcal{N}(\mathbf{Y}|\mathbf{0}, \Sigma_{\mathbf{Y}}),$$

where:

$$\Sigma_{\mathbf{Y}} = \beta^{-1} \mathbf{I}_N + \mathbf{X}^T \mathbf{A}^{-1} \mathbf{X}.$$

With:

$$l(\alpha, \beta) = \frac{1}{2} \log |\Sigma_{\mathbf{Y}}| + \frac{1}{2} \mathbf{Y}^T \Sigma_{\mathbf{Y}}^{-1} \mathbf{Y} + \sum_j (a \cdot \log \alpha_j - b \cdot \alpha_j) + c \cdot \log \beta - d \cdot \beta.$$

We have (using the matrix inverse lemma):

$$\begin{aligned} \frac{\partial \mathbf{Y}^T \Sigma_{\mathbf{Y}}^{-1} \mathbf{Y}}{\partial \alpha_j} &= \frac{\partial}{\partial \alpha_j} \beta^{-2} (\mathbf{X} \mathbf{Y})^T \Sigma_E \mathbf{X} \mathbf{Y} \\ &= \frac{1}{2} \mu_{E,j} \mu_{E,j}, \end{aligned}$$

on the other hand, to find:

$$\frac{\partial \log |\Sigma_Y|}{\partial \alpha_j},$$

note that:

$$\Sigma_{\mathbf{Y}, m, n} = \sum_{j=1}^D x_{mj} x_{nj} \alpha_j^{-1} + \beta^{-1} \mathbb{I}(m = n).$$

Thus the term $\frac{\partial \log |\Sigma_Y|}{\partial \alpha_j}$ can be boiled down to $\gamma \cdot \alpha_j^{-1}$, hence the update is:

$$\alpha_j = \frac{2\gamma + 2a}{\mu_{E,j}^2 + 2b}.$$

For β , the procedure is similar.

13.4 Marginal likelihood for linear regression

This is straightforward algebra:

$$\begin{aligned} p(\mathcal{D}|\gamma) &= \int \int \mathcal{N}(\mathbf{Y}|\mathbf{X}^T \mathbf{w}, \sigma^2 \mathbf{I}) \mathcal{N}(\mathbf{w}|\mathbf{0}, \sigma^2 \Sigma) p(\sigma^2) d\mathbf{w} d\sigma^2 \\ &\propto \int \int (\sigma^2)^{-\frac{N}{2} - \frac{D}{2} - 1} \exp \left\{ -\frac{1}{2\sigma^2} [\mathbf{w}^T (\mathbf{X}^T \mathbf{X} + \Sigma^{-1}) \mathbf{w} - 2\mathbf{w}^T \mathbf{X} \mathbf{Y} + \mathbf{Y}^T \mathbf{Y}] \right\} d\mathbf{w} d\sigma^2. \end{aligned}$$

We firstly integrate out \mathbf{w} :

$$\int \exp \left\{ -\frac{1}{2\sigma^2} [\mathbf{w}^T (\mathbf{X}^T \mathbf{X} + \Sigma^{-1}) \mathbf{w} - 2\mathbf{w}^T \mathbf{X} \mathbf{Y} + \mathbf{Y}^T \mathbf{Y}] \right\} d\mathbf{w} \propto [\sigma^2 (\mathbf{X}^T \mathbf{X} + \Sigma^{-1})]^{-\frac{1}{2}} \exp \left\{ -\frac{S(\gamma)}{2\sigma^2} \right\}.$$

(It seems that two factors are lost in S given by (13.17). Finally, we integrate out σ^2 :

$$\int (\sigma^2)^{-\frac{N+D}{2}} \exp \left\{ -\frac{S(\gamma)}{2\sigma^2} \right\} d\sigma^2,$$

what left is tedious calculus. The plugging-in of the g -prior is trivial.

13.5 Reducing elastic net to lasso

Expanding the l.h.s. of (13.196) yields:

$$\begin{aligned} J_1(c\mathbf{w}) &= (\mathbf{y} - c\mathbf{X}\mathbf{w})^T (\mathbf{y} - c\mathbf{X}\mathbf{w}) + c^2 \lambda_2 \mathbf{w}^T \mathbf{w} + \lambda_1 |\mathbf{w}|_1 \\ &= \mathbf{y}^T \mathbf{y} + c^2 \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + c^2 \lambda_2 \mathbf{w}^T \mathbf{w} + \lambda_1 |\mathbf{w}|_1. \end{aligned}$$

While for its r.h.s.:

$$\begin{aligned}
 J_2(\mathbf{w}) &= \begin{pmatrix} \mathbf{y} - c\mathbf{X}\mathbf{w} \\ -c\sqrt{\lambda_2}\mathbf{w} \end{pmatrix}^T \begin{pmatrix} \mathbf{y} - c\mathbf{X}\mathbf{w} \\ -c\sqrt{\lambda_2}\mathbf{w} \end{pmatrix} + c\lambda_1|\mathbf{w}|_1 \\
 &= (\mathbf{y} - c\mathbf{X}\mathbf{w})^T(\mathbf{y} - c\mathbf{X}\mathbf{w}) + c^2\lambda_2\mathbf{w}^T\mathbf{w} + c\lambda_1|\mathbf{w}|_1 \\
 &= \mathbf{y}^T\mathbf{y} + c^2\mathbf{w}^T\mathbf{X}^T\mathbf{X}\mathbf{w} - 2\mathbf{y}^T\mathbf{X}\mathbf{w} + c^2\lambda_2\mathbf{w}^T\mathbf{w} + c\lambda_1|\mathbf{w}|_1.
 \end{aligned}$$

Hence (13.192) and (13.193) are equal, this proceeds to prove (13.195).

This shows elastic net regularization, which pick a regularizing term as a linear combination of l_1 and l_2 equals a lasso one. (The design matrix used in this exercise collects data as rows.)

13.6 Shrinkage in linear regression

For the ordinary least square, the loss is defined as:

$$\text{RSS}(\mathbf{w}) = (\mathbf{Y} - \mathbf{X}^T\mathbf{w})^T(\mathbf{Y} - \mathbf{X}^T\mathbf{w}).$$

Since $\mathbf{X}\mathbf{X}^T = \mathbf{I}$:

$$\text{RSS}(\mathbf{w}) = \mathbf{Y}^T\mathbf{Y} + \mathbf{w}^T\mathbf{w} - 2\mathbf{Y}^T\mathbf{X}^T\mathbf{w}.$$

Take its derivative w.r.t. w_k , where we note that all D weights has been decoupled in the RSS:

$$\frac{\partial}{\partial w_k} \text{RSS}(\mathbf{w}) = 2w_k - 2 \sum_{n=1}^N y_n x_{nk}.$$

Therefore we ends up with:

$$\hat{w}_k^{\text{OLS}} = \sum_{n=1}^N y_n x_{nk} = \frac{c_k}{2}.$$

For the ridge regression:

$$\text{RSS}(\mathbf{w}) = (\mathbf{Y} - \mathbf{X}^T\mathbf{w})^T(\mathbf{Y} - \mathbf{X}^T\mathbf{w}) + \lambda_2\mathbf{w}^T\mathbf{w}.$$

Take its derivative and set it to zero:

$$(2 + 2\lambda_2)w_k = 2 \sum_{n=1}^N y_n x_{nk}.$$

Thus

$$\hat{w}_k^{\text{ridge}} = \frac{\sum_{n=1}^N y_n x_{nk}}{1 + \lambda_2} = \frac{c_k}{2(1 + \lambda_2)}.$$

Finally, recall that

$$\hat{w}_k^{\text{lasso}} = \text{sign}(\hat{w}_k^{\text{OLS}})(|\hat{w}_k^{\text{OLS}}| - \frac{\lambda_1}{2})_+.$$

Observe Figure 13.24, it is easy to address the black line as OLS, the gray one Ridge and the dotted one lasso. Obviously $\lambda_1 = \lambda_2 = 1$. It is noticeable that ridge cause a shrinkage to horizontal axis while lasso cause a sharp shrinkage to zero under certain threshold.

13.7 Prior for the Bernoulli rate parameter in the spike and slab model

Recall that the prior takes the following decomposition:

$$p(\gamma|\alpha_1, \alpha_2) = \prod_{d=1}^D p(\gamma_d|\alpha_1, \alpha_2) = \prod_{d=1}^D p(\gamma_d|\pi_d) \cdot \text{Beta}(\pi_d|\alpha_1, \alpha_2).$$

We now integrate out π_d for the d -th parameter:

$$\begin{aligned} p(\gamma_d|\alpha_1, \alpha_2) &= \frac{1}{B(\alpha_1, \alpha_2)} \int \pi_d^{\gamma_d} (1 - \pi_d)^{(1-\gamma_d)} \pi_d^{\alpha_1-1} (1 - \pi_d)^{\alpha_2-1} d\pi_d \\ &= \frac{1}{B(\alpha_1, \alpha_2)} \int \pi_d^{\alpha_1+\gamma_d-1} (1 - \pi_d)^{\alpha_2+1-\gamma_d-1} d\pi_d \\ &= \frac{B(\alpha_1 + \gamma_d, \alpha_2 + 1 - \gamma_d)}{B(\alpha_1, \alpha_2)} = \frac{\Gamma(\alpha_1 + \alpha_2)}{\Gamma(\alpha_1)\Gamma(\alpha_2)} \cdot \frac{\Gamma(\alpha_1 + \gamma_d)\Gamma(\alpha_2 + 1 - \gamma_d)}{\Gamma(\alpha_1 + \alpha_2 + 1)}. \end{aligned}$$

Therefore(N_1 marks the number of activa parameters in γ):

$$\begin{aligned} p(\gamma|\alpha_1, \alpha_2) &= \binom{N}{N_1} \cdot \frac{\Gamma(\alpha_1 + \alpha_2)^N}{\Gamma(\alpha_1)^N \Gamma(\alpha_2)^N} \cdot \frac{\Gamma(\alpha_1 + 1)^{N_1} \Gamma(\alpha_2 + 1)^{N-N_1}}{\Gamma(\alpha_1 + \alpha_2 + 1)^N} \\ &= \binom{N}{N_1} \cdot \frac{\alpha_1^{N_1} \alpha_2^{N-N_1}}{(\alpha_1 + \alpha_2)^N}. \end{aligned}$$

Hence $p(\gamma)$ in this case is a binomial distribution.

13.8 Deriving E step for GSM prior

The hints contains typos and are unnecessary. By definition, we are to calculate:

$$\begin{aligned}\mathbb{E}\left[\frac{1}{\tau_j^2}|w_j\right] &= \int \frac{1}{\tau_j^2} p(\tau_j^2|w_j) d\tau_j^2 = \int \frac{1}{\tau_j^2} \frac{p(w_j|\tau_j^2)p(\tau_j^2)}{p(w_j)} d\tau_j^2 \\ &= \frac{1}{p(w_j)} \int \frac{1}{\tau_j^2} \mathcal{N}(w_j|0, \tau_j^2) p(\tau_j^2) d\tau_j^2 \\ &= \int \frac{1}{\sqrt{2\pi}} (\tau_j^2)^{-1.5} \exp\left(-\frac{w_j^2}{2\tau_j^2}\right) p(\tau_j^2) d\tau_j^2.\end{aligned}$$

While we also have:

$$p'(w_j) = \int \frac{1}{\sqrt{2\pi}} (\tau_j^2)^{-1.5} \exp\left(-\frac{w_j^2}{2\tau_j^2}\right) p(\tau_j^2) \cdot w_j d\tau_j^2.$$

Plugging both side into (13.197) finishes the proof.

13.9 EM for sparse probit regression with Laplace prior

The ordinary Probit regression involves no latent variable. Introducing Laplace prior for the linear weight \mathbf{w} results in its lasso version. Since Laplace distribution is a continuous mixture of Gaussian according to (13.86), a latent variable τ^2 with the same dimension as \mathbf{w} is introduced. For each component w_j of \mathbf{w} , there is a corresponding latent variable τ_j^2 to guide its variance. The PGM for this Probit regression looks like:

$$\gamma \rightarrow \tau^2 \rightarrow \mathbf{w} \rightarrow \mathbf{y} \leftarrow \mathbf{X}.$$

The joint distribution is:

$$p(\tau^2, \mathbf{w}, \mathbf{y}|\mathbf{X}, \gamma) = \left(\prod_{d=1}^D p(\tau_d^2|\gamma) p(w_d|\tau_d^2) \right) \cdot \left(\prod_{n=1}^N \Phi(\mathbf{w}^T \mathbf{x}_n)^{y_n} (1 - \Phi(\mathbf{w}^T \mathbf{x}_n))^{1-y_n} \right),$$

where Φ is the c.d.f. for a unit Gaussian. According to (13.86):

$$p(\tau^2|\gamma) = \text{Ga}(\tau_d^2|1, \frac{\gamma^2}{2}),$$

$$p(w_d|\tau_d^2) = N(w_d|0, \tau_d^2).$$

Hence:

$$p(\tau^2, \mathbf{w}, \mathbf{y} | \mathbf{X}, \gamma) \propto \exp \left\{ -\frac{1}{2} \sum_{d=1}^D \left(\gamma^2 \tau_d^2 + \frac{w_d^2}{\tau_d^2} \right) \right\} \cdot \prod_{d=1}^D \frac{1}{\tau_d} \\ \cdot \prod_{n=1}^N \Phi(\mathbf{w}^T \mathbf{x}_n)^{y_n} (1 - \Phi(\mathbf{w}^T \mathbf{x}_n))^{1-y_n}.$$

To build the auxiliary function, we assumed \mathbf{w} as the parameter to be estimated and τ^2 as latent variable, thus:

$$Q(\mathbf{w}, \mathbf{w}^{\text{old}}) = \mathbb{E}_{p(\tau^2 | \mathbf{w}^{\text{old}}, \mathbf{y}, \mathbf{X}, \gamma)} [\log p(\tau^2, \mathbf{w}, \mathbf{y} | \mathbf{X}, \gamma)].$$

We now extract terms involving \mathbf{w} from $\log p(\tau^2, \mathbf{w}, \mathbf{y} | \mathbf{X}, \gamma)$:

$$Q(\mathbf{w}, \mathbf{w}^{\text{old}}) = c - \frac{1}{2} \sum_{d=1}^D \frac{w_d^2}{\tau_d^2} + \sum_{n=1}^N y_n \log \Phi(\mathbf{w}^T \mathbf{x}_n) + (1 - y_n) \log (1 - \Phi(\mathbf{w}^T \mathbf{x}_n)).$$

Thus we only need to calculate the conditional expectation:

$$\mathbb{E} \left[\frac{1}{\tau_d^2} | \mathbf{w}^{\text{old}} \right]$$

for the E-step. Whose result is already given in exercise 13.8. The M-step is the same as Gaussian-prior Probit regression and hence is omitted.

13.10 GSM representation of group lasso

Note that the prior over \mathbf{w}_g takes the form:

$$u_g^{\frac{1}{2}} (\rho u_g)^{-\frac{1}{2}} \exp(-\rho u_g).$$

Therefore its logarithm takes the form:

$$c - \|\mathbf{w}_g\|_2,$$

where c is a term that is independent from \mathbf{w}_g and the data. Combining this term with the ordinary data-dependent NLL yields the loss function of the group lasso.

13.11 Projected gradient descent for l1 regularized least squares

Generally, we take the gradient of \mathbf{w} and optimize. When some constraint on \mathbf{w} is broken by the gradient descent, the increment is moderated so that the constraint remains valid. For the following loss function:

$$\min_{\mathbf{w}} \{ \text{NLL}(\mathbf{w}) + \lambda \|\mathbf{w}\|_1 \},$$

consider under a linear regression context:

$$\text{NLL}(\mathbf{w}) = \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$

For $\lambda \|\mathbf{w}\|_1$ which is not differentiable, it is suggest:

$$\mathbf{w} = \mathbf{u} - \mathbf{v},$$

where

$$\begin{aligned} u_i &= (x_i)_+ = \max \{0, x_i\}, \\ v_i &= (-x_i)_+ = \max \{0, -x_i\}. \end{aligned}$$

With $\mathbf{u} \geq \mathbf{0}, \mathbf{v} \geq \mathbf{0}$, we have:

$$\|\mathbf{w}\|_1 = \mathbf{1}_n^T \mathbf{u} + \mathbf{1}_n^T \mathbf{v}.$$

Hence the original problem is translated into:

$$\begin{aligned} \min_{\mathbf{w}} \left\{ \frac{1}{2} \|\mathbf{y} - \mathbf{X}(\mathbf{u} - \mathbf{v})\|_2^2 + \lambda \mathbf{1}_n^T \mathbf{u} + \lambda \mathbf{1}_n^T \mathbf{v} \right\} \\ \text{s.t. } \mathbf{u} \geq \mathbf{0}, \mathbf{v} \geq \mathbf{0}. \end{aligned}$$

Denote:

$$\mathbf{z} = \begin{pmatrix} \mathbf{u} \\ \mathbf{v} \end{pmatrix},$$

then we can rewrite the original target to be optimized into:

$$\begin{aligned} \min_{\mathbf{z}} \left\{ f(\mathbf{z}) = \mathbf{c}^T \mathbf{z} + \frac{1}{2} \mathbf{z}^T \mathbf{A} \mathbf{z} \right\}, \\ \text{s.t. } \mathbf{z} \geq \mathbf{0}, \end{aligned}$$

where:

$$\mathbf{c} = \begin{pmatrix} \lambda \mathbf{1}_n - \mathbf{yX} \\ \lambda \mathbf{1}_n + \mathbf{yX} \end{pmatrix},$$

$$\mathbf{A} = \begin{pmatrix} \mathbf{X}^T \mathbf{X} & -\mathbf{X}^T \mathbf{X} \\ -\mathbf{X}^T \mathbf{X} & \mathbf{X}^T \mathbf{X} \end{pmatrix}.$$

Now we have changed the problem to a quadratic problem with a simple bound constraint. The gradient is given by:

$$\nabla f(\mathbf{z}) = \mathbf{c} + \mathbf{Az}.$$

An ordinary gradient descent step is:

$$\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha \nabla f(\mathbf{z}^k).$$

For projected case, take \mathbf{g}^k :

$$\mathbf{g}_i^k = \min \{ \mathbf{z}_i^k, \alpha \nabla f(\mathbf{z}^k)_i \}.$$

And:

$$\mathbf{z}^{k+1} = \mathbf{z}^k - \mathbf{g}^k,$$

hence \mathbf{z} is constrained as a legal weight candidate.

The original paper suggest more delicate method to moderate the learning rate, refer to *Gradient Projection for Sparse Reconstruction: Application to Compressed Sensing and Other Inverse Problems*, Mario A.T.Figueiredo.

13.12 Subderivative of the hinge loss function

When $\theta < 1$, we have:

$$\partial f(\theta) = \{-1\}.$$

When $\theta = 1$, we have:

$$\partial f(\theta) = [-1, 0],$$

as for $\theta > 1$:

$$\partial f(\theta) = \{0\}.$$

13.13 Lower bounds to convex functions

Since f is a convex function, then any hyperplane that is tangent to an arbitrary point $\mathbf{x}' \in \text{dom}(f)$ has the entire f above it. Denote the normal of the hyperplane at \mathbf{x}' by $\mathbf{n}'_{\mathbf{x}}$, then the hyperplane intersect with \mathbf{x} at $(\mathbf{x}, \phi'_{\mathbf{x}})$ with:

$$(\mathbf{x}' - \mathbf{x}, f(\mathbf{x}') - \phi'_{\mathbf{x}}) \cdot \mathbf{n}'_{\mathbf{x}} = 0.$$

And we have:

$$f(\mathbf{x}) \geq \phi'_{\mathbf{x}}.$$

This assertion holds for arbitrary \mathbf{x}' , hence:

$$f(\mathbf{x}) \geq \sup_{\mathbf{x}'} (\phi'_{\mathbf{x}}).$$

For more details, refer to *Rigorous Affine Lower Bound Functions for Multivariate Polynomials and Their Use in Global Optimisation*.

14 Kernels

The textbook *Learning with Kernels*, Scholkopf and Smola is perhaps the most comprehensive introduction to kernel series models. This branch of machine learning is more involved with functional analysis and operator algebra than with probability, for which reason is a little bitter for probabilistic machine learning students. But kernel machines remain an indispensable tool even when deep models prevail. Therefore one is still encouraged to have an intensive tour into the kernel school.

14.1 Fitting an SVM classifier by hand

The mapped samples are $[1, 0, 0]^T$ and $[1, 2, 2]^T$.

For question (a), the best decision plane is the one with normal $[0, 1, 1]^T$ and passes the point $[1, 1, 1]^T$. Hence \mathbf{w} can be chosen as $[0, 1, 1]^T$.

For question (b), the marginal is the distance between $[1, 1, 1]^T$ and $[1, 2, 2]^T$, hence is $\sqrt{2}$.

For question (c), we have:

$$\mathbf{w} = [0, \frac{1}{2}, \frac{1}{2}]^T.$$

For question (d), we have:

$$\frac{1}{2} \cdot 0 + \frac{1}{2} \cdot 0 + w_0 = -1,$$

$$\frac{1}{2} \cdot 2 + \frac{1}{2} \cdot 2 + w_0 = 1,$$

hence $w_0 = -1$.

For question (e), the discriminant function is:

$$f(x) = -1 + \frac{1}{2}x_2 + \frac{1}{2}x_3.$$

14.2 Linear separability

This is generally false. Since it is possible that the margin is increased by softening the threshold for some samples.

15 Gaussian processes

Refer to *Gaussian Process for Machine Learning*.

15.1 Reproducing property

We have known that the inner product of feature vectors can yield a semi-positive definite kernel matrix. This exercise introduces one method of deducing feature vectors from a legal kernel matrix. Denote $\kappa(\mathbf{x}_1, \mathbf{x})$ by $f(\mathbf{x})$ and $\kappa(\mathbf{x}_2, \mathbf{x})$ by $g(\mathbf{x})$. From the definition:

$$f(\mathbf{x}) = \sum_{i=1}^{\infty} f_i \phi_i(\mathbf{x}),$$

$$\kappa(\mathbf{x}_1, \mathbf{x}) = \sum_{i=1}^{\infty} \lambda_i \phi_i(\mathbf{x}_1) \phi_i(\mathbf{x}).$$

Since \mathbf{x} can be chosen arbitrarily, we have the properties hold (the one for g is obtained similarly):

$$f_i = \lambda_i \phi_i(\mathbf{x}_1),$$

$$g_i = \lambda_i \phi_i(\mathbf{x}_2).$$

Therefore:

$$\begin{aligned} \langle \kappa(\mathbf{x}_1, \cdot), \kappa(\mathbf{x}_2, \cdot) \rangle &= \langle f, g \rangle \\ &= \sum_{i=1}^{\infty} \frac{f_i g_i}{\lambda_i} \\ &= \sum_{i=1}^{\infty} \lambda_i \phi_i(\mathbf{x}_1) \phi_i(\mathbf{x}_2) \\ &= \kappa(\mathbf{x}_1, \mathbf{x}_2). \end{aligned}$$

This completes the proof.

16 Adaptive basis function models

16.1 Nonlinear regression for inverse dynamics

The data provided is in MATLAB form. So practice by yourself.

17 Markov and hidden Markov models

HMM forms the basis of many time-series models. Modern deep models as RNN, LSTM or the Transformer family claim to have better performance. However, they are uniformly combined with ordinary losses that originate from the analysis of the HMM or similar models.

17.1 Derivation of Q function for HMM

Recall that for the auxiliary function, we are to calculate the log-likelihood w.r.t. θ with $\mathbf{z}_{1:T}$ replaced by their expectation w.r.t. θ^{old} :

$$\begin{aligned}
Q(\theta, \theta^{\text{old}}) &= \mathbb{E}_{p(\mathbf{z}_{1:T}|\mathbf{x}_{1:T}, \theta^{\text{old}})} [\log p(\mathbf{z}_{1:T}, \mathbf{x}_{1:T}|\theta)] \\
&= \mathbb{E}_p \left[\log \left\{ \prod_{i=1}^N \left\{ p(z_{i,1}|\pi) \cdot \left(\prod_{t=2}^{T_i} p(z_{i,t}|z_{i,t-1}, \mathbf{A}) \right) \cdot \left(\prod_{t=1}^{T_i} p(x_{i,t}|z_{i,t}, \mathbf{B}) \right) \right\} \right\} \right] \\
&= \mathbb{E}_p \left[\sum_{i=1}^N \sum_{k=1}^K \mathbb{I}[z_{i,1} = k] \cdot \log \pi_k + \sum_{i=1}^N \sum_{t=2}^{T_i} \sum_{j=1}^K \sum_{k=1}^K \mathbb{I}[z_{i,t} = k, z_{i,t-1} = j] \cdot \log \mathbf{A}(j, k) \right. \\
&\quad \left. + \sum_{i=1}^N \sum_{t=1}^{T_i} \sum_{k=1}^K \mathbb{I}[z_{i,t} = k] \cdot \log p(x_{i,t}|z_{i,t} = k, \mathbf{B}) \right].
\end{aligned}$$

From which we obtain (17.98), (17.99) and (17.100), hence (17.97).

17.2 Two filter approach to smoothing in HMMs

For $r_t(i) = p(z_t = i|\mathbf{x}_{t+1:T})$, we have:

$$\begin{aligned}
p(z_t = i|\mathbf{x}_{t+1:T}) &= \sum_j p(z_t = i, z_{t+1} = j|\mathbf{x}_{t+1:T}) \\
&= \sum_j p(z_{t+1} = j|\mathbf{x}_{t+1:T}) \cdot p(z_t = i|z_{t+1} = j, \mathbf{x}_{t+1:T}) \\
&= \sum_j p(z_{t+1} = j|\mathbf{x}_{t+1:T}) \cdot p(z_t = i|z_{t+1} = j) \\
&= \sum_j p(z_{t+1} = j|\mathbf{x}_{t+1:T}) \cdot \Psi^-(j, i),
\end{aligned}$$

where Ψ^- denotes the transform matrix in an inverse sense. In the third step we make use of the conditional independence property within a HMM.

We further have:

$$\begin{aligned}
p(z_{t+1} = j | \mathbf{x}_{t+1:T}) &= p(z_{t+1} = j | \mathbf{x}_{t+1}, \mathbf{x}_{t+2:T}) \\
&\propto p(z_{t+1} = j, x_{t+1}, x_{t+2:T}) \\
&\propto p(z_{t+1} = j | x_{t+2:T}) \cdot p(x_{t+1} | z_{t+1} = j, x_{t+2:T}) \\
&= r_{t+1}(j) \cdot p(x_{t+1} | z_{t+1} = j) \\
&= r_{t+1}(j) \cdot \psi_{t+1}(j).
\end{aligned}$$

Therefore we can calculate $r_t(i)$ recursively:

$$r_t(i) \propto \sum_j r_{t+1}(j) \cdot \psi_{t+1}(j) \cdot \Psi^-(j, i).$$

Finally, it is straightforward to rewrite $\gamma_t(i)$ in terms of new factors:

$$\begin{aligned}
\gamma_t(i) &= p(z_t = i | \mathbf{x}_{1:T}) \\
&\propto p(z_t = i, \mathbf{x}_{1:T}) \\
&= p(z_t = i) \cdot p(\mathbf{x}_{1:T} | z_t = i) \\
&= p(z_t = i) \cdot p(\mathbf{x}_{1:t} | z_t = i) \cdot p(\mathbf{x}_{t+1:T} | z_t = i, \mathbf{x}_{1:t}) \\
&= p(z_t = i) \cdot p(\mathbf{x}_{1:t} | z_t = i) \cdot p(\mathbf{x}_{t+1:T} | z_t = i) \\
&= \frac{1}{p(z_t = i)} \cdot p(\mathbf{x}_{1:t}, z_t = i) \cdot p(\mathbf{x}_{t+1:T}, z_t = i) \\
&\propto \frac{1}{p(z_t = i)} \cdot p(z_t = i | \mathbf{x}_{1:t}) \cdot p(z_t = i | \mathbf{x}_{t+1:T}) \\
&= \frac{\alpha_t(i) \cdot r_t(i)}{\prod_t(i)}.
\end{aligned}$$

This finishes the deduction. Note that in all deduction in HMM we can freely use \propto to drop out the evidence and make use of the graphical conditional independence.

17.3 EM for HMMs with mixture of Gaussian observations

The complete likelihood takes the form:

$$p(\mathbf{z}_{1:T}, \mathbf{x}_{1:T} | \pi, \mathbf{A}, \mathbf{W}, \mu, \Sigma) = \left(\prod_{j=1}^J \pi_j^{\mathbb{I}[z_1=j]} \right) \cdot \left(\prod_{t=2}^T \prod_{i=1}^J \prod_{j=1}^J \mathbf{A}_{i,j}^{\mathbb{I}[z_{t-1}=i, z_t=j]} \right) \cdot \left[\prod_{t=1}^T \prod_{j=1}^J \left(\sum_{k=1}^K w_{jk} \cdot \mathcal{N}(\mathbf{x}_t | \mu_{jk}, \Sigma_{jk}) \right)^{\mathbb{I}[z_t=j]} \right],$$

where π is the initial distribution of the hidden state, \mathbf{A} is a $J * J$ matrix with transition probability, \mathbf{W} is a $J * K$ matrix for w_{jk} , μ and Σ are tensors whose (j, k) -th component denote μ_{jk} and Σ_{jk} respectively. Now its logarithm reads (we temporarily drop the condition on parameters for conciseness):

$$\begin{aligned} \log p(\mathbf{z}_{1:T}, \mathbf{x}_{1:T}) &= \sum_{j=1}^J \mathbb{I}[z_1 = j] \cdot \log \pi_j + \sum_{t=2}^T \sum_{i=1}^J \sum_{j=1}^J \mathbb{I}[z_{t-1} = i, z_t = j] \cdot \log \mathbf{A}_{i,j} \\ &\quad + \sum_{t=1}^T \sum_{j=1}^J \mathbb{I}[z_t = j] \cdot \log \left(\sum_{k=1}^K w_{jk} \cdot \mathcal{N}(\mathbf{x}_t | \mu_{jk}, \Sigma_{jk}) \right). \end{aligned}$$

When being taken expectation w.r.t. $p(\mathbf{z}_{1:T} | \mathbf{x}_{1:T})$, the only terms that matter are identical to those in exercise 17.1. This finishes the E-step for this model.

For the M-step, the update for \mathbf{A} and π are identical to those for an ordinary HMM since their gradients remain the same. To update \mathbf{W}, μ, Σ , note that their dependence on the auxiliary function is through:

$$\sum_{n=1}^N \sum_{t=1}^{T_n} \mathbb{E}[\mathbb{I}[z_{n,t} = j]] \cdot \log \left(\sum_{k=1}^K w_{jk} \cdot \mathcal{N}(\mathbf{x}_t | \mu_{jk}, \Sigma_{jk}) \right).$$

This is tantamount to estimate the parameters for $J * K$ Gaussian components independently, with extra weight on each sample. Let us denote

$$\alpha_{n,t}(j) = \mathbb{E}[\mathbb{I}[z_{n,t} = j]],$$

so $\alpha_{n,t}(j)$ is determined from the old set of parameters, then consider the likelihood, in which we use a one-hot vector of length $J * K$ to embed the

new latent variables:

$$p(\mathbf{x}_{n,t} | \mathbf{h}_{n,t}) = \mathcal{N}(\mathbf{x} | \mu_{jk}, \Sigma_{jk})^{\mathbb{I}[h_{n,t,j,k}=1]}.$$

Although tedious, one can consider (n, t) as the complex index for data, (j, k) as the complex index for Gaussian components. Now the internal auxiliary function reads (the evidence of the latent variables does not depend on the new parameters and is thus omitted):

$$Q^I(\theta, \theta^{\text{old}}) = \sum_{n,t} \sum_{j,k} \mathbb{E}[\mathbb{I}[h_{n,t,j,k} = 1]] \cdot \log \mathcal{N}(\mathbf{x} | \mu_{jk}, \Sigma_{jk}),$$

where:

$$\mathbb{E}[\mathbb{I}[h_{n,t,j,k} = 1]] = p(h_{n,t,j,k} = 1) = p(z_{n,t} = j) \cdot p(h_{n,t,j,k} = 1 | z_{n,t} = j),$$

in which $p(z_{n,t} = j) = \alpha_{n,t}(j)$ and $p(h_{n,t,j,k} = 1 | z_{n,t} = j)$ can be computed as in an ordinary GMM by focusing on the K Gaussian components under the hidden state j . This concludes the E-step for the second auxiliary function. The M-step for the Q^I is thus similar to that for an ordinary, except for the introduction of an extra factor. This also completes the rest M-step for the first auxiliary function.

17.4 EM for HMMs with tied mixtures

In this case on term in the auxiliary function reads:

$$\begin{aligned} \log p(\mathbf{z}_{1:T}, \mathbf{x}_{1:T}) &= \sum_{j=1}^J \mathbb{I}[z_1 = j] \cdot \log \pi_j + \sum_{t=2}^T \sum_{i=1}^J \sum_{j=1}^J \mathbb{I}[z_{t-1} = i, z_t = j] \cdot \log \mathbf{A}_{i,j} \\ &\quad + \sum_{t=1}^T \sum_{j=1}^J \mathbb{I}[z_t = j] \cdot \log \left(\sum_{k=1}^K w_{jk} \cdot \mathcal{N}(\mathbf{x}_t | \mu_k, \Sigma_k) \right). \end{aligned}$$

Following the same method as in exercise 17.4, we build an internal auxiliary function with new latent variables so it reads:

$$Q^I(\theta, \theta^{\text{old}}) = \sum_{n,t} \sum_{j,k} \mathbb{E}[\mathbb{I}[h_{n,t,j,k} = 1]] \cdot \log \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k),$$

what left is the same as in an ordinary GMM model.

18 State space models

18.1 Derivation of EM for LG-SSM

We directly work on the auxiliary function, in which θ denotes the collection of all six matrices. The dependence on \mathbf{u} is omitted for simplicity. The first subscript for variables \mathbf{z} and \mathbf{y} is the index of the sequence. :

$$\begin{aligned}
Q(\theta, \theta^{\text{old}}) &= \mathbb{E}_{p(\mathbf{Z}|\mathbf{Y}, \theta^{\text{old}})} \left[\log \prod_{n=1}^N p(\mathbf{z}_{n,1:T_n}, \mathbf{y}_{n,1:T_n} | \theta) \right] \\
&= \mathbb{E} \left[\log \left(\prod_{n=1}^N p(\mathbf{z}_{n,1} | \theta) \cdot \left(\prod_{i=2}^{T_n} p(\mathbf{z}_{n,i} | \mathbf{z}_{n,i-1}, \theta) \right) \cdot \left(\prod_{i=1}^{T_n} p(\mathbf{y}_{n,i} | \mathbf{z}_{n,i}, \theta) \right) \right) \right] \\
&= \mathbb{E} \left[\sum_{n=1}^N \log \mathcal{N}(\mathbf{z}_{n,1} | \mu_0, \Sigma_0) + \sum_{n=1}^N \sum_{i=2}^{T_n} \log \mathcal{N}(\mathbf{z}_{n,i} | \mathbf{A}_i \mathbf{z}_{n,i-1} + \mathbf{B}_i \mathbf{u}_i, \mathbf{Q}_i) \right. \\
&\quad \left. + \sum_{n=1}^N \sum_{i=1}^{T_n} \log \mathcal{N}(y_{n,i} | \mathbf{C}_i \mathbf{z}_{n,i} + \mathbf{D}_i \mathbf{u}_i, \mathbf{R}_i) \right] \\
&= \mathbb{E} \left[-\frac{N}{2} \cdot \log |\Sigma_0| + \left\{ -\frac{1}{2} \sum_{n=1}^N (\mathbf{z}_{n,1} - \mu_0)^T \Sigma_0^{-1} (\mathbf{z}_{n,1} - \mu_0) \right\} \right. \\
&\quad - \frac{1}{2} \sum_{n=1}^N \sum_{i=2}^{T_n} \log |\mathbf{Q}_i| \\
&\quad + \left\{ -\frac{1}{2} \sum_{n=1}^N \sum_{i=2}^{T_n} (\mathbf{z}_{n,i} - \mathbf{A}_i \mathbf{z}_{n,i-1} - \mathbf{B}_i \mathbf{u}_i)^T \mathbf{Q}_i^{-1} (\mathbf{z}_{n,i} - \mathbf{A}_i \mathbf{z}_{n,i-1} - \mathbf{B}_i \mathbf{u}_i) \right\} \\
&\quad - \frac{1}{2} \sum_{n=1}^N \sum_{i=1}^{T_n} \log |\mathbf{R}_i| \\
&\quad \left. + \left\{ -\frac{1}{2} \sum_{n=1}^N \sum_{i=1}^{T_n} (\mathbf{y}_{n,i} - \mathbf{C}_i \mathbf{z}_{n,i} - \mathbf{D}_i \mathbf{u}_i)^T \mathbf{R}_i^{-1} (\mathbf{y}_{n,i} - \mathbf{C}_i \mathbf{z}_{n,i} - \mathbf{D}_i \mathbf{u}_i) \right\} \right].
\end{aligned}$$

For the E-step, we are to compute:

$$\begin{aligned}
&\mathbb{E}_{p(\mathbf{Z}|\mathbf{Y}, \theta^{\text{old}})} [\mathbf{z}_{n,i}], \\
&\mathbb{E}_{p(\mathbf{Z}|\mathbf{Y}, \theta^{\text{old}})} [\mathbf{z}_{n,i}^T \mathbf{E} \mathbf{z}_{n,i}], \\
&\mathbb{E}_{p(\mathbf{Z}|\mathbf{Y}, \theta^{\text{old}})} [\mathbf{z}_{n,i}^T \mathbf{F} \mathbf{z}_{n,i-1}],
\end{aligned}$$

respectively, in which \mathbf{E} can be Σ_0^{-1} , \mathbf{Q}_i^{-1} , $\mathbf{A}_i^T \mathbf{Q}_i^{-1} \mathbf{A}_i$ and $\mathbf{C}_i^T \mathbf{R}_i^{-1} \mathbf{C}_i$, while \mathbf{F} is $\mathbf{Q}_i^{-1} \mathbf{A}_i$. We then evoke the linearity of the trace operator:

$$\mathbb{E}_{p(\mathbf{Z}|\mathbf{Y}, \theta^{\text{old}})} [\mathbf{z}^T \mathbf{E} \mathbf{z}] = \mathbb{E} [\text{tr}(\mathbf{E} \mathbf{z} \mathbf{z}^T)] = \mathbf{E} \mathbb{E}[\mathbf{z} \mathbf{z}^T].$$

Therefore we only need the conditional expectation of $\mathbf{z}_{n,i}$, $\mathbf{z}_{n,i} \mathbf{z}_{n,i}^T$ and $\mathbf{z}_{n,i} \mathbf{z}_{n,i-1}^T$. The first two variables have been covered by (18.56)-(18.59). For the last variable, we simply need to borrow the form from (18.63). Adding all these together completes the E-step.

For the M-step, we should note that the auxiliary function after the E-step has become the form such that for μ_0 :

$$\frac{\partial Q}{\partial \mu_0} = \frac{\partial}{\partial \mu_0} \left(-\frac{N}{2} \mu_0^T \Sigma_0^{-1} \mu_0 + \left(\sum_{n=1}^N \mathbb{E}[\mathbf{z}_{n,1}] \right)^T \Sigma_0^{-1} \mu_0 \right),$$

whose form is close to that for the mean in an ordinary MVN. For Σ_0 , \mathbf{Q}_i and \mathbf{R}_i , the conclusion is similar. As for matrices such as \mathbf{A}_i , \mathbf{B}_i , \mathbf{C}_i and \mathbf{D}_i , note that $\frac{\partial Q}{\partial \mathbf{A}_i}$ is composed of a series with form:

$$\frac{\partial}{\partial \mathbf{A}_i} \mathbf{z}^T \mathbf{A}_i \mathbf{v} = \mathbf{v} \mathbf{z}^T,$$

and

$$\frac{\partial}{\partial \mathbf{A}_i} \mathbf{z}^T \mathbf{A}_i^T \mathbf{Q}_i^{-1} \mathbf{A}_i \mathbf{z} = 2 \mathbf{z} \mathbf{z}^T \mathbf{A}_i^T \mathbf{Q}_i^{-1}.$$

(assuming symmetry!) Hence we end up with an involving update formula between \mathbf{A}_i and \mathbf{B}_i , so is that for \mathbf{C}_i and \mathbf{D}_i . This finishes the M-step for the EM for LG-SSM.

18.2 Seasonal LG-SSM model in standard form

From Fig.18.6(a), we have:

$$\mathbf{A} = \begin{pmatrix} 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & -1 & -1 & -1 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \end{pmatrix},$$

$$\mathbf{Q} = \begin{pmatrix} \lambda & 0 & 0 & 0 & 0 \\ 0 & \lambda & 0 & 0 & 0 \\ 0 & 0 & Q_c & 0 & 0 \\ 0 & 0 & 0 & \lambda & 0 \\ 0 & 0 & 0 & 0 & \lambda \end{pmatrix},$$

where $\lambda \rightarrow \infty$. The trick here (by using $c_t = -\sum_{s=1}^{S-1} c_{t-s}$ to preserve periodicity) is widely adopted in circuit design and embedded systems. The latent state has to include $S - 1$ states as the memory for the last $S - 1$ states that have appeared.

19 Undirected graphical models(Markov random fields)

19.1 Derivation of the log partition function

Recall the definition of the partition function:

$$Z(\theta) = \sum_{\mathbf{y}} \prod_{c \in \mathcal{C}} \psi_c(\mathbf{y}_c | \theta_c),$$

where \mathcal{C} is the collection of all maximal cliques. The rest is straightforward algebra:

$$\begin{aligned} \frac{\partial \log Z(\theta)}{\partial \theta_{c'}} &= \frac{\partial}{\partial \theta_{c'}} \log \sum_{\mathbf{y}} \prod_{c \in \mathcal{C}} \psi_c(\mathbf{y}_c | \theta_c) \\ &= \frac{1}{Z(\theta)} \sum_{\mathbf{y}} \frac{\partial}{\partial \theta_{c'}} \prod_{c \in \mathcal{C}} \psi_c(\mathbf{y}_c | \theta_c) \\ &= \frac{1}{Z(\theta)} \sum_{\mathbf{y}} \left(\prod_{c \in \mathcal{C}, c \neq c'} \psi_c(\mathbf{y}_c | \theta_c) \right) \cdot \frac{\partial}{\partial \theta_{c'}} \psi_{c'}(\mathbf{y}_{c'} | \theta_{c'}) \\ &= \frac{1}{Z(\theta)} \sum_{\mathbf{y}} \left(\prod_{c \in \mathcal{C}, c \neq c'} \psi_c(\mathbf{y}_c | \theta_c) \right) \cdot \frac{\partial}{\partial \theta_{c'}} \exp \{ \theta_{c'}^T \phi_{c'}(\mathbf{y}_{c'}) \} \\ &= \frac{1}{Z(\theta)} \sum_{\mathbf{y}} \left(\prod_{c \in \mathcal{C}} \psi_c(\mathbf{y}_c | \theta_c) \right) \cdot \phi_{c'}(\mathbf{y}_{c'}) \\ &= \sum_{\mathbf{y}} \phi_{c'}(\mathbf{y}_{c'}) \cdot \left(\frac{1}{Z(\theta)} \prod_{c \in \mathcal{C}} \psi_c(\mathbf{y}_c | \theta_c) \right) \\ &= \sum_{\mathbf{y}} \phi_{c'}(\mathbf{y}_{c'}) \cdot p(\mathbf{y} | \theta) \\ &= \mathbb{E}[\phi_{c'}(\mathbf{y}_{c'}) | \theta]. \end{aligned}$$

19.2 CI properties of Gaussian graphical models


For question (a), we have:

$$\Sigma = \begin{pmatrix} 0.75 & 0.5 & 0.25 \\ 0.5 & 1.0 & 0.5 \\ 0.25 & 0.5 & 0.75 \end{pmatrix},$$


and:

$$\Lambda = \Sigma^{-1} = \begin{pmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{pmatrix}.$$

Thus we have independency: $X_1 \perp X_2 | X_3$. This introduces a MRF like:



For question (b): The inverse of Σ contains no zero element, hence no conditional independency inhabits in this model. Therefore there has to be edge between any two nodes.



This model cancels the marginal independency $X_1 \perp X_3$. But it is possible to model this property by Bayesian network with two directed edges $X_1 \rightarrow X_2$ and $X_3 \rightarrow X_2$. The UGM is obtained by moralizing X_1 and X_3 from the directed version.


For question (c), we only have to consider the terms inside the exponential:

$$-\frac{1}{2} \{x_1^2 + (x_2 - x_1)^2 + (x_3 - x_2)^2\},$$

from which it is easy to see the precision matrix and covariance matrix are:

$$\Lambda = \begin{pmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 1 \end{pmatrix}, \Sigma = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 2 \\ 1 & 2 & 3 \end{pmatrix}.$$

For question (d), the only conditional independency is $X_1 \perp X_3 | X_2$, which is in accordance with the following model:



19.3 Independencies in Gaussian graphical models

For question (a), the model implies that:

- Given X_2 , X_1 is independent from X_3 .

Hence the inverse of its covariance matrix should be zero at $(1, 3)$ -th entry.

Thus A and D are candidates.

For question (b), the candidates are C and D .

For question (c), the PGM tells that:

- X_1 is independent from X_3 .

Hence its covariance should be zero at its $(1, 3)$ -th entry. Therefore C and D could be the covariance matrix.

For question (d), the candidates are A and B .

For question (e), recall (4.68). So A is true while B is not. It is safe to derive the marginal distribution from the joint Gaussian by partitioning the covariance, but not its inverse.

19.4 Cost of training MRFs and CRFs

The answers are generally:

$$\mathcal{O}(r(Nc + 1))$$

for MRFs and:

$$\mathcal{O}(r(Nc + N))$$

for CRFs.

19.5 Full conditional in an Ising model

We derive an expression for (19.126) straightforwardly:

$$\begin{aligned}
p(x_k = 1 | \mathbf{x}_{-k}, \theta) &= \frac{p(x_k = 1, \mathbf{x}_{-k} | \theta)}{p(\mathbf{x}_{-k} | \theta)} \\
&= \frac{p(x_k = 1, \mathbf{x}_{-k} | \theta)}{p(x_k = 0, \mathbf{x}_{-k} | \theta) + p(x_k = 1, \mathbf{x}_{-k} | \theta)} \\
&= \frac{\frac{1}{Z(\theta)} \exp \left\{ \sum_{j \neq k} x_j J_{kj} + h_k + f(\mathbf{x}_{-k}) \right\}}{\frac{1}{Z(\theta)} \exp \{f(\mathbf{x}_{-k})\} + \frac{1}{Z(\theta)} \exp \left\{ \sum_{j \neq k} x_j J_{kj} + h_k + f(\mathbf{x}_{-k}) \right\}} \\
&= \frac{1}{1 + \exp \left\{ - \left(\sum_{j \neq k} x_j J_{kj} + h_k \right) \right\}} \\
&= \sigma \left(h_k + \sum_{j \neq k} x_j J_{kj} \right),
\end{aligned}$$

where $f(\mathbf{x}_{-k})$ are the sum of terms independent from x_k inside the exponential.

As for $p(x_k = 1 | \mathbf{x}_{\text{nb}_k}, \theta)$, the reason is similar, since the summation $\sum_{j \neq k} x_j J_{kj}$ is conducted on x_k 's neighbour.

In case we use the notation $x_i \in \{-1, +1\}$, the deduction becomes:

$$\begin{aligned}
p(x_k = 1 | \mathbf{x}_{-k}, \theta) &= \frac{\frac{1}{Z(\theta)} \exp \left\{ \sum_{j \neq k} x_j J_{kj} + h_k + f(\mathbf{x}_{-k}) \right\}}{\frac{1}{Z(\theta)} \exp \left\{ - \sum_{j \neq k} x_j J_{kj} - h_k + f(\mathbf{x}_{-k}) \right\} + \frac{1}{Z(\theta)} \exp \left\{ \sum_{j \neq k} x_j J_{kj} + h_k + f(\mathbf{x}_{-k}) \right\}} \\
&= \frac{1}{1 + \exp \left\{ -2 \cdot \left(\sum_{j \neq k} x_j J_{kj} + h_k \right) \right\}} \\
&= \sigma \left(2 \cdot \left(h_k + \sum_{j \neq k} x_j J_{kj} \right) \right).
\end{aligned}$$

20 Exact inference for graphical models

20.1 Variable elimination

The figure index is mistaken. The graphical model takes the following structure:




Figure. 8. Exercise 20.1.

For question (a), we use the following order for variable elimination:

$$x_1, x_2, x_3, x_4, x_5, x_6.$$

We detail the entire process as follows:

$$\begin{aligned}
 Z &= \sum_{x_6} \sum_{x_5} \sum_{x_4} \sum_{x_3} \sum_{x_2} \sum_{x_1} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \psi_{24}(x_2, x_4) \psi_{34}(x_3, x_4) \psi_{45}(x_4, x_5) \psi_{56}(x_5, x_6) \\
 &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_4} \psi_{45} \sum_{x_3} \psi_{34} \sum_{x_2} \psi_{24} \sum_{x_1} \psi_{12} \psi_{13} \\
 &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_4} \psi_{45} \sum_{x_3} \psi_{34} \sum_{x_2} \psi_{24} \tau_{23} \\
 &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_4} \psi_{45} \sum_{x_3} \psi_{34} \tau_{34} \\
 &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_4} \psi_{45} \tau_4 \\
 &= \sum_{x_6} \sum_{x_5} \psi_{56} \tau_5 \\
 &= \sum_{x_6} \tau_6.
 \end{aligned}$$

The largest intermediate factor is:

$$\tau_{23} \sum_{x_1} \psi_{12} \psi_{13},$$

hence for (b) the largest clique consists of three nodes.

For question (c), we have:

$$\begin{aligned} Z &= \sum_{x_6} \sum_{x_5} \sum_{x_3} \sum_{x_2} \sum_{x_1} \sum_{x_4} \psi_{12} \psi_{13} \psi_{24} \psi_{34} \psi_{45} \psi_{56} \\ &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_3} \sum_{x_2} \sum_{x_1} \psi_{12} \psi_{13} \sum_{x_4} \psi_{24} \psi_{34} \psi_{45} \\ &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_3} \sum_{x_2} \sum_{x_1} \psi_{12} \psi_{13} \tau_{235} \\ &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_3} \sum_{x_2} \tau'_{235} \\ &= \sum_{x_6} \sum_{x_5} \psi_{56} \sum_{x_3} \tau_{35} \\ &= \sum_{x_6} \sum_{x_5} \psi_{56} \tau_5 \\ &= \sum_{x_6} \tau_6, \end{aligned}$$

where we observe a big clique $\{x_2, x_3, x_4, x_5\}$.

20.2 Gaussian times Gaussian is Gaussian

For this exercise, we have:

$$\begin{aligned} &\mathcal{N}(x|\mu_1, \lambda_1^{-1}) \times \mathcal{N}(x|\mu_2, \lambda_2^{-1}) \\ &= \frac{\sqrt{\lambda_1 \lambda_2}}{2\pi} \cdot \exp \left\{ -\frac{\lambda_1}{2} (x - \mu_1)^2 - \frac{\lambda_2}{2} (x - \mu_2)^2 \right\} \\ &= \frac{\sqrt{\lambda_1 \lambda_2}}{2\pi} \cdot \exp \left\{ -\frac{\lambda_1 + \lambda_2}{2} \cdot x^2 + (\lambda_1 \mu_1 + \lambda_2 \mu_2) \cdot x - \frac{\lambda_1 \mu_1^2 + \lambda_2 \mu_2^2}{2} \right\}. \end{aligned}$$

We now complete the square:

$$\begin{aligned} &\exp \left\{ -\frac{\lambda_1 + \lambda_2}{2} x^2 + (\lambda_1 \mu_1 + \lambda_2 \mu_2) x - \frac{\lambda_1 \mu_1^2 + \lambda_2 \mu_2^2}{2} \right\} \\ &= c \cdot \exp \left\{ -\frac{\lambda}{2} (x - \mu)^2 \right\}, \end{aligned}$$

where:

$$\begin{aligned}\lambda &= \lambda_1 + \lambda_2, \\ \mu &= \lambda^{-1}(\lambda_1\mu_1 + \lambda_2\mu_2).\end{aligned}$$

The constant factor c is obtained by computing the constant terms inside the exponential, which ends up with (20.20).

20.3 Message passing on a tree

The complete likelihood for this model is:

$$p(X_1, X_2, X_3, G_1, G_2, G_3) = p(G_1)p(G_2|G_1)p(G_3|G_1)p(X_1|G_1)p(X_2|G_2)p(X_3|G_3).$$

For question (a), we have the following variable elimination:

$$\begin{aligned}p(X_2 = 50) &= \int_{X_1} \int_{X_3} \sum_{G_1} \sum_{G_2} \sum_{G_3} p(X_1, X_2 = 50, X_3, G_1, G_2) \\ &= \sum_{G_1} p(G_1) \sum_{G_2} p(G_2|G_1)p(X_2 = 50|G_2) \sum_{G_3} p(G_3|G_1) \int_{X_3} p(X_3|G_3) \int_{X_1} p(X_1|G_1) \\ &= \sum_{G_1} p(G_1) \sum_{G_2} p(G_2|G_1)p(X_2 = 50|G_2) \\ &= \end{aligned}$$

On the other hand:

$$\begin{aligned}p(X_2 = 50, G_1 = 1) &= \int_{X_1} \int_{X_3} \sum_{G_2} \sum_{G_3} p(X_1, X_2 = 50, X_3, G_1 = 1, G_2) \\ &= p(G_1 = 1) \sum_{G_2} p(G_2|G_1 = 1)p(X_2 = 50|G_2).\end{aligned}$$

Hence:

$$\begin{aligned}p(G_1 = 1|X_2 = 50) &= \frac{p(G_1 = 1, X_2 = 50)}{p(X_2 = 50)} \\ &= \frac{0.45 + 0.05 \times e^{-5}}{0.5 + 0.5 \times e^{-5}} \\ &\approx 0.9.\end{aligned}$$

For question (b), we have:

$$p(X_2 = 50, X_3 = 50) = \sum_{G_1} p(G_1) \sum_{G_2} p(G_2|G_1)p(X_2 = 50|G_2) \sum_{G_3} p(G_3|G_1)p(X_3 = 50|G_3) \int_{X_1} p(X_1|G_1),$$

similar for $p(G_1 = 1, X_2 = 50, X_3 = 50)$, hence:

$$\begin{aligned} p(G_1 = 1 | X_2 = 50, X_3 = 50) &= \frac{p(G_1 = 1, X_2 = 50, X_3 = 50)}{p(X_2 = 50, X_3 = 50)} \\ &= \frac{0.405 + 0.09 \times e^{-5} + 0.005 \times e^{-10}}{0.41 + 0.18 \times e^{-5} + 0.41 \times e^{-10}} \\ &\approx 0.99. \end{aligned}$$

Since both X_2, X_3 equal fifty implies that both G_2, G_3 are approximately unity, this confirms the belief that $G_1 = 1$.

For question (c), the setting is symmetric to (b), hence:

$$p(G_1 = 0 | X_2 = 60, X_3 = 60) \approx 0.01.$$

For question (d), we follow the analysis as before:

$$p(G_1 = 1 | X_2 = 50, X_3 = 60) = \frac{p(G_1 = 1)p(X_2 = 50, X_3 = 60) | G_1)}{p(X_2 = 50, X_3 = 60)}.$$

Note that:

$$p(X_2 = 50, X_3 = 60) | G_1) = p(X_2 = 50, X_3 = 60) | G_{10})$$

by symmetry, hence:

$$p(G_1 = 1 | X_2 = 50, X_3 = 60) = \frac{1}{2}.$$

Since two contradictive evidence cancell each other.

This is a soft/fuzzy version of the ordinary logical reasoning, where we could have set $p(G_{2,3} | G_1)$ to then identity matrices and the variance for generating X_i to zero.

20.4 Inference in 2D lattice MRFs

For question (a), we consider only two options:

- Eliminate by column, i.e., eliminate m variables each round.
- Eliminate by row, i.e., eliminate n variables each round.

Let $\{X_{i,j}\}_{i=1,j=1}^{m,n}$ denotes the collection of all hidden variables. Note that the local evidence $\{Y_{i,j}\}_{i=1,j=1}^{m,n}$ can be absorbed into the local energy, what

being broken is only the normalization condition. It is easy to see that the bottleneck in computing $p(\mathbf{X}|\mathbf{Y})$ is the evidence:

$$p(\mathbf{Y}) = \int p(\mathbf{X}, \mathbf{Y}) d\mathbf{X}.$$

The order by which we eliminate all hidden variables determines the efficiency of the overall algorithm.

For the first option, we begin by cancelling the last row, $\{X_{i,n}\}_{i=1}^m$, this step involves altogether:

$$3m - 1$$

terms. In which are $(m - 1)$ edges within this column, m edges with the $(n - 1)$ -th column and m local evidence. To perform the integral, let us assume that a sampling-based paradigm is adopted, and each variable in $X_{i,n}$ and $X_{i,n-1}$ is sampled for K times. This elimination ends up with a table with K^m entries. Hence the integral is performed under complexity:

$$K^{2m} \times (3m - 1) = \mathcal{O}(m \cdot K^{2m}).$$

This procedure shall be iterated along the MRF for n times, results in the overall complexity of:

$$\mathcal{O}(m \cdot n \cdot K^{2m}).$$

Symmetrically, that for the second option is:

$$\mathcal{O}(m \cdot n \cdot K^{2n}).$$

For question (b), with $m \ll n$, we should adopt the first method.

For question (c), since $m = n$, either method will do. We can adopt the Viterbi algorithm for HMM in this case, using a table of size 2^{2n} to save the propagated information (as dynamic programming).

21 Variational inference

The *variational* methods refer to a broader family of models and algorithms. Variation is the functional version of differential, the later is a slight change in the domain of one variable, while variation is a slight change in the space of function. Before variational inference in machine learning, the variational methods have been adopted by physicists for centuries.

At any first course for classical mechanics, the variational method should be introduced where the space to be studied is that of all possible trajectories (see, e.g., *Mechanics, Landau*). Minimizing the action w.r.t. the trajectory would yield the actual physical behavior of the system. In quantum mechanics, the variational method is widely used to derive an upper bound for the energy of the base state for a complex particle system, whose wave functions' forms are too complex to be solved analytically (see, e.g., *Modern Quantum Mechanics, J.J.Sakurai*).

In machine learning, the (local) variational inference has a similar motivation as their physics counterpart, especially that for quantum mechanics: *To bound an analytic intractable object with a tractable model.*

Variational inference is arguably one of the most involved and complicated branches for orthodox probabilistic machine learning.

21.1 Laplace approximation to $p(\mu, \log \sigma | \mathcal{D})$ for a univariate Gaussian

Laplace approximation aims at representing $f(\mu, l) = \log p(\mu, l = \log \sigma | \mathcal{D})$ from its first and second-order gradients. We have:

$$\begin{aligned}
 \log p(\mu, l | \mathcal{D}) &= \log p(\mu, l, \mathcal{D}) - \log p(\mathcal{D}) \\
 &= \log p(\mu, l) + \log p(\mathcal{D} | \mu, l) + \text{const} \\
 &= \log p(\mathcal{D} | \mu, l) + \text{const} \\
 &= \sum_{n=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp \left\{ -\frac{(y_n - \mu)^2}{2\sigma^2} \right\} + \text{const} \\
 &= -N \log \sigma + \sum_{n=1}^N -\frac{(y_n - \mu)^2}{2\sigma^2} + \text{const} \\
 &= -N \cdot l - \frac{\exp(-2l)}{2} \sum_{n=1}^N (y_n - \mu)^2 + \text{const}.
 \end{aligned}$$

Thus we can take gradients:

$$\begin{aligned}
 \frac{\partial \log p(\mu, l | \mathcal{D})}{\partial \mu} &= \frac{-\exp(-2l)}{2} \sum_{n=1}^N 2 \cdot (y_n - \mu) \\
 &= -\frac{N}{\sigma^2} \cdot (\bar{y} - \mu), \\
 \frac{\partial \log p(\mu, l | \mathcal{D})}{\partial l} &= -N - \frac{-2 \cdot \exp(-2l)}{2} \sum_{n=1}^N (y_n - \mu)^2 \\
 &= -N + \frac{1}{\sigma^2} \sum_{n=1}^N (y_n - \mu)^2, \\
 \frac{\partial^2 \log p(\mu, l | \mathcal{D})}{\partial \mu^2} &= -\frac{N}{\sigma^2}, \\
 \frac{\partial^2 \log p(\mu, l | \mathcal{D})}{\partial l^2} &= -\frac{2}{\sigma^2} \sum_{n=1}^N (y_n - \mu)^2, \\
 \frac{\partial^2 \log p(\mu, l | \mathcal{D})}{\partial \mu \partial l} &= N \cdot (\bar{y} - \mu) \cdot (-2) \cdot \frac{1}{\sigma^2}.
 \end{aligned}$$

Finally, to conduct the Laplace approximation, recall that:

$$f(\mu, l) \approx \text{const} + \begin{pmatrix} \frac{\partial f}{\partial \mu} & \frac{\partial f}{\partial l} \end{pmatrix} \begin{pmatrix} \mu \\ l \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \mu & l \end{pmatrix} \mathbf{H} \begin{pmatrix} \mu \\ l \end{pmatrix}.$$

With this expansion in mind, we have:

$$p(\mu, l) = \exp(l(\mu, l)) = \text{const} \cdot \exp \left\{ -\frac{1}{2} \begin{pmatrix} \mu - m_1 & l - m_2 \end{pmatrix} \Lambda \begin{pmatrix} \mu - m_1 \\ l - m_2 \end{pmatrix} \right\}.$$

Calibrating the coefficient, we know that the covariance and the mean for this approximation are:

$$\Sigma = - \begin{pmatrix} \frac{\partial^2 \log p(\mu, l|D)}{\partial \mu^2} & \frac{\partial^2 \log p(\mu, l|D)}{\partial l^2} \\ \frac{\partial^2 \log p(\mu, l|D)}{\partial l^2} & \frac{\partial^2 \log p(\mu, l|D)}{\partial \mu \partial l} \end{pmatrix}^{-1},$$

$$m = \begin{pmatrix} \frac{\partial^2 \log p(\mu, l|D)}{\partial \mu^2} & \frac{\partial^2 \log p(\mu, l|D)}{\partial l^2} \\ \frac{\partial^2 \log p(\mu, l|D)}{\partial l^2} & \frac{\partial^2 \log p(\mu, l|D)}{\partial \mu \partial l} \end{pmatrix} \begin{pmatrix} \frac{\partial f}{\partial \mu} \\ \frac{\partial f}{\partial l} \end{pmatrix}.$$

21.2 Laplace approximation to normal-gamma

This is the same with exercise 21.1 when the prior is uninformative. Note that:

$$\begin{aligned} \sum_{n=1}^N (y_n - \mu)^2 &= \sum_{n=1}^N ((y_n - \bar{y}) - (\mu - \bar{y}))^2 \\ &= \sum_{n=1}^N (y_n - \bar{y})^2 + \sum_{n=1}^N (\mu - \bar{y})^2 + 2(\mu - \bar{y}) \cdot \sum_{n=1}^N (y_n - \bar{y}) \\ &= N s^2 + N(\mu - \bar{y})^2, \end{aligned}$$

where $s^2 = \frac{1}{N} \sum_{n=1}^N (y_n - \bar{y})^2$

Conclusions for all problems (a), (b) and (c) have been included in those of the previous exercise.

21.3 Variational lower bound for VB for univariate Gaussian

Recall that the variational lower bound for VB is defined by:

$$\mathcal{L}(q) = \mathbb{KL}(q \| \tilde{p}),$$

where \tilde{p} is the (unnormalized) posterior and q is the variational distribution. For the univariate Gaussian case, we already arrive in (21.84)-(21.87) and

(21.92)-(21.97), thus we only have to fill the gaps (21.88)-(21.91), among which the last three equations have been illustrated in Chapter 2.

Here we derive (21.88). The Gamma distribution is an exponential family distribution:

$$\begin{aligned}\text{Ga}(x|a, b) &= \frac{b^a}{\Gamma(a)} x^{a-1} \exp \{-b \cdot x\} \\ &\propto \exp \{-b \cdot x + (a-1) \ln x\} \\ &= \exp \{\phi(x)^T \theta\},\end{aligned}$$

whose sufficient statistics is $\phi(x) = (x, \ln x)^T$ and natural parameter is $\theta = (-b, a-1)^T$. Its cumulant function is:

$$\begin{aligned}A(\theta) &= \log Z(\theta) \\ &= \log \frac{\Gamma(a)}{b^a} \\ &= \log \Gamma(a) - a \log b.\end{aligned}$$

Recall that the expectation of one sufficient statistics is given by the derivative of the cumulant function, therefore:

$$\mathbb{E}[\ln x] = \frac{\partial A}{\partial (a-1)} = \frac{\Gamma'(a)}{\Gamma(a)} - \log b.$$

Finally, according to the definition $\psi(a) = \frac{\Gamma'(a)}{\Gamma(a)}$, we arrive in:

$$\mathbb{E}[\ln x] = \psi(a) - \log b.$$

By using the property of the exponential family, we can get rid of tedious calculus. The rest steps have already been detailed in 21.5.1.6.

21.4 Variational lower bound for VB for GMMs

For variational GMM, the lower bound is given by:

$$\begin{aligned}
\mathbb{E}_q \left[\log \frac{p(\theta, \mathcal{D})}{q(\theta)} \right] &= \mathbb{E}_q [\log p(\theta, \mathcal{D})] - \mathbb{E}_q [\log q(\theta)] \\
&= \mathbb{E}_q [\log p(\mathcal{D}|\theta)] + \mathbb{E}_q [\log p(\theta)] - \mathbb{E}_q [\log q(\theta)] \\
&= \mathbb{E} [\log p(\mathbf{x}|\mathbf{z}, \mu, \Lambda, \pi)] + \mathbb{E} [\log p(\mathbf{z}, \mu, \Lambda, \pi)] \\
&\quad - \mathbb{E} [\log q(\mathbf{z}, \mu, \Lambda, \pi)] \\
&= \mathbb{E} [\log p(\mathbf{x}|\mathbf{z}, \mu, \Lambda, \pi)] + \mathbb{E} [\log p(\mathbf{z}|\pi)] + \mathbb{E} [\log p(\pi)] + \mathbb{E} [\log p(\mu, \Lambda)] \\
&\quad - \mathbb{E} [\log q(\mathbf{z})] - \mathbb{E} [\log q(\pi)] - \mathbb{E} [\log q(\mu, \Lambda)].
\end{aligned}$$

Which involved Bayes rules and the decomposition form of the posterior and the variational distribution q . We now proceed to prove (21.209) through (21.215).

For (21.209):

$$\begin{aligned}
\mathbb{E} [\log p(\mathbf{x}|\mathbf{z}, \mu, \Lambda)] &= \mathbb{E}_{q(\mathbf{z})q(\mu, \Lambda)} [\log p(\mathbf{x}|\mathbf{z}, \mu, \Lambda)] \\
&= \sum_n \sum_k \mathbb{E}_{q(\mathbf{z})q(\mu, \Lambda)} \left[z_k \left(-\frac{D}{2} \log 2\pi + \frac{1}{2} \log |\Lambda_k| - \frac{1}{2} (\mathbf{x}_n - \mu_k)^T \Lambda_k (\mathbf{x}_n - \mu_k) \right) \right].
\end{aligned}$$

Plugging (21.131) and (21.132) into this form and we ends up with (21.209).

For (21.210):

$$\begin{aligned}
\mathbb{E} [\log p(\mathbf{z}|\pi)] &= \mathbb{E}_{q(\mathbf{z})q(\pi)} [\log p(\mathbf{z}|\pi)] \\
&= \mathbb{E}_{q(\mathbf{z})q(\pi)} \left[\log \prod_{n=1}^N \prod_{k=1}^K \pi_k^{z_{nk}} \right] \\
&= \sum_{n=1}^N \sum_{k=1}^K \mathbb{E}_{q(\mathbf{z})q(\pi)} [z_{nk} \log \pi_k] \\
&= \sum_{n=1}^N \sum_{k=1}^K \mathbb{E}_{q(\mathbf{z})} [z_{nk}] \cdot \mathbb{E}_{q(\pi)} [\log \pi_k] \\
&= \sum_{n=1}^N \sum_{k=1}^K r_{nk} \log \bar{\pi}_k,
\end{aligned}$$

where we used the fact the the expectation of the product of independent random variables is the product of their expectations. The notations follows (21.129).

For (21.211):

$$\begin{aligned}
\mathbb{E}[\log p(\pi)] &= \mathbb{E}_{q(\pi)}[\log p(\pi)] \\
&= \mathbb{E}_{q(\pi)}[\log(\text{const} \cdot \prod_{k=1}^K \pi_k^{\alpha_0-1})] \\
&= \ln \text{const} + (\alpha_0 - 1) \sum_{k=1}^K \log \bar{\pi}_k,
\end{aligned}$$

plugging (21.216) finishes the proof.

For (21.212):

$$\begin{aligned}
\mathbb{E}[\log p(\mu, \Lambda)] &= \mathbb{E}_{q(\mu, \Lambda)}[\log p(\mu, \Lambda)] \\
&= \mathbb{E}_{q(\mu, \Lambda)} \left[\log \prod_{k=1}^K \text{Wi}(\Lambda_k | L_0, v_0) \cdot \mathcal{N}(\mu_k | m_0, (\beta_0 \Lambda_k)^{-1}) \right] \\
&= \sum_{k=1}^K \mathbb{E}_{q(\mu, \Lambda)} \left[\log C + \frac{1}{2}(v_0 - D - 1) \log |\Lambda_k| - \frac{1}{2} \text{tr} \{ \Lambda_k L_0^{-1} \} \right. \\
&\quad \left. - \frac{D}{2} \log 2\pi - \frac{1}{2} \log |\beta_0 \Lambda_k| - \frac{1}{2} (\mu_k - m_0)^T (\beta_0 \Lambda_k) (\mu_k - m_0) \right].
\end{aligned}$$

Where we have used (21.131) to expand the expected value of the quadratic form and used the fact that the mean of a Wi distribution is $v_k L_k$.

For (21.213):

$$\begin{aligned}
\mathbb{E}[\log q(\mathbf{z})] &= \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \\
&= \mathbb{E}_{q(\mathbf{z})} \left[\sum_i \sum_k z_{ik} \log r_{ik} \right] \\
&= \sum_i \sum_k \mathbb{E}_{q(\mathbf{z})}[z_{ik}] \cdot \log r_{ik} \\
&= \sum_i \sum_k r_{ik} \log r_{ik}.
\end{aligned}$$

We only have to recall (21.124).

For (21.214):

$$\begin{aligned}
\mathbb{E}[\log q(\pi)] &= \mathbb{E}_{q(\pi)}[\log q(\pi)] \\
&= \mathbb{E}_{q(\pi)} \left[\log C + \sum_{k=1}^K (\alpha_k - 1) \log \pi_k \right] \\
&= \log C + \sum_k (\alpha_k - 1) \log \bar{\pi}_k.
\end{aligned}$$

Finally, for (21.215):

$$\begin{aligned}\mathbb{E}[\log q(\mu, \Lambda)] &= \mathbb{E}_{q(\mu, \Lambda)}[\log q(\mu, \Lambda)] \\ &= \sum_k \mathbb{E}_{q(\mu, \Lambda)}[\log q(\Lambda_k) - \frac{D}{2} \log 2\pi + \frac{1}{2} \log |\beta_k \Lambda_k| \\ &\quad - \frac{1}{2} (\mu_k - m_k)^T (\beta_k \Lambda_k) (\mu_k - m_k)].\end{aligned}$$

Using (21.132) to expand the quadratic gives $\mathbb{E}[(\mu_k - m_k)^T (\beta_k \Lambda_k) (\mu_k - m_k)] = D$.

21.5 Derivation of $\mathbb{E}[\log \pi_k]$ under a Dirichlet distribution

Dirichlet distribution is an exponential family distribution with:

$$\begin{aligned}\phi(\pi) &= (\log \pi_1, \log \pi_2, \dots, \log \pi_K)^T, \\ \theta &= \alpha.\end{aligned}$$

Its cumulant function is:

$$A(\alpha) = \log B(\alpha) = \sum_{i=1}^K \log \Gamma(\alpha_i) - \log \Gamma\left(\sum_{i=1}^K \alpha_i\right).$$

Therefore we have:

$$\mathbb{E}[\log \pi_k] = \frac{\partial A(\alpha)}{\partial \alpha_k} = \frac{\Gamma'(\alpha_k)}{\Gamma(\alpha_k)} - \frac{\Gamma'\left(\sum_{i=1}^K \alpha_k\right)}{\Gamma\left(\sum_{i=1}^K \alpha_k\right)} = \psi(\alpha_k) - \psi\left(\sum_{i=1}^K \alpha_i\right).$$

Take exponential on both sides yields the desired form:

$$\exp(\mathbb{E}[\log \pi_k]) = \exp\left(\psi(\alpha_k) - \psi\left(\sum_{i=1}^K \alpha_k\right)\right) = \frac{\exp(\psi(\alpha_k))}{\exp\left(\psi\left(\sum_{i=1}^K \alpha_i\right)\right)}.$$

21.6 Alternative derivation of the mean field updates for the Ising model

The energy for the Ising model is given by:

$$E(\mathbf{x}) = - \sum_{i=1}^D \sum_{j \in \text{nbr}_i} W_{ij} x_i x_j - \sum_{i=1}^D L_i(x_i),$$

where the first term corresponds to the interaction field and the second to the external field. The variational free energy (VFE) is the expectation of the energy after replacing each x_i with a Bernoulli random variable whose mean is μ_i . The dependency of the VEF on μ_i is:

$$-\mu_i \sum_{j \in \text{nbr}_i} W_{ij} \mu_j - \frac{1 - \mu_i}{2} L_i^+ - \frac{\mu_i - 1}{2} L_i^-,$$

from which we can derive the posterior margin as well as the new mean, which are identical to (21.46)-(21.48). Thus the update (21.50) remains the same.

21.7 Forwards vs reverse KL divergence

To formulate $\mathbb{KL}(p||q)$ with $q(x, y) = q(x)q(y)$, we have:

$$\begin{aligned} \mathbb{KL}(p(x, y)||q(x, y)) &= \mathbb{E}_{p(x, y)} \left[\log \frac{p(x, y)}{q(x, y)} \right] \\ &= \sum_{x, y} p(x, y) \log p(x, y) - \sum_{x, y} p(x, y) \log q(x) - \sum_{x, y} p(x, y) \log q(y) \\ &= \sum_{x, y} p(x, y) \log p(x, y) - \sum_x \left(\sum_y p(x, y) \right) \log q(x) - \sum_y \left(\sum_x p(x, y) \right) \log q(y) \\ &= \text{const} + \mathbb{KL}(p(x)||q(x)) + \mathbb{KL}(p(y)||q(y)). \end{aligned}$$

The non-negativity of KL divergence implies that we should set $q(x) = p(x)$ and $q(y) = p(y)$.

For the tabular case, let $q = \{q_{ij}\}_{i,j=1}^4$ be the variational distribution s.t.:

$$q(x, y) = q_x(x) \cdot q_y(y),$$

with six free parameters: $q_x(1)$, $q_x(2)$, $q_x(3)$, $q_y(1)$, $q_y(2)$ and $q_y(3)$. The reverse KL is:

$$\mathbb{KL}(q||p) = \sum_{x, y} q(x, y) \log p(x, y) - H_q(x, y).$$

Since the decomposition assumes the independence between x and y , we have:

$$H_q(x, y) = H_q(x) + H_q(y).$$

Meanwhile, the symmetry within the table indicates that the roles of $q_x(1)$ and $q_x(2)$ are interchangeable, so are $q_x(3)$ and $q_x(4)$, $q_y(1)$ and $q_y(2)$, $q_y(3)$ and $q_y(4)$. This observation, together with the property of the entropy (and the Cauchy inequality), implies that q_x takes the form:

$$(\alpha, \alpha, \frac{1}{2} - \alpha, \frac{1}{2} - \alpha), \alpha \in [0, \frac{1}{2}].$$

So is q_y , hence the KL takes the value:

$$\min_{\alpha \in [0, \frac{1}{2}]} \left\{ -16\alpha^2 + 4\alpha - 2 \cdot H \left(\alpha, \alpha, \frac{1}{2} - \alpha, \frac{1}{2} - \alpha \right) \right\},$$

where α takes some value between $\frac{1}{8}$ and $\frac{1}{4}$. In this case the minimum is not necessarily unique since the function to be minimized is the sum of a convex and a concave function w.r.t. α .

21.8 Derivation of the structured mean field updates for FHMM

According to the conclusion from mean-field variational methods, we resort to the moment matching:

$$E(\mathbf{x}_m) = \mathbb{E}_{q/m}[E(\bar{p}(\mathbf{x}_m))],$$

where the l.r.s. equals:

$$- \sum_{t=1}^T \mathbf{x}_{t,m}^T \tilde{\xi}_{t,m},$$

and the r.h.s. is:

$$\frac{1}{2} \mathbb{E} \left[\sum_{t=1}^T \left(\mathbf{y}_t - \sum_{l \neq m}^M \mathbf{W}_l \mathbf{x}_{t,m} \right)^T \Sigma^{-1} \left(\mathbf{y}_t - \sum_{l \neq m}^M \mathbf{W}_l \mathbf{x}_{t,m} \right) \right] + \text{const.}$$

From which we have dropped the irrelevant terms. Calibrating the coefficients for \mathbf{x}_{tm} yields (21.62)-(21.64).

Comparing the coefficient of $x_{t,m,k}$ (i.e. setting $x_{t,m,k}$ to 1) ends in:

$$\tilde{\epsilon}_{t,m,k} = W_m^T \Sigma^{-1} (\mathbf{y}_t - \sum_{l \neq m} W_l \mathbb{E}[\mathbf{x}_{t,l}]) - \frac{1}{2} (W_m^T \Sigma^{-1} W_m)_{k,k}$$

21.9 Variational EM for binary FA with sigmoid link

We begin with the likelihood:

$$p(\mathbf{Z}, \mathbf{X} | \mathbf{W}) = \prod_{n=1}^N \prod_{j=1}^D \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n)^{x_{nj}} (1 - \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n))^{(1-x_{nj})}.$$

The prior for the hidden variables is assumed to be:

$$\forall n, p(\mathbf{z}_n) = \mathcal{N}(\mathbf{0}, \mathbf{I}).$$

Assume the factorized variational distribution:

$$p(\mathbf{W}, \mathbf{Z} | \mathbf{X}) \approx q(\mathbf{W}) \prod_{n=1}^N q(\mathbf{z}_n).$$

For the variational E-step, our goal is to match the logarithm of the variational distribution on the hidden variables:

$$\log q(\mathbf{z}),$$

with:

$$\begin{aligned} \mathbb{E}_{q(\mathbf{W})} [\log p(\mathbf{Z}, \mathbf{X}, \mathbf{W})] &= \mathbb{E}_{q(\mathbf{W})} \left[\sum_n \sum_j x_{nj} \log \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n) + (1 - x_{nj}) \log(1 - \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n)) \right] \\ &= \sum_{n,j} \mathbb{E}_{q(\mathbf{W})} \left[x_{nj} \log \frac{\text{sigm}(\mathbf{w}_j^T \mathbf{z}_n)}{1 - \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n)} + \log(1 - \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n)) \right] \\ &= \sum_{n,j} \mathbb{E}_{q(\mathbf{W})} [x_{nj} \cdot \mathbf{w}_j^T \mathbf{z}_n + \log(1 - \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n))]. \end{aligned}$$

We can see that this form cannot painlessly reduce to an exponential family, hence approximation needs to be conducted to transfer $\log(1 - \text{sigm}(\mathbf{w}_j^T \mathbf{z}_n))$ to a linear function of \mathbf{z}_n and optionally $\mathbf{z}_n^T \mathbf{z}_n$ (e.g., the Laplace approximation). Then we can see that $\mathbb{E}_{q(\mathbf{W})} [\log p(\mathbf{Z}, \mathbf{X}, \mathbf{W})]$ is a quadratic function in \mathbf{z}_n , hence the E-step reduces $q(\mathbf{z})$ to a Gaussian.

For the variational M-step:

$$\mathbb{E}_{q(\mathbf{Z})} [\log p(\mathbf{Z}, \mathbf{X}, \mathbf{W})]$$

can again be approximated as a quadratic function w.r.t. \mathbf{W} , where the expectation of \mathbf{z} shall be replaced by their counterpart in the E-step.

21.10 VB for binary FA with probit link

To attack this question, we assume the following variational distribution:

$$q(\mathbf{X}, \mathbf{Z}, \mathbf{W}) = \left(\prod_{l=1}^L q(\mathbf{w}_l) \right) \cdot \left(\prod_{i=1}^N q(\mathbf{x}_i) q(z_i) \right).$$

Since we have three independent family of variables to be estimated, this pseudo VB procedure consists of three steps (not simple an E-step and an M-step). We assume $L = 1$ w.l.o.g.

For the first step, we update for \mathbf{x}_i , this is done by collecting terms relevant to \mathbf{x}_i from the expectation of the log likelihood w.r.t. $q(\mathbf{w}, \mathbf{Z})$:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{w}, \mathbf{Z})} [\log p(\mathbf{x}_i, \mathbf{w}, z_i, y_i)] \\ &= \mathbb{E}_{q(\mathbf{w})q(z_i)} [\log p(\mathbf{x}_i) + \log p(z_i | \mathbf{x}_i, \mathbf{w})] + \text{const} \\ &= \mathbb{E}_{q(\mathbf{w})q(z_i)} \left[\log p(\mathbf{x}_i) - \frac{(z_i - \mathbf{w}^T \mathbf{x}_i)^2}{2} \right] + \text{const}. \end{aligned}$$

From which we observe that $q(\mathbf{x}_i)$ takes the form of a Gaussian, the update for its precision is:

$$\Lambda_{\mathbf{x}_i} \rightarrow \Lambda_{\mathbf{x}_i} + \mathbb{E}_{q(\mathbf{w})} [\mathbf{w} \mathbf{w}^T].$$

For its mean:

$$\mu_{\mathbf{x}_i} \rightarrow \Lambda_{\mathbf{x}_i}^{-1} (\Lambda_{\mathbf{x}_i} \mu_{\mathbf{x}_i} + \mathbb{E}[\mathbf{w}] \mathbb{E}[z_i]).$$

In the second step, we update \mathbf{w} , note that:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{X}, \mathbf{Z})} [\log p(\mathbf{X}, \mathbf{w}, \mathbf{Z}, \mathbf{Y})] \\ &= \mathbb{E}_{q(\mathbf{X})q(\mathbf{Z})} [\log p(\mathbf{w}) + \sum_i \log p(z_i | \mathbf{x}_i, \mathbf{w})] + \text{const} \\ &= \mathbb{E}_{q(\mathbf{X})q(\mathbf{Z})} \left[\log p(\mathbf{w}) - \sum_i \frac{(z_i - \mathbf{w}^T \mathbf{x}_i)^2}{2} \right] + \text{const}. \end{aligned}$$

Therefore the variational distribution for \mathbf{w} should better be a Gaussian, whose update takes the form:

$$\begin{aligned} \Lambda_{\mathbf{w}} &\rightarrow \Lambda_{\mathbf{w}} + \sum_i \mathbf{x}_i \mathbf{x}_i^T, \\ \mu_{\mathbf{w}} &\rightarrow \Lambda_{\mathbf{w}}^{-1} (\Lambda_{\mathbf{w}} \mu_{\mathbf{w}} + \sum_i \mathbb{E}[z_i] \mathbb{E}[\mathbf{x}_i]). \end{aligned}$$

Finally, we update for z_i :

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}, \mathbf{w})} [\log p(\mathbf{x}_i, \mathbf{w}, z_i, y_i)] \\ &= \mathbb{E}_{q(\mathbf{x}_i)q(\mathbf{w})} \left[-\frac{(z_i - \mathbf{w}^T \mathbf{x}_i)^2}{2} \right] + \text{const}, \end{aligned}$$

moreover, the domain for z_i is confined in $(0, \infty)$ if $y_i = 1$ and $(-\infty, 0]$ otherwise. Therefore the variational distribution for z_i is truncated, yet it has a quadratic form w.r.t. z_i in its exponential, hence it is a truncated Gaussian as what has to be proven. The variance for this distribution (if not truncated) is uniformly unity, and its mean is:

$$\mathbb{E}[\mathbf{w}]^T \mathbb{E}[\mathbf{x}_i].$$

All three update steps form a fixed-point fashion learning.

22 More variational inference*

The exercises are from the textbook *Probabilistic Graphical Models: Principles and Techniques*, Koller and Friedman. This textbook provides a solid and comprehensive introduction to PGM. One needs some graph knowledge besides probability theory before entering that book. References for exercises in this chapter are left to that textbook.

23 Monte Carlo inference

23.1 Sampling from a Cauchy

Recall that the p.d.f. of a standard Cauchy $\mathcal{T}(x|0, 1, 1)$ is:

$$\mathcal{T}(x|0, 1, 1) = \frac{1}{Z} \cdot \frac{1}{1+x^2},$$

where:

$$Z = \int_{-\infty}^{\infty} \frac{1}{1+x^2} dx = \arctan x|_{-\infty}^{\infty} = \pi.$$

Thus the c.d.f. is:

$$F(x) = \int_{-\infty}^x \frac{1}{\pi(1+t^2)} dt = \frac{\arctan x + \frac{\pi}{2}}{\pi},$$

whose inverse is:

$$\tan \left[\pi \left(p - \frac{1}{2} \right) \right].$$

So we can sample p from $\text{Uniform}[-\frac{1}{2}, \frac{1}{2}]$ and transform it by $\tan(\pi p)$, the result would follow a standard Cauchy.

23.2 Rejection sampling from a Gamma using a Cauchy proposal

The p.d.f. of a Cauchy is:

$$\mathcal{T}(x|\mu, \sigma^2, 1) = \frac{1}{Z} \cdot \frac{1}{1 + \left(\frac{x-\mu}{\sigma} \right)^2},$$

where $Z = \pi\sigma$.

That of a Gamma is:

$$\text{Ga}(x|a, b) = \frac{b^a}{\Gamma(a)} x^{a-1} e^{-xb}.$$

Finally, we determine M from:

$$M = \max_x \left\{ \frac{\text{Ga}(x|a, b)}{\mathcal{T}(x|\mu, \sigma^2, 1)} \right\},$$

where x takes value from the zeros of the polynomial:

$$f(x) = (a-1-bx) \left[1 + \left(\frac{x-\mu}{\sigma} \right)^2 \right] + \frac{2x(x-\mu)}{\sigma^2},$$

and can be solved analytically. The optimal value for μ and σ^2 should be:

$$\arg \min_{\mu, \sigma^2} \{M(\mu, \sigma^2 | a, b)\}$$

in order to increase the rate of acceptance.

23.3 Optimal proposal for particle filtering with linear-Gaussian measurement model

The model is specified by:

$$p(\mathbf{z}_t | \mathbf{z}_{t-1}) = \mathcal{N}(\mathbf{z}_t | f(\mathbf{z}_{t-1}), \mathbf{Q}_{t-1}),$$

$$p(\mathbf{y}_t | \mathbf{z}_t) = \mathcal{N}(\mathbf{y}_t | \mathbf{H}_t \mathbf{z}_t, \mathbf{R}_t).$$

For $p(\mathbf{z}_t | \mathbf{z}_{t-1}, \mathbf{y}_t)$, we have:

$$\begin{aligned} p(\mathbf{z}_t | \mathbf{z}_{t-1}, \mathbf{y}_t) &= \frac{p(\mathbf{z}_t, \mathbf{z}_{t-1}, \mathbf{y}_t)}{p(\mathbf{z}_{t-1}, \mathbf{y}_t)} \\ &\propto p(\mathbf{z}_t | \mathbf{z}_{t-1}) \cdot p(\mathbf{y}_t | \mathbf{z}_t) \\ &\propto \exp \left\{ -\frac{1}{2} (\mathbf{z}_t - f(\mathbf{z}_{t-1}))^T \mathbf{Q}_{t-1}^{-1} (\mathbf{z}_t - f(\mathbf{z}_{t-1})) \right\} \\ &\quad \cdot \exp \left\{ -\frac{1}{2} (\mathbf{y}_t - \mathbf{H}_t \mathbf{z}_t)^T \mathbf{R}_t^{-1} (\mathbf{y}_t - \mathbf{H}_t \mathbf{z}_t) \right\}, \end{aligned}$$

from which we can readily read that the posterior covariance and mean for \mathbf{z}_t are:

$$\begin{aligned} \Sigma &= (\mathbf{Q}_t^{-1} + \mathbf{H}_t^T \mathbf{R}_t^{-1} \mathbf{H}_t)^{-1}, \\ \mu &= \Sigma (\mathbf{Q}_t^{-1} f(\mathbf{z}_{t-1}) + \mathbf{H}_t^T \mathbf{R}_t^{-1} \mathbf{y}_t). \end{aligned}$$

On the other hand:

$$\begin{aligned} p(\mathbf{y}_t | \mathbf{z}_{t-1}) &= \int p(\mathbf{y}_t, \mathbf{z}_t | \mathbf{z}_{t-1}) d\mathbf{z}_t \\ &= \int p(\mathbf{y}_t | \mathbf{z}_t) \cdot p(\mathbf{z}_t | \mathbf{z}_{t-1}) d\mathbf{z}_t \\ &\propto \exp \left\{ -\frac{1}{2} \mathbf{y}_t^T \mathbf{R}_t^{-1} \mathbf{y}_t \right\} \\ &\quad \cdot \exp \left\{ \frac{1}{2} (f^T \mathbf{Q}_t^{-1} + \mathbf{y}_t^T \mathbf{R}_t^{-1} \mathbf{H}_t) \Sigma (\mathbf{Q}_t^{-1} f + \mathbf{H}_t^T \mathbf{R}_t^{-1} \mathbf{y}_t) \right\}, \end{aligned}$$

where Σ is the posterior covariance for \mathbf{z}_t . Therefore the posterior covariance and mean for \mathbf{y}_t given \mathbf{z}_{t-1} are:

$$\Sigma' = (\mathbf{R}_t^{-1} - \mathbf{R}_t^{-1} \mathbf{H}_t \Sigma \mathbf{H}_t^T \mathbf{R}_t^{-1})^{-1},$$

and finally:

$$\mu' = \Sigma' \mathbf{R}_t^{-1} \mathbf{H}_t \Sigma \mathbf{Q}_t^{-1} f(\mathbf{z}_{t-1}).$$

24 Markov chain Monte Carlo (MCMC) inference

24.1 Gibbs sampling from a 2D Gaussian

The results follow (4.69):

$$p(x_1|x_2) = \mathcal{N}(x_1|\mu, \sigma^2),$$

where:

$$\sigma^2 = 1 - (-0.5) \times (-0.5) = 0.75,$$

$$\mu = 1 + (-0.5) \times (x_2 - 1).$$

The symmetric result holds for $p(x_2|x_1)$.

24.2 Gibbs sampling for a 1D Gaussian mixture model

The likelihood for this model is (assume that there are K Gaussian components, we discuss this problem for MVN):

$$\begin{aligned} p(\mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma) &= p(\pi) \cdot \left(\prod_{k=1}^K p(\mu_k) \cdot p(\Sigma_k) \right) \cdot \left(\prod_{n=1}^N p(\mathbf{z}_n|\pi) \cdot p(\mathbf{x}_n|\mathbf{z}_n, \mu, \Sigma) \right) \\ &= \frac{1}{B(\alpha)} \cdot \prod_{k=1}^K \pi_k^{\alpha_k-1} \\ &\quad \cdot \prod_{k=1}^K \mathcal{N}(\mu_k|\mathbf{m}_0, \mathbf{V}_0) \text{IW}(\Sigma_k|\mathbf{S}_0, v_0) \\ &\quad \cdot \prod_{n=1}^N \prod_{k=1}^K (\pi_k \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k))^{z_{nk}}. \end{aligned}$$

For Gibbs sampling, we derive the conditional distribution of \mathbf{Z} , π , μ and Σ conditioned on all other variables.

For the latent variables:

$$\begin{aligned} p(z_{nk} = 1|\mathbf{X}, \pi, \mu, \Sigma) &= \frac{p(z_{nk} = 1, \mathbf{X}, \pi, \mu, \Sigma)}{p(\mathbf{X}, \pi, \mu, \Sigma)} \\ &\propto p(z_{nk} = 1|\pi) \cdot p(\mathbf{x}_n|z_{nk} = 1, \mu, \Sigma) \\ &= \pi_k \cdot \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k). \end{aligned}$$

For the weights:

$$\begin{aligned}
p(\pi|\mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma) &= \frac{p(\pi, \mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma)}{p(\mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma)} \\
&\propto p(\pi) \cdot p(\mathbf{Z}|\pi) \\
&\propto \left(\prod_{k=1}^K \pi_k^{\alpha_k - 1} \right) \cdot \left(\prod_{n=1}^N \prod_{k=1}^K \pi_k^{z_{nk}} \right) \\
&= \prod_{k=1}^K \pi_k^{\alpha_k + (\sum_{n=1}^N z_{nk}) - 1}.
\end{aligned}$$

Hence the conditional distribution for π follows (24.11).

For the means:

$$\begin{aligned}
p(\mu_k|\mathbf{X}, \mathbf{Z}, \pi, \mu_{-k}, \Sigma) &= \frac{p(\mu_k, \mathbf{X}, \mathbf{Z}, \pi, \mu_{-k}, \Sigma)}{p(\mathbf{X}, \mathbf{Z}, \pi, \mu_{-k}, \Sigma)} \\
&\propto p(\mu_k) \cdot p(\mathbf{X}|\mu_k, \mathbf{Z}, \Sigma, \mu_{-k}) \\
&= \mathcal{N}(\mu_k|\mathbf{m}_0, \mathbf{V}_0) \cdot \prod_{n=1}^N (\pi_k \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma_k))^{z_{nk}}.
\end{aligned}$$

At this point it is better to observe in the logarithm perspective:

$$\begin{aligned}
\log p(\mu_k|\mathbf{X}, \mathbf{Z}, \pi, \mu_{-k}, \Sigma) &= -\frac{1}{2}(\mu_k - \mathbf{m}_0)^T \mathbf{V}_0^{-1}(\mu_k - \mathbf{m}_0) \\
&\quad + \sum_{n=1}^N -\frac{z_{nk}}{2}(\mu_k - \mathbf{x}_n)^T \Sigma_k^{-1}(\mu_k - \mathbf{x}_n) + \text{const},
\end{aligned}$$

from which we observe that the coefficient for $\mu^T \mu$ and μ are:

$$\begin{aligned}
&-\frac{1}{2} \left(\mathbf{V}_0^{-1} + \Sigma_k^{-1} \sum_{n=1}^N z_{nk} \right), \\
&\mathbf{V}_0^{-1} \mathbf{m}_0 + \sum_{n=1}^N z_{nk} \Sigma_k^{-1} \mathbf{x}_n.
\end{aligned}$$

Therefore the conditional distribution for μ_k is a Gaussian with covariance:

$$\mathbf{V}_k = \left(\mathbf{V}_0^{-1} + \Sigma_k^{-1} \sum_{n=1}^N z_{nk} \right)^{-1}$$

and mean:

$$\mathbf{m}_k = \mathbf{V}_k \left(\mathbf{V}_0^{-1} \mathbf{m}_0 + \sum_{n=1}^N z_{nk} \Sigma_k^{-1} \mathbf{x}_n \right).$$

Finally, for the covariance:

$$\begin{aligned} p(\Sigma_k | \mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma_{-k}) &= \frac{p(\Sigma_k, \mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma_{-k})}{p(\mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma_{-k})} \\ &\propto p(\Sigma_k) \cdot p(\mathbf{X} | \Sigma_k, \mathbf{Z}, \mu, \Sigma_{-k}). \end{aligned}$$

This is tantamount to:

$$\begin{aligned} \log p(\Sigma_k | \mathbf{X}, \mathbf{Z}, \pi, \mu, \Sigma_{-k}) &= -\frac{v_0}{2} \cdot \log |\Sigma_k| - \frac{1}{2} \text{tr}[\mathbf{S}_0^{-1} \Sigma_k^{-1}] \\ &\quad + \sum_{n=1}^N z_{nk} \left(-\frac{1}{2} \cdot \log |\Sigma_k| - \frac{1}{2} (\mathbf{x}_n - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_n - \mu_k) \right) + \text{const} \\ &= -\frac{v_0 + \sum_{n=1}^N z_{nk}}{2} \cdot \log |\Sigma_k| \\ &\quad - \frac{1}{2} \text{tr} \left[\left(\mathbf{S}_0^{-1} + \sum_{n=1}^N z_{nk} (\mathbf{x}_n - \mu_k)^T (\mathbf{x}_n - \mu_k) \right) \Sigma_k^{-1} \right] + \text{const}. \end{aligned}$$

Hence the conditional distribution of Σ_k is the inverse-Wishart distribution with:

$$\begin{aligned} \mathbf{v}_k &= v_0 + \sum_{n=1}^N z_{nk}, \\ \mathbf{S}_k &= \left(\mathbf{S}_0^{-1} + \sum_{n=1}^N z_{nk} (\mathbf{x}_n - \mu_k)^T (\mathbf{x}_n - \mu_k) \right)^{-1}. \end{aligned}$$

The difference from (24.18) is in the definition of \mathbf{S}_0 .

24.3 Gibbs sampling from the Potts model

Modify by yourself.

24.4 Full conditionals for hierarchical model of Gaussian means

Recall that the likelihood for the Gaussian-Gaussian model is:

$$p(\theta, \mathcal{D} | \mu, \tau^2, \sigma^2) = \left(\prod_{j=1}^D \mathcal{N}(\theta_j | \mu, \tau^2) \right) \cdot \left(\prod_{j=1}^D \prod_{i=1}^{N_j} \mathcal{N}(x_{ij} | \theta_j, \sigma^2) \right).$$

We proceed from the conditional distribution of μ :

$$\begin{aligned} p(\mu|\theta, \mathcal{D}, \tau^2, \sigma^2) &= \frac{p(\mu, \theta, \mathcal{D}, \tau^2, \sigma^2)}{p(\theta, \mathcal{D}, \tau^2, \sigma^2)} \\ &\propto p(\mu) \cdot p(\theta|\mu, \tau^2). \end{aligned}$$

Taking exponential yields:

$$\log p(\mu|\theta, \mathcal{D}, \tau^2, \sigma^2) = -\frac{(\mu - \mu_0)^2}{2\gamma_0^2} - \sum_{j=1}^D \frac{(\mu - \theta_j)^2}{2\tau^2}.$$

Hence the conditional variance and mean are:

$$\begin{aligned} \gamma_c &= \left(\frac{1}{\gamma_0^2} + \frac{D}{\tau^2} \right)^{-1}, \\ \mu_c &= \gamma_c \left(\frac{\mu_0}{\gamma_0^2} + \sum_{j=1}^D \frac{\theta_j}{\tau^2} \right), \end{aligned}$$

which are tantamount to (24.115).

For the centroids:

$$\begin{aligned} p(\theta_k|\mathcal{D}, \theta_{-k}, \mu, \tau^2, \sigma^2) &= \frac{p(\theta_k, \mathcal{D}, \theta_{-k}, \mu, \tau^2, \sigma^2)}{p(\mathcal{D}, \theta_{-k}, \mu, \tau^2, \sigma^2)} \\ &\propto p(\theta_k|\mu, \tau^2) \cdot p(\mathcal{D}|\theta_k, \theta_{-k}, \sigma^2) \\ &\propto \exp \left\{ -\frac{(\theta_k - \mu)^2}{2\tau^2} \right\} \\ &\quad \cdot \exp \left\{ -\sum_{i=1}^{N_k} \frac{(\theta_k - x_{ik})^2}{2\sigma^2} \right\}. \end{aligned}$$

We ends up with a Gaussian with variance:

$$\tau_k^2 = \left(\frac{1}{\tau^2} + \frac{N_k}{\sigma^2} \right)^{-1},$$

and mean:

$$\mu_j = \tau_k^2 \left(\frac{\mu}{\tau^2} + \sum_{i=1}^{N_k} \frac{x_{ij}}{\sigma^2} \right).$$

For the variance on the centroids:

$$\begin{aligned}
p(\tau^2|\theta, \mathcal{D}, \mu, \sigma^2) &= \frac{p(\tau^2, \theta, \mathcal{D}, \mu, \sigma^2)}{p(\theta, \mathcal{D}, \mu, \sigma^2)} \\
&\propto p(\tau^2) \cdot \prod_{j=1}^D p(\theta_j|\mu, \tau^2) \\
&\propto (\tau^2)^{-\frac{\eta_0}{2}-1} \cdot \exp\left\{-\frac{\eta_0\tau_0^2}{2\tau^2}\right\} \\
&\quad \cdot (\tau^2)^{-\frac{D}{2}} \cdot \exp\left\{-\frac{\sum_{j=1}^D (\theta_j - \mu)^2}{2\tau^2}\right\}.
\end{aligned}$$

Therefore the parameters for the conditionl IG distribution are:

$$\begin{aligned}
\eta_c &= \eta_0 + D, \\
\tau_c^2 &= \frac{\eta_0\tau_0^2 + \sum_{j=1}^D (\theta_j - \mu)^2}{\eta_c}.
\end{aligned}$$

Finally, we handle the variance on the data:

$$\begin{aligned}
p(\sigma^2|\theta, \mathcal{D}, \mu, \tau^2) &\propto p(\sigma^2) \cdot p(\mathcal{D}|\theta, \sigma^2) \\
&\propto (\sigma^2)^{-\frac{v_0}{2}-1} \cdot \exp\left\{-\frac{v_0\sigma_0^2}{2\sigma^2}\right\} \\
&\quad \cdot (\sigma^2)^{-\frac{N}{2}} \cdot \exp\left\{\sum_{j=1}^D \sum_{i=1}^{N_j} -\frac{(\theta_j - x_{ij})^2}{2\sigma^2}\right\},
\end{aligned}$$

where $N = \sum_{j=1}^D N_j$. This is sufficient for (24.118).

24.5 Gibbs sampling for robust linear regression with a Student t likelihood

Recall that in robust linear regression, each sample y_i is assigned an extra latente variable z_i as the estimation of the noise. The complete likelihood is:

$$\begin{aligned}
p(\mathbf{Y}, \mathbf{X}, \mathbf{Z}, \mathbf{w}, \sigma^2, v) &= p(\mathbf{X}, \mathbf{w}, \sigma^2, v) \cdot p(\mathbf{Z}|v) \cdot p(\mathbf{Y}|\mathbf{Z}, \mathbf{X}, \mathbf{w}, \sigma^2) \\
&\propto p(\sigma^2) \cdot \left(\prod_{n=1}^N z_n^{\frac{v}{2}-1} e^{-\frac{z_n \cdot v}{2}}\right) \cdot \left(\prod_{n=1}^N (\sigma^2)^{-\frac{1}{2}} z_n^{\frac{1}{2}} e^{-\frac{z_n(y_n - \mathbf{w}^T \mathbf{x}_n)^2}{2\sigma^2}}\right).
\end{aligned}$$

From which we can read that the conditional distribution for z_n takes the form:

$$z_n^{\frac{v+1}{2}-1} e^{-z_n(\frac{v}{2} + \frac{(y_n - \mathbf{w}^T \mathbf{x}_n)^2}{2\sigma^2})},$$

hence is again a Gamma distribution.

For the variance σ^2 , its conditional posterior is:

$$p(\sigma^2) \cdot \left(\prod_{n=1}^N (\sigma^2)^{-\frac{1}{2}} z_n^{\frac{1}{2}} e^{-\frac{z_n(y_n - \mathbf{w}^T \mathbf{x}_n)^2}{2\sigma^2}} \right),$$

which can be absorbed into the prior with an Inverse-Gamma form.

Finally, for the weight \mathbf{w} , the conditional posterior takes the Gaussian form, so using a Gaussian conjugate prior is sufficient for Gibbs sampling update.

24.6 Gibbs sampling for probit regression

With:

$$p(\mathbf{Y}, \mathbf{Z}, \mathbf{w} | \mathbf{X}) = p(\mathbf{w}) \cdot p(\mathbf{Z} | \mathbf{X}, \mathbf{w}) \cdot p(\mathbf{Y} | \mathbf{Z}),$$

we can derive all conditional posterior as in exercise 24.5, where the logarithm of $p(\mathbf{Y} | \mathbf{Z})$ truncates the Gaussian for \mathbf{Z} .

24.7 Gibbs sampling for logistic regression with the Student approximation

The complete likelihood with (24.88)-(24.91) is:

$$\begin{aligned} p(\lambda, \mathbf{Z}, \mathbf{w}, \mathbf{Y} | \mathbf{X}, v) &= p(\lambda | v) \cdot p(\mathbf{w}) \cdot p(\mathbf{Z} | \mathbf{X}, \mathbf{w}, \lambda) \cdot p(\mathbf{Y} | \mathbf{Z}) \\ &= \prod_{n=1}^N \text{Ga}(\lambda_n | \frac{v}{2}, \frac{v}{2}) \\ &\quad \cdot p(\mathbf{w}) \cdot \prod_{n=1}^N \mathcal{N}(z_n | \mathbf{w}^T \mathbf{x}_n, \lambda_n^{-1}) \\ &\quad \cdot \prod_{n=1}^N \mathbb{I}[y_n \sim z_n]. \end{aligned}$$

For λ , we have:

$$p(\lambda_n | \dots) \propto \lambda_n^{\frac{v}{2}-1} \cdot \exp \left\{ -\frac{v\lambda_n}{2} \right\} \cdot \lambda_n^{\frac{1}{2}} \cdot \exp \left\{ -\frac{\lambda_n(z_n - \mathbf{w}^T \mathbf{x}_n)^2}{2} \right\}.$$

So the conditional posterior is a Gamma distribution with parameters:

$$\frac{v+1}{2},$$

$$\frac{v}{2} + \frac{(z_n - \mathbf{w}^T \mathbf{x}_n)^2}{2}.$$

For \mathbf{Z} , we have:

$$p(z_n | \dots) \propto \exp \left\{ -\frac{\lambda_n (z_n - \mathbf{w}^T \mathbf{x}_n)^2}{2} \right\},$$

s.t. $y_n \sim z_n$, therefore z_n follows a truncated Gaussian.

Finally, \mathbf{w} follows a Gaussian, as in exercise 24.5 and 24.6.

25 Clustering

Probability is not the only ghost behind clustering analysis. There are other possible models for the membership in clustering. For example, the *fuzzy clustering* models the membership as fuzzy linguistic term. Under this assumption different regularization and normalization conditions are exerted to the memberships, leads to many crucial algorithms for the *fuzzy theory* and *fuzzy systems*. Besides fuzzy clustering, there are other methods such as heuristic clustering (in which each cluster grows in a heuristic manner), possibilistic clustering, etc.

26 Graphical model structure learning

26.1 Causal reasoning in the sprinkler network

For question (a) and (b), a perfect intervention cuts out links to W , so the probability that the sprinkler is on is:

$$0.5 \times 0.9 + 0.5 \times 0.1 = 0.5,$$

so is that in case the sprinkler is off.

For question (c), the probability is 0.1, since interfering with the root node of a PGM affects neither the graphical structure, nor the information propagation within.

27 Latent variable models for discrete data

27.1 Partition function for an RBM

The partition function for a binary RBM takes the form:

$$Z(\theta) = \sum_{\mathbf{v}} \sum_{\mathbf{h}} \exp \{ \mathbf{v}^T \mathbf{W} \mathbf{h} + \mathbf{v}^T \mathbf{b} + \mathbf{h}^T \mathbf{c} \}.$$

For a fixed \mathbf{v} , the terms to be summed up are:

$$\exp \{ \mathbf{v}^T \mathbf{b} \} \cdot \sum_{\mathbf{h}} \exp \{ \mathbf{h}^T \mathbf{u} \},$$

where:

$$\mathbf{u} = \mathbf{W}^T \mathbf{v} + \mathbf{c}.$$

The summation $\sum_{\mathbf{h}} \exp \{ \mathbf{h}^T \mathbf{u} \}$ shall be written as:

$$\begin{aligned} & \exp(0) + \exp(u_1) + \exp(u_2) + \exp(u_2 + u_1) + \cdots \\ &= \exp(0) + \exp(u_1) + \exp(u_2)[\exp(0) + \exp(u_1)] + \cdots, \end{aligned}$$

Hence:

$$\sum_{\mathbf{h}} \exp \{ \mathbf{h}^T \mathbf{u} \} = \prod_{k=1}^K (1 + \exp(u_k)),$$

whose computation cost is $\mathcal{O}(K)$. The overall complexity is thus $\mathcal{O}(K \cdot 2^R)$.

Since the summation order for \mathbf{v} and \mathbf{h} is interchangeable, the complexity is reduced to:

$$\mathcal{O} \left(\min_{R,K} \{ K \cdot 2^R, R \cdot 2^K \} \right).$$

28 Deep learning

This is where people believe that machine learning can lead to artificial intelligence.